

UNIVERSITY OF AMSTERDAM

MSC ARTIFICIAL INTELLIGENCE
MASTER THESIS

A Plug-and-Play Approach to Age-Adaptive Dialogue Generation

by
LENNERT JANSEN
10488952

November 16, 2021

48EC
January 11, 2021 - November 16, 2021

Supervisors:
Dr SANDRO PEZZELLE
Dr RAQUEL FERNÁNDEZ
Dr ARABELLA SINCLAIR

Assessor:
???



INSTITUTE FOR LOGIC, LANGUAGE AND COMPUTATION

Contents

1	Introduction	4
2	Literature Review	7
2.1	Background	7
2.1.1	Language Models	8
2.1.2	(Controllable) Text Generation	9
2.1.3	Dialogue, Dialogue Systems, and Dialogue Generation	9
2.1.4	Dialogue Response Generation as a Language Modeling Task	11
2.1.5	Controllable Dialogue Generation	12
2.1.6	Transformers	12
2.1.7	Causal Language Modeling with Transformers for Dialogue Response Generation	14
2.1.8	Plug-and-Play Language Modeling	15
2.1.9	Language and Age	18
2.2	Related Work	20
2.2.1	Automated Age Detection	20
2.2.2	Controllable Language Generation	21
2.2.3	Text Style Transfer	22
2.2.4	Dialogue Generation	23
2.2.5	Controlled Dialogue Generation	24
3	Experiment 1: Detecting Age-Related Linguistic Patterns in Dialogue and Discourse	27
3.1	Data	28
3.1.1	Dialogue Dataset	28
3.1.2	Discourse Dataset	29
3.2	Methodology and Experimental Setup	30
3.3	Detecting Age-Related Linguistic Patterns in Dialogue	31

3.3.1	Classification Performance on Discourse	32
3.3.2	Classification Performance on Dialogue	32
3.4	Age Detection Analyses	34
3.4.1	Performance Against Topic	34
3.4.2	Comparing Model Predictions	35
3.4.3	Most Informative N-grams	36
4	Experiment 2: Age-Adaptive Dialogue Generation using PPLM	38
4.1	Introduction	38
4.1.1	Recap of previous chapter and Connection to Experiment 1	38
4.1.2	Research Objectives, Hypotheses, and Contributions	40
4.2	Methodology and Experimental Setup	41
4.2.1	Attribute Model Development	42
4.2.2	Prompts	43
4.2.3	Generation	44
4.2.4	Evaluation	45
4.3	Controlled Dialogue Generation Results	47
4.4	Controlled Dialogue Generation Analyses	50
4.4.1	The Relationship between Perplexity and Target Probability	50
4.4.2	The Relationship between Response Length and Various Evaluation Metrics	53
4.4.3	The Effects of Prompt Style on PPLM Control	57
4.4.4	BERT _{FT} Attention Patterns	61
4.4.5	Qualitative Analyses	65
5	Discussion	70
5.1	Summary of Key Findings and Interpretations	70
5.2	Contributions and Implications	72
5.3	Limitations	73
5.4	Future research directions	74
5.5	Ethical and environmental considerations	75
6	Conclusion	76

A Supplementary material	82
A.1 Prompts used for controlled dialogue response generation	82
A.2 Wordlists for BoW-based Approaches	82
A.3 Supplementary Material about the Effects of Prompt Style on PPLM Output	84
A.4 Where to put these?	84
A.5 Age Discrimination on the Imbalanced British National Corpus [L: Do we even need these plots?]	85
A.6 Placeholders for Final CTG Results Tables	87
A.7 CTG Results for Unprompted Setup [L: Redundant?]	89

List of Footnotes

1.1 https://github.com/lennertjansen/msc-ai-thesis	6
2.1 https://jalammar.github.io/illustrated-transformer/	13
3.1 They are left with roughly 511K datapoints after pre-processing, while we experiment with around 677K.	32
3.2 In particular, it represents one of the utterance's topic, i.e., the one most frequently used in the whole data.	34
4.1 https://pytorch.org/docs/stable/generated/torch.optim.Adam.html# torch.optim.Adam	43
4.2 https://www.statista.com/statistics/272365/ age-distribution-of-internet-users-worldwide/	47
4.3 https://www.statista.com/statistics/1125159/ reddit-us-app-users-age/	48
A.1 https://www.statista.com/statistics/1125159/ reddit-us-app-users-age/	90

Chapter 1

Introduction

In recent years, we have witnessed promising advances in natural language processing (NLP) tasks, such as language modeling, reading comprehension, machine translation, controllable text generation, and conversational response generation [Radford et al., 2019, Bahdanau et al., 2015, Dathathri et al., 2020, Madotto et al., 2020]. Vaswani et al. [2017]’s Transformer architecture plays a central role in many of the state of the art (SotA) solutions to these problems. Transformer-based language models (LMs) pre-trained on massive amounts of textual data, most famously OpenAI’s GPT-2 (Generative Pre-trained Transformer-2), have demonstrated their usefulness for several of the aforementioned NLP tasks [Radford et al., 2019]. For instance, controllable text generation and producing dialogue responses have improved greatly because of GPT-based hybrid models.

[L: CTG comes out of the blue in the following paragraph. Introduce it a little bit by describing what it is, and why/how it is an important task.]

- Controllable text generation entails generating text samples that possess a predefined textual property, like having a positive sentiment, or being about a certain topic.
- Controlling more fine-grained linguistic properties, like resemblance of age-specific vernacular, still poses an important, yet unsolved/insufficiently studied (?) challenge.
- Personalized interaction between humans and AI systems is crucial to obtain systems that can be trusted by users and are perceived as natural.
- (Age-)adaptive language generation can be used to personalize AI-powered personal assistants like Siri and Alexa, improving user experience and trust.

- It is important for AI-power conversational agents to be accessible to varying user profiles, rather than targeted at one particular user group.
- In this work, I/we focus on one aspect that may influence successful personalization of conversational agents: user age profile.

Controllable text generation (CTG) aims to enforce abstract properties, like writing style, on the passages being produced. Fine-tuning large-scale LMs for writing-style adaptation is extremely expensive, but Dathathri et al. [2020] and Li et al. [2020] propose methods that both excel at the task, while bypassing significant retraining costs. Dialogue response generation is the task of producing replies to a conversational agent’s prompts, in a manner that is ideally both non-repetitive and relevant to the course of the conversation. With DialoGPT, Zhang et al. [2020] also manage to leverage GPT-2’s powerful fluency for dialogue tasks, by framing them as language modeling tasks where multi-turn dialogue sessions are seen as long texts.

[L: Introduce dialogue response generation a bit more. Also emphasize its importance. And then introduce the combined task and its importance.] A blend of CTG and dialogue response generation, i.e., controllable dialogue response generation, is an interesting and only partially explored route. It ties closely to one of Artificial Intelligence’s long-standing goals of achieving human-like conversation with machines, as humans are known to adapt their language use to the characteristics of their interlocutor [Gallois and Giles, 2015]. Adaptive dialogue generation is difficult due to the challenge of representing traits, like age, gender, or other persona-labeled traits via language expression [Zheng et al., 2019].

In this thesis, I investigate the problem of controllable dialogue generation, with a focus on adapting responses to users’ age. As a preliminary research objective, I aim to study to what extent a classifier can detect age-related linguistic differences in natural language, and which features are most helpful in age-group detection. Do they (i.e., the linguistic or latent features exploited by the classifier) match the age-related informative features reported in previous work? After empirically confirming that speaker age detection is possible, I explore whether large-scale LMs, e.g. GPT-2, can be leveraged for text generation, controlled for age-groups. And what role does the used data play in the differences in output and performance between regular GPT-2 and controllable GPT-2? Finally, my research focuses on the degree to which such a CTG model is successful in generating dialogue that is adaptive w.r.t. age, such that it has a detectable effect on the perception of the user.

[L: Fix this.]The remainder of this thesis is structured as follows: Chapter 2.1 positions the subject of controlled text generation in its theoretical background, and and 2.2 compares it to the most relevant related work. The methodology in Section 3.2 gives detailed explanations of the most important modeling methods and techniques used for this research. [L: The code used to produce the results can be found on GitHub¹. (Is this the best place to mention this?)]

- When introducing your own work and proposing your hypothesis, use the following argument: *This idea that age prediction from text is more challenging than topic or sentiment prediction could be an indication that controlled language generation for age-differences is also a more nuanced problem than topical steered text generation.*

¹<https://github.com/lennertjansen/msc-ai-thesis>

Chapter 2

Literature Review

This chapter is a two-part literature review. The first section, Background or Section 2.1, provides an overview of this thesis' central problem of controllable dialogue generation, and the components involved, i.e., dialogue, (controllable) language generation, dialogue response generation, and age modeling. In the second section, Related work or Section 2.2, I discuss previous approaches, relevant to my work, that have been proposed to tackle each of these components, either separately or jointly. Approaches to different, but strongly related, problems, like text style transfer, are also described in Section 2.2.

2.1 Background

[L: Keep in mind the following distinction between Background and Related Work - The Background section should give an overview of the problem and the components involved: dialogue, language generation, dialogue response generation, age modelling, etc., without focusing on one or the other approach — in Related Work, you describe approaches that have been proposed to tackle each of these components, separately or jointly, and which are related or relevant to your own work for some reason]

We focus on controllable language generation, i.e., endowing automatically produced text with certain desired linguistic characteristics, and apply it to dialogue. This naturally involves having a model for language generation that can be controlled to write texts passages with different linguistic styles. In what follows, we introduce the crucial concepts behind this: models for language generation, methods to control the output, dialogue, and age modeling.

[L: isn't this redundant, given the chapter's introduction?]

2.1.1 Language Models

Generally speaking, language modeling is central to many NLP tasks. A language model (LM) is a probability distribution over words in a sentence or document. Language models are trained to predict the probability of the next word in an sentence, given the preceding sequence of words. The language modeling task is formulated as an unsupervised distribution estimation problem of datapoints $\{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ (e.g., documents), each representing sequences (of e.g., symbols or tokens) of varying lengths $(s_{i,1}, \dots, s_{i,n}), i \in \{1, \dots, N\}$. Note that N denotes the corpus size, and n the sequence length of datapoint i . To avoid cluttered notation, the subscript i will sometimes be omitted when discussing an arbitrary datapoint. The probability distribution over an observation \mathbf{x} (i.e., the joint probability of an ordered sequence) can then be factorized as the product of its constituent conditionals [Radford et al., 2019]:

$$p(\mathbf{x}) = \prod_{j=1}^n p(s_j | s_1, \dots, s_{j-1}). \quad (2.1)$$

This formulation allows language models to detect and learn patterns in language. The learned representations of these patterns can then be used for a plethora of applications, such as classification, and text generation. Moreover, this results in a framework for tractable sampling from the unconditional language model $p(\mathbf{x})$. $p(\mathbf{x})$ can therefore be seen as a base generative model that can generate sample sentences [Dathathri et al., 2020].

In recent years, the attention-based models, Transformers [Vaswani et al., 2017], have replaced recurrent neural networks (RNNs) as the dominant architecture for LMs, with major improvements in distribution estimation, long-range dependency handling, sample diversity, and parallel processing. Another recent development in language modeling is that of pre-training LMs on massive corpora. So-called large-scale general purpose LMs have demonstrated significant improvements in downstream tasks, i.e., other NLP tasks for which the model was not specifically trained or fine-tuned. Most famously the OpenAI's series of Generative Pre-trained Transformer (GPT) models have improved numerous NLP benchmarks [Radford et al., 2018, 2019, ?].

2.1.2 (Controllable) Text Generation

In text generation, a language model $p(\mathbf{x})$ is asked to produce text \mathbf{x} given a prompt by sampling from the distribution of words that are assigned the highest likelihood of following the primer text [Radford et al., 2019]. Text generation in itself is the task of generating a piece of text given an input text. This process can be seen as sampling from a conditional distribution. Controllable text generation refers to the more restrictive problem of enforcing higher-level linguistic features on the generated text during sampling [Dathathri et al., 2020, Prabhumoye et al., 2020]. This can be seen as a sub-problem of vanilla text generation, because the conditioning factor for the output text is further constrained to also include some predefined textual attribute. This attribute represents a linguistic characteristic of the text, like sentiment, topic, or writing style.

Controllable text generation or CTG is a more challenging problem than vanilla text generation for a number of reasons. First, defining the desired attribute to be controlled for in a manner that it is intelligible for a machine is a challenge in itself [Zheng et al., 2019]. Second, like many NLP problems, there are not many parallel corpora [Dai et al., 2019]. In the context of controllable generation, parallel corpora are datasets of target and source texts that only differ with respect to some attribute. Furthermore, the measure of attribute adherence is a very vague and ambiguous concept [Dathathri et al., 2020, Dai et al., 2019]. Namely, a text can be written to have an extremely positive sentiment in multiple formulations, all of which adhere to the positive sentiment. Another important hurdle for controllable text generation, especially when CTG is combined to leverage the linguistic power of large-scale language models, is that the cost of fine-tuning or pre-training a model to control for a linguistic attribute can be very high [Dathathri et al., 2020, Madotto et al., 2020].

2.1.3 Dialogue, Dialogue Systems, and Dialogue Generation

[L:

- Think of better title.
- Flesh out this section more.
- Add more references. E.g., Welleck et al. [2019] have good definitions for Dialogue Generation and Persona-Based Dialogue.
- Properly introduce the concept of a dialogue with references.

- Properly introduce Dialogue Systems and their general scheme. E.g., Kushneryk et al. [2019] provide good figures for these concepts.
- Emphasize that the focus of this thesis is on Open-Domain dialogue systems. Talk about open-domain dialogue systems
- *Open-domain conversational systems are a special case of language models where the prefix is the dialogue history and the continuation is a humanlike response.* [Madotto et al., 2020]
- Very briefly mention that the focus of your research is on the natural language generation part of a dialogue system.

]

A dialogue a written or spoken conversational exchange between two or more interlocutors [L: Are you sure it requires two or more? And not just one or more?]. Generally speaking, its purpose is to exchange information or build relationships between interlocutors [Bohm and Nichol, 2013]. A dialogue typically consists of interlocutors exchanging utterances in turns.

[L: **TODO** - Give an example of a dialogue snippet (from the BNC)]

Dialogue distinguishes itself from discourse in that it necessarily involves two or more participants exchanging information and contributing to the conversation, whereas discourse can be a one-way exchange of information, like a lecture or blog-post. [L: Verify if this is true. Is that really the difference? Why does dialogue require at least two participants?]

Developing dialogue systems is a very complex task and many approaches have been proposed for this [McTear, 2020]. [L: Find specific examples of these approaches, mention a few along with references.]

The focus of my thesis is on dialogue generation, a particular operationalization of dialogue in NLP.

It is the task of automatically generate a response given a prompt by the user [Madotto et al., 2020].

Dialogue generation can be framed as a next utterance prediction problem [Welleck et al., 2019]. In this framework, an utterance is a sequence of tokens representing a dialogue turn. The next utterance is predicted conditioned on a dialogue prefix or prompt (e.g., a single previous

dialogue turn, or the entire conversation history). Continuing with the framework of Welleck et al. [2019], a sequence of utterances can be interpreted as a dialogue between agents.

Computational dialogue modeling distinguishes itself from most NLP domains due to the challenges associated with modeling human conversation: informal, noisy, unstructured, and even erroneous real-world responses, possibly competing goals of interlocutors, or an inherently more diverse set of acceptable responses.

[L: TODO - work out best way to merge previous and next paragraph(s)]

Text generation is suitable to tackle tasks such as machine translation, abstractive summarization, and paraphrasing. Dialogue response generation is also a special case of language generation. It can be seen as language generation where the prompt is a turn in a dialogue session. Conversational response generation shares open-domain text generation’s overarching objective of producing grammatically correct fluent text, while remaining relevant to the prompt.

2.1.4 Dialogue Response Generation as a Language Modeling Task

[L: Is this the best place for this subsection?]

Dialogue response generation can be formulated as a language modeling task. Zeng et al. [2020] suggest to formulate it as a language modeling task, where the next token or utterance is conditioned on the dialogue history. More formally, it consists of concatenating all dialogue turns in a multi-turn dialogue session into a long text: x_1, \dots, x_N . Denote the source sentence or dialogue history as $S = x_1, \dots, x_m$ and the target sentence (ground truth response) as $T = x_{m+1}, \dots, x_N$. The conditional probability of dialogue continuation given its history $P(T|S)$ can be written as

$$p(T|S) = \prod_{n=m+1}^N p(x_n|x_1, \dots, x_{n-1}). \quad (2.2)$$

A multi-turn dialogue session T_1, \dots, T_K can be written as $p(T_K, \dots, T_2|T_1)$ which is essentially the product of all source-target pairs probabilities $p(T_i|T_1, \dots, T_{i-1})$. This formulation also shows that optimising the single objective $p(T_K, \dots, T_2|T_1)$ is equivalent to optimising all source-target pair probabilities.

2.1.5 Controllable Dialogue Generation

Endowing a dialogue system with personality traits to generate human-like conversation is a long-standing goal in AI [Edlund et al., 2008, Scheutz et al., 2011, Madotto et al., 2020]. Zheng et al. [2019] argue that this objective is difficult to reach because of the challenge of representing personality traits via language expression and the lack of large-scale persona-labeled dialogue datasets. Assuming an encoder-decoder setup, the same authors postulate that most personalized neural conversation models can be classified as one of two types: implicit and explicit personalisation models. For implicit personalization models, each speaker has its own vector representation, which implicitly captures the speaking style of the speaker in the decoding process [Kottur et al., 2017, Li et al., 2016]. These models enjoy the benefit of having a more granular and realistic representation of speaking style, as opposed to a simple discrete set of traits (as is the case for explicit personalization models). On the other hand, it is unclear how speaker style is captured and should be interpreted, as all the information about a speaker’s style is encoded in a real-valued vector. Furthermore, these methods suffer from a data sparsity issue, because each dialogue should be tagged with a speaker identifier and there should be sufficient dialogues from each trait-group to train a reliable trait-adaptive model.

When generating responses, *explicit* personalization models are conditioned either on a given personal profile [Qian et al., 2018], text-described persona [Zhang et al., 2018a], or simply an attribute label [Madotto et al., 2020]. That is, speaker traits are represented as key-value pairs or descriptions about age, gender, etc. This can be seen as conditioning the decoder’s output on an attribute a . Speakers with same set of personality traits can share attribute representations, so it does not require a speaker-specific representation vector. Such structured character descriptions are more explicit, straight-forward, and interpretable. However, explicit personalization models require manually labeled or crowdsourced datasets for development, making it difficult to scale these models to large-scale dialogue datasets Zheng et al. [2019], Madotto et al. [2020].

2.1.6 Transformers

The Transformer architecture plays a central role in most of the recent advances in NLP. The same holds for the methods used in this thesis to investigate controlled dialogue generation and automated detection of age-related linguistic patterns in dialogue utterances. For a more detailed

review of the model architecture, the reader is referred to the original paper [Vaswani et al., 2017] or this blog post¹.

The Transformer, like most neural sequence processing models, has an encoder-decoder structure. On a high level, the encoder receives an input sequence $\mathbf{x} = (x_1, \dots, x_n)$ (e.g., a sentence), and maps this to a sequence of latent continuous variables $\mathbf{z} = (z_1, \dots, z_n)$. The decoder then takes \mathbf{z} as input, and maps this to an output sequence $\mathbf{y} = (y_1, \dots, y_m)$. Note that the use of positional encodings of the input and output embeddings enables the Transformer to process and generate sequences in arbitrary order, allowing for a high degree of parallelization. The generation of \mathbf{y} happens element-by-element in an auto-regressive fashion, where at step t , element y_{t-1} is also taken as input.

Both the encoder and decoder are comprised of N identical layers (denoted by the ‘ $N \times$ ’ in the left part of Figure 2.1). Every sub-layer performs a succession of transformations using multi-head self-attention mechanisms and point-wise, fully connected layers, along with residual connections [He et al., 2016] around every sub-layer followed by layer normalization [Ba et al., 2016]. The decoder’s first self-attention sub-layer is masked to ensure that the output predictions at sequence position i cannot depend on output positions greater than i . Finally, the decoder passes its output through a linear and softmax layer to produce a probability distribution over the problem space (e.g., the vocabulary) from which the most likely symbols for the generated output sequence \mathbf{y} can be sampled.

A key aspect of the Transformer architecture is its use of attention [Bahdanau et al., 2015]. This allows the encoder-decoder architecture to selectively focus on parts of the input sequence to produce a more informative hidden representation. Vaswani et al. [2017] formulate an attention function as a mapping of queries and sets of key-value pairs to an attention output, where matrices represent the queries Q , keys K , and values V . The attention output is a weighted sum of the values, based on the relevance of the corresponding keys to a query. In particular, they employ scaled dot-product attention:

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V. \quad (2.3)$$

Furthermore, Vaswani et al. [2017] propose to use multi-head attention by using learned linear projections to project the queries, keys and values h times, and apply the aforementioned

¹<https://jalammar.github.io/illustrated-transformer/>

attention function to these projections in parallel. The concatenation of these attention outputs, passed through a linear layer, ultimately produces the final output of the Transformer’s attention sub-layers. This allows the model to attend to the relevant information from all representation sub-spaces at various sequence positions. See Figure 2.1 for a schematic illustration of the Transformer’s structure described above.

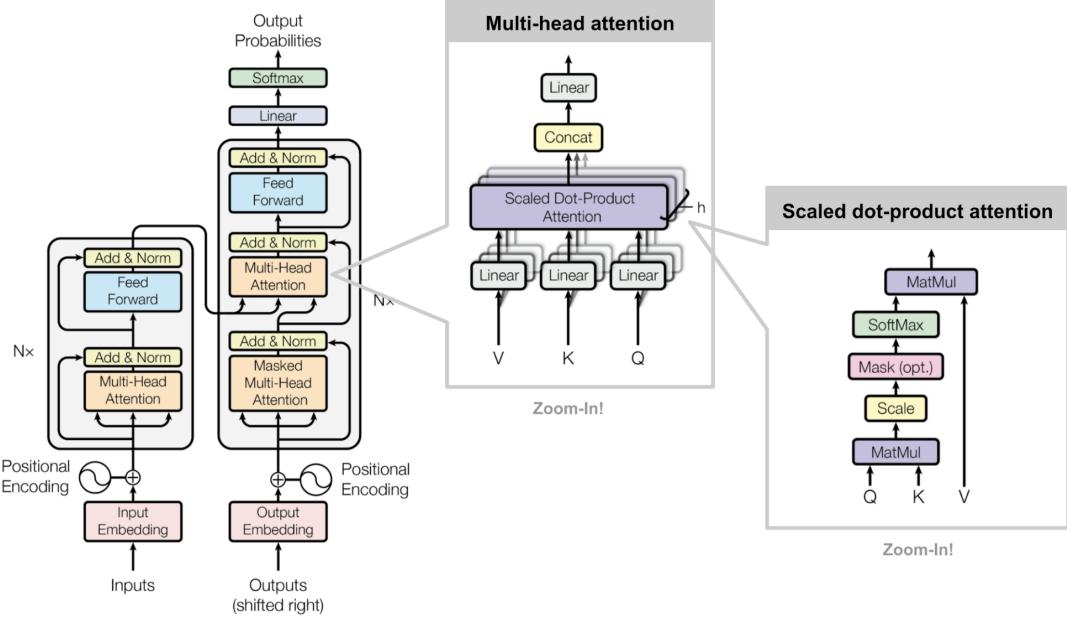


Figure 2.1: An overview of the full Transformer model architecture. *Collated image source:* Fig. 17 in this blog post <https://lilianweng.github.io/lil-log/2018/06/24/attention-attention.html>. *Original image source:* Figures 1 and 2 in Vaswani et al. [2017]

2.1.7 Causal Language Modeling with Transformers for Dialogue Response Generation

As described in Section 2.1.3, here we focus on dialogue response generation. A formulation of dialogue response generation as a language modeling task is described in Section 2.1.4. Following the conventions of Dathathri et al. [2020] and Madotto et al. [2020], a dialogue is comprised of multiple alternating turns (sometimes referred to as utterances) between more than one speaker. For simplicity, this project only focuses on dialogues between two speakers. The conversation history at turn t is defined as $\mathcal{D}_t = \{S_1^{(1)}, S_1^{(2)}, \dots, S_t^{(1)}\}$, where $S_t^{(j)}$ is speaker j ’s utterance at time t . Madotto et al. [2020] denote speaker 1 as the user U , and speaker 2 as the conversational system S , yielding dialogue history $\mathcal{D}_t = \{U_1, S_1, \dots, U_t\}$. This notational convention will also be used for the user-system experiments later on in this thesis.

A Transformer-based language model (denoted LM) is used in this thesis to model the distribution of dialogues, using dialogue history at time t , \mathcal{D}_t , as a prompt to auto-regressively generate the dialogue continuation S_t . More specifically, let the concatenation of the dialogue history at t and its continuation, $\{\mathcal{D}_t, S_t\}$, be represented as a sequence of tokens $\mathbf{x} = \{x_0, \dots, x_n\}$. Then, by recursively applying the product rule of probability (Bishop [2006]), the unconditional probability of the sequence $p(\mathbf{x})$ can be expressed as:

$$p(\mathbf{x}) = \prod_{i=1}^n p(x_i | x_0, \dots, x_{i-1}). \quad (2.4)$$

Dathathri et al. [2020] and Madotto et al. [2020] define the Transformer's decoding process in a recursive fashion. Let H_t denote the conversation history's key-value pairs, i.e., $H_t = [(K_t^{(1)}, V_t^{(1)}), \dots, (K_t^{(l)}, V_t^{(l)})]$, with $(K_t^{(i)}, V_t^{(i)})$ representing the key-value pairs from the LM's i -th layer generated at all time steps 0 through t . This results in the recurrent decoding process being expressed as:

$$o_{t+1}, H_{t+1} = \text{LM}(x_t, H_t), \quad (2.5)$$

where o_{t+1} is the hidden state of the last layer. Finally, after applying a softmax transformation, the next token x_{t+1} is sampled from the resulting probability distribution, i.e., $x_{t+1} \sim p_{t+1} = \text{softmax}(W o_{t+1})$, where W is a linear mapping from the model's last hidden state to a vector of vocabulary size. This recursive formulation allows for efficient text generation by leveraging cached memories, without repeated forward passes.

2.1.8 Plug-and-Play Language Modeling

The Plug-and-Play Language Model (PPLM) Dathathri et al. [2020] works by using a text classifier, referred to as an attribute model, to control the text generated by a language model. Let $p(X)$ denote the output distribution of a Transformer-based language model (e.g., GPT-2 or DialoGPT), where X represents the generated text. And $p(a|X)$ denotes the attribute model (e.g., a single-layer or BoW classifier) that represents the degree of adherence of text X to a certain attribute a (e.g., topic or sentiment (previous work), or age-group characteristics (this work)). Then PPLM can be seen as modeling the conditional distribution of generated text X given attribute a , i.e., $p(X|a)$. Note that Bayes' theorem ties these three definitions together as follows:

$$p(X|a) \stackrel{\text{Bayes' theorem}}{\hat{\equiv}} \frac{p(X)p(a|X)}{p(a)} \propto p(X)p(a|X). \quad (2.6)$$

To control the generated text, PPLM shifts the aforementioned history H_t (i.e., all Transformer key-value pairs generated up to time t) in the direction of the sum of two gradients:

1. Ascending $\nabla \log p(a|X)$: maximizing the log-likelihood of the desired attribute a under the conditional attribute model. This enforces attribute control.
2. Ascending $\nabla \log p(X)$: maximizing the log-likelihood of the generated language under the original (possibly conversational) language model. This promotes fluency of the generated text.

These two incentive-representing gradients are combined with various coefficients, yielding a set of tunable parameters to steer the generated text in the direction of the desired fluency, attribute control, and length.

Let's first focus on the first of the two gradients, i.e., the attribute control promoting $\nabla \log p(a|X)$. ΔH_t represents the update to history H_t that pushes the distribution of the generated text X in the direction that has a higher likelihood of adhering to desired attribute a . The gradient update rule can be expressed as:

$$\Delta H_t \leftarrow \Delta H_t + \alpha \frac{\nabla_{\Delta H_t} \log p(a|H_t + \Delta H_t)}{\|\nabla_{\Delta H_t} \log p(a|H_t + \Delta H_t)\|^\gamma} \quad (2.7)$$

where α is the step size, and γ denotes the normalization term's scaling coefficient. Both step size (α) and the scaling coefficient (γ) influence attribute control. Attribute control can be softened by either decreasing α or increasing γ and vice versa. Note that $\alpha = 0$ recovers the original uncontrolled underlying language model (e.g., GPT-2 or DialoGPT). In practice, ΔH_t is initialized at zero, and the update rule in Equation 2.7 is applied m times (usually 3 to 10), resulting in the updated key-value pair history $\tilde{H}_t = H_t + \Delta H_t$. Then the updated history \tilde{H}_t is passed through the language model, yielding the updated logits (final Transformer-layer): $\tilde{o}_{t+1}, H_t = \text{LM}(x_t, \tilde{H}_t)$. And finally the shifted \tilde{o}_{t+1} is linearly mapped through a softmax layer to produce a new, more attribute-adherent, distribution from which to sample, i.e., $x_{t+1} \sim \tilde{p}_{t+1} = \text{softmax}(W\tilde{o}_{t+1})$.

The method described until now will generate attribute-adherent text, but will likely yield fooling examples [Nguyen et al., 2015] that are gibberish to humans, but get assigned high $p(a|x)$ by

the attribute model [Dathathri et al., 2020]. That is why Dathathri et al. [2020] apply two methods to ensure fluency of the generated text. The first is to update ΔH_t such to minimize the Kullback-Leibler (KL) divergence (denoted D_{KL}) between the shifted and original distributions. In practice, D_{KL} is scaled by a coefficient λ_{KL} , typically found to work well for most tasks when set to 0.01. Repetitive text generation (i.e., high $p(a|x)$ but low $p(x)$) can therefore sometimes be avoided by increasing λ_{KL} . The second method to ensure fluency is Post-norm Geometric Mean Fusion [Stahlberg et al., 2018] which, instead of directly influencing ΔH_t like minimizing D_{KL} , fuses the altered generative distribution \tilde{p}_{t+1} with the unconditional language distribution $p(x)$. This is done during generation by sampling the next token as follows:

$$x_{t+1} \sim \frac{1}{\beta} \left(\tilde{p}_{t+1}^{\gamma_{gm}} p_{t+1}^{1-\gamma_{gm}} \right) \quad (2.8)$$

where β is a normalization constant, p_{t+1} and \tilde{p}_{t+1} denote the original and modified distributions, respectively, and γ_{gn} is a scaling term that interpolates between the two distributions. Because the new sampling distribution in Equation 2.8 converges towards the unconditional language model as $\gamma_{gm} \rightarrow 0$, repetitive text generation can be avoided by decreasing the scaling term.

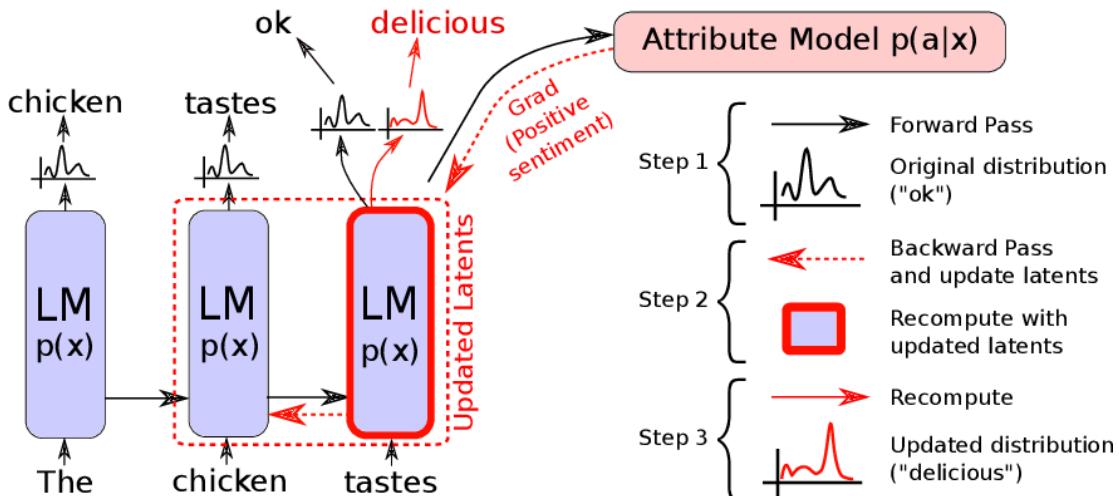


Figure 2.2: A schematic overview of the plug-and-play interaction between attribute model $p(a|x)$ and language model $p(x)$. *Original image source:* Figure 1 of Dathathri et al. [2020]

It is worth mentioning that the plug-and-play method applied by Dathathri et al. [2020] and Madotto et al. [2020] is different from fine-tuning. Note that in Equation 2.7 the gradient updates are restricted to the history H_t , and do not affect the model's parameters. Because the key-value pairs $(K_t^{(i)}, V_t^{(i)})$ that comprise H_t are activations and not model-weights, the updates only take

place in the activation-space. This means that PPLM leaves the underlying (conversational) language model untouched.

Contrary to fine-tuning often massive LMs, PPLM does not incur a significant training cost (depending of course on the complexity of the discriminator or attribute model). However, Madotto et al. [2020] show that PPLM needs a fixed number of m update-steps to for every generated token. This makes the original PPLM setup unsuitable for online interactive applications, like conversational systems. Addressing this problem, they introduce plug-and-play conversational model (PPCM), which extends PPLM by using the original model setup to generate dialogue datasets with the desired attribute a , and then use optimized residual adapters [Bapna and Firat, 2019] to control LM’s output distribution. More specifically, Madotto et al. [2020] use PPLM to generate n attribute-adherent dialogue datasets $\mathcal{D}^a = \{\mathcal{D}^1, \dots, \mathcal{D}^n\}$, for attribute a . These generated dialogue datasets are then used to train the residual adapter modules that control the language model’s output distribution.

2.1.9 Language and Age

The relationship between a person’s age and use of language is a thoroughly studied subject with a decades long history and inherent challenges [Pennebaker and Stone, 2003, Nguyen et al., 2014, Zheng et al., 2019]. A number factors like community membership (e.g., gender, socioeconomic status, or political affiliation), experimental condition (e.g., emotional versus non-emotional disclosure), mode of disclosure (writing versus talking), and other confounding variables complicate the study of age’s relation to language [Nguyen et al., 2011]. The relatively recent advent of widely available computational resources and vast amounts of textual data made it possible to leverage machine learning methods to help detect patterns in language that eluded conventional sociolinguistic research. Early computational investigations into the connection between a person’s age and use of language is typically a combination of qualitative and statistical methods. For instance, using a mix between their proprietary count-based text analysis framework, Linguistic Inquiry and Word Count (LIWC) and sociolinguistic theory, Pennebaker and Stone [2003] study the changes in written and spoken language use with increasing age. They discuss four important areas of a person’s character that have been found to change with age: emotional experience and expression, identity and social relationships, time orientation, and cognitive abilities. These four axes and their hypothesized relationships with language use and age can be interpreted in the following ways:

1. *Emotional experience and expression*: This is the relationship between increasing age and linguistically observable manifestations of a person's experienced emotions. In practical terms, this is framed as detectable instances of positive and negative affect in language. This complex relationship between age and emotional expression is characterized by decreased levels of negative affect and slightly non-decreasing levels of positive affect. This is also confirmed by the findings of Schler et al. [2006].
2. *Sense of identity and social relationships*: These terms refer to developmental trends in one's relation to self and others, as expressed in their language, e.g., as references to self (*I, me, my, and we, us, our*) or others (*they, them, theirs*). Pennebaker and Stone [2003] report that the *quantity* of social connections decreases and the *quality* of remaining relationships increases with age.
3. *Time orientation*: This relationship describes how people express their perception of and orientation towards time. For instance, this can be indicated by the use of time-related verb tenses. The authors suggest that older individuals tend to be more past-oriented than their younger future-oriented counterparts.
4. *Cognitive abilities*: This refers to markers of cognitive capacity in language. Aging is expected to be associated with less use of cognitively complex words after a certain mid-adulthood peak. Specifically, the relationship between markers of cognitive complexity in natural language (cognitive mechanisms, causal insight, and exclusive words) and age is hypothesized to be curvilinear. And because verbal ability does not decline until very late in life, markers of verbal ability (e.g., use of big words) are not expected to show changes with age.

Pennebaker and Stone [2003] consider the following variables: positive and negative emotions, first-person singular and first-person plural pronouns, social references, time-related words (past-tense, present-tense, and future-tense verbs), big words (> 6 letters), cognitive mechanisms, causal insight, and exclusive words. Their main findings suggest that increasing with age, people use more positive and fewer negative affect words, use fewer self-references, use more future-tense and fewer past-tense verbs, and exhibit a general pattern of increasing cognitive complexity.

Detectable linguistic differences between age-groups can often be attributed to the use of language fads or references to age-specific popular culture. For instance, Schler et al. [2006] find that the use of slang and neologisms (such as *lol* and *ur*) are strong indicators of youth. Similarly, words

like ‘facebook’, ‘instagram’, and ‘netflix’ appear in the most frequently used words by younger participants of conversational data collection efforts, like that of the British National Corpus’ spoken component [Love et al., 2017].

2.2 Related Work

[L: Keep in mind the following distinction between Background and Related Work - The Background section should give an overview of the problem and the components involved: dialogue, language generation, dialogue response generation, age modeling, etc., without focusing on one or the other approach — in Related Work, you describe approaches that have been proposed to tackle each of these components, separately or jointly, and which are related or relevant to your own work for some reason]

2.2.1 Automated Age Detection

[L: Consider adding a small sub-section about automated age detection from text, because you often bring up the problem and other researchers’ approaches to solving it in the Background section.]

...

[L: work this into this subsection. Taken from workshop paper submission]

Previous work on age detection in dialogue has focused on speech features, which are known to systematically vary across age groups. For example, Wolters et al. [2009] learn logistic regression age classifiers from a small dialogue dataset using different acoustic cues supplemented with a small set of hand-crafted lexical features, while Li et al. [2013] develop SVM classifiers using acoustic and prosodic features extracted from scripted utterances spoken by participants interacting with an artificial system. In contrast to this line of work, we investigate whether different age groups can be detected from textual linguistic information rather than voice-related cues. We explore whether, and to what extent, various state-of-the-art NLP models are able to capture such differences in dialogue data as a preliminary step to age-group adaptation by conversational agents. We build on the work of Schler et al. [2006], who focus on age detection in written discourse using a corpus of blog posts. The authors learn a Multi-Class Real Winnow classifier leveraging a set of pre-determined style- and content-based features, including part-of-speech categories, function words, and the 1000 unigrams with the highest information gain in the training set. They find that content features (lexical unigrams) yield higher accuracy

($\sim 74\%$) than style features ($\sim 72\%$), while their best results are obtained with their combination ($\sim 76\%$). We extend this investigation in several key ways: (1) we leverage state-of-the-art NLP models that allow us to learn representations end-to-end, without the need to specify concrete features in advance; (2) we apply this approach to dialogue data, using a large-scale dataset of transcribed, spontaneous open-domain dialogues; (3) we show that text-based models can indeed detect age-related differences in both discourse and dialogue, even in the case of very sparse signal at the level of dialogue utterances; and finally (4) we carry out an in-depth analysis of the models’ predictions to gain insight on which elements of language use are most informative.

...

More recent studies, like that of Nguyen et al. [2011], Zheng et al. [2019], and Abdallah et al. [2020], frame age prediction from text as traditional machine learning problems, like linear regression, support vector machines, or neural architectures. These modeling approaches tend to reveal that strong indicators of age lie at the syntactic or structural level of language use, as opposed to the more content-based lexical level. Furthermore, this could explain why automatic detection from text of more content-based traits, like topic or sentiment, tend to be easier problems to solve than age prediction from text. To emphasize one such complicating factor, Nguyen et al. [2014] argue that differences in language use are often relation to the speaker’s social identity, which could differ from their biological identity.

2.2.2 Controllable Language Generation

Previous approaches to controlled language generation require fine-tuning large Transformer-based language models or training conditional generative LMs from scratch. Most notably CTRL [Keskar et al., 2019], which achieves controllable generation by training a generative Transformer for a number of control codes. CTG models that require fine-tuning for control, like CTRL, can produce high quality fluent text because they are specifically trained to maximize the likelihood of generated sequences, given an attribute (denoted $p(\mathbf{x}|a)$), but require training massive language models with computational costs.

Other recent examples of controllable language generation models that are not Transformer-based also exist. Li et al. [2020] introduce OPTIMUS, a large pre-trained Variational Autoencoder (VAE) [Kingma and Welling, 2014] that can be fine-tuned for specific natural language tasks, like guided sentence generation. They demonstrate OPTIMUS’ ability to perform controlled text generation from latent style-embeddings, with fluency at par with GPT-2. They also

show how OPTIMUS generalizes better for low-resource languages than BERT [Devlin et al., 2019]. Nevertheless, much like the previously mentioned CTG models, OPTIMUS still incurs a significant computational cost for fine-tuning per NLP task.

[L: TODO: Where does the following sentence fit best? “The plug-and-play setup of PPLM forms one of the main theoretical foundations of this work.”]

The plug-and-play language model (PPLM) [Dathathri et al., 2020] is a recent solution to the problem of high re-training costs of controlled language generation. This approach, inspired by a similar technique for style-control of generated images [Nguyen et al., 2017], leverages the fluency of large-scale language models when controlling them for a specific linguistic attribute, while avoiding incurring significant costs of fine-tuning these massive language models. The main benefit of this setup is its low-cost extensibility. Namely, such large-scale language models are often open-source and available online, and can now be tailored to users’ specific needs using a significantly easier to train attribute model. The original architecture proposed by Dathathri et al. uses GPT-2 as a base language model which provides grammatical fluency, combined with a significantly easier to train attribute model (i.e., a simple BoW or single-layer classifier). Using gradient updates to the activation space of the much smaller attribute model, they manage to generate language that combines (some of) the fluency of GPT-2 with the stylistic control of the attribute model, without the cost of retraining a specialised architecture. They demonstrate that PPLM achieves desirable fluency (i.e., perplexity measured with GPT(-1) [Radford et al., 2018]), as well as measurable attribute control. Their architecture’s applicability is also demonstrated on tasks such as controlled story writing and language detoxification. They also show a clear trade-off between attribute control and grammatical correctness and diversity.

2.2.3 Text Style Transfer

Text style transfer is the task of changing a text’s stylistic properties, while retaining its style-independent properties, like content and fluency [Dai et al., 2019]. Text style transfer is a closely related problem to controllable language generation. Its similarity lies in trying to modify the output distribution of a language generation model, such that stylistic characteristics of the produced text are controllable, keeping content and fluency preserved. It involves rewriting an input text with a specific style. More formally, given a text \mathbf{x} , its corresponding style-representing vector $\mathbf{s}^{(i)}$, the number of different styles K over which there exists a distribution, and a desired

style $\hat{\mathbf{s}} \in \{\mathbf{s}^{(i)}\}_{i=1}^K$, the goal of text style transfer is to produce output text $\hat{\mathbf{x}}$ with style $\hat{\mathbf{s}}$, and the style-independent properties of \mathbf{x} .

Previous approaches to text style transfer involve passing input text through an RNN-based encoder, yielding a style-dependent latent representation \mathbf{z} [Zhang et al., 2018b]. Typically, these approaches then attempt to “disentangle” \mathbf{z} into a style-independent content representation and a latent representation of the stylistic properties of the input text. The subsequent decoder then receives the content representation and a new latent style variable as input, to ultimately produce a style-altered output text with unchanged content. This style-disentanglement approach has a number of drawbacks: **(1)** It is difficult to evaluate the quality of disentanglement of the latent space. **(2)** It is hard to capture rich semantic information in the latent representation due to limited capacity of vector representations (especially for long texts). **(3)** To disentangle style and content in the latent representations, all previous approaches have to assume all input texts can be encoded by a fixed-size latent vector. **(4)** Since most previous approaches use RNN-based encoder-decoder frameworks, they have problems capturing long-range dependencies in the input sentences. Furthermore, disentanglement might be unnecessary, as Lample et al. [2019] have shown a proper decoder can perform controllable text generation from an entangled latent representation by “overwriting” the original style.

To address these drawbacks, Dai et al. [2019] propose Style Transformer, a Transformer-based alternative encoder-decoder framework for text style transfer. The authors’ approach does not require any manipulation (i.e., disentanglement) of the latent space, eliminates the need for a fixed-size vector representation of the input, and handles long-range dependencies better due to Transformers’ attention mechanism. Aside from this being the first application of Transformers for text style transfer, Dai et al. [2019] contribute a novel training algorithm for such models, that boasts significant improvements of results on two text style transfer datasets.

2.2.4 Dialogue Generation

Dialogue generation is task of automatically generating a response given a user’s prompt. Zhang et al. [2020] introduce DialoGPT, a tunable large-scale language model for generation of conversational responses, trained on Reddit discussion chain data. DialoGPT therefore extends GPT-2 [Radford et al., 2019] to address a more restrictive sub-category of text generation, i.e., conversational response generation. DialoGPT inherits from GPT-2 a 12-to-48 layer transformer with layer normalization, a custom initialization scheme that accounts for model depth, and byte

pair encodings [Sennrich et al., 2016] as a tokenizer. The generation task remains framed as language modeling, where a multi-turn dialogue session is modeled as a long text.

To address the well-known problem of open-domain text generation models producing bland and uninformative samples, Zhang et al. [2020] implement a maximum mutual information (MMI) scoring function. MMI uses a pre-trained backward model to predict $p(\text{source}|\text{target})$: i.e., the source sentences (dialogue history) given the target (responses, dialogue continuation). First, top-K sampling is used to generate a set of hypotheses. Then the probability $p(\text{source}|\text{hypothesis})$ is used to re-rank all hypotheses. As frequent and repetitive hypotheses can be associated with many possible queries/sources (i.e., a hypothesis that frequently occurs is one that is apparently applicable to many queries), and maximizing backward model likelihood penalizes repetitive hypotheses, MMI yields a lower probability for highly frequent hypotheses, thereby reducing blandness and promoting diversity.

DialoGPT is evaluated on the Dialog System Technology Challenge (DSTC) 7 track, an end-to-end conversational modeling task in which the goal is to generate conversation responses that go beyond chitchat by injecting information that is grounded in external knowledge. The model achieves state-of-the-art results on both the human and automatic evaluation results, by achieving near human-like responses that are diverse, relevant to the prompt, much like GPT-2 for open-domain language generation. They train 3 models of small (117M), medium (345M), and large (762M) parameter sizes. The medium-sized 345M model achieves the best automatic evaluation results across most metrics, and is used as one of the baselines in later experiments in this thesis. Their Hugging Face PyTorch implementation can be tested here: <https://huggingface.co/microsoft/DialoGPT-medium>.

Dialogue generation is the essential precursor to this thesis' ultimate task of controllable dialogue generation.

[L: The previous sentence feels out of the blue. Consider removing it or think of a way to create a natural flow towards it.]

2.2.5 Controlled Dialogue Generation

Controlled dialogue generation is the task of steering automatically generated conversational responses to possess desired attributes, like sentiment, topic, or more abstract writing style characteristics. Zeng et al. [2020] explore the applications of fine-tuning large language models, like GPT, on (Mandarin and English) medical consultation data. The resulting dialogue systems

succeed at generating clinically correct and human-like responses to patients’ medical questions. Medical dialogue systems like these can help make healthcare services more accessible and aid medical doctors to improve patient care.

Zheng et al. [2019] investigate the problem of incorporating explicit personal characteristics in dialogue generation to deliver personalized conversation. They introduce a dataset **PersonalDialog**, which is a large-scale multi-turn dialogue dataset with personality trait labeling (i.e., Age, Gender, Location, Interest Tags, etc.) for a large number of speakers. Zheng et al. [2019] also propose persona-aware models that include a trait fusion module in the encoder-decoder framework to capture and address personality traits in dialogue generation. Persona-aware attention mechanisms and bias are used to incorporate personality information in the decoding process. All their tested classification and dialogue generation models are either variations of RNNs (such as LSTMs or gated recurrent units (GRUs)), convolutional neural networks (CNNs), or hybrids of these systems (LSTM-outputs fed into a CNN, known as recurrent convolutional neural networks (RCNNs)). The authors study the influence of age, gender, and location on dialogue classification and generation, and use both automatic (perplexity, trait accuracy, and generated response diversity measures) and human evaluation. They find dialogues to be distinguishable by gender (about 90.61% test accuracy), then age (78.32% test accuracy), and finally location (62.04% test accuracy). Both automatic and human evaluation of the generated responses show that the best performing models benefit greatly from the persona-aware attention mechanism, possibly making a case to consider more attention-based architectures instead of RNNs.

Although the previously mentioned architectures are able to produce human-like conversational responses, sometimes even leveraging the fluency of large pre-trained LMs, they all suffer from the same computational drawback. They all require massive amounts of computational power to adapt their language styles, because in their cases, guided generation implies fine-tuning (or even retraining) large attribute-specific dialogue datasets. For general controlled language generation, this obstacle is overcome by Dathathri et al. [2020]’s previously mentioned PPLM setup. The conversational analog of this idea, plug-and-play conversational model (PPCM), is proposed by Madotto et al. [2020]. Similar to PPLM, PPCM achieves guided dialogue generation via activation-space perturbations using easy to train attribute models. Due to the computational complexity of PPLM’s decoding process, PPLM is unusable as practical conversational system. PPCM solves this problem by using residual adapters [Bapna and Firat, 2019] to tweak the

decoding procedure such that it does not require more computational resources. See Section 2.1.8 for a detailed explanation of the mechanisms behind PPLM and PPCM. Madotto et al. [2020] show, using both human and automatic evaluation, that PPCM can balance grammatical fluency and high degrees of attribute-adherence in its generated responses. PPCM uses DialoGPT as its base language model, and is tested for topical or sentimental attributes (i.e., positive, negative, sports, business, or science & tech). Previous work on controllable language generation focuses on content (e.g., topical attributes, or sentiment), rather than more abstract linguistic features, which I hypothesize are more challenging to model and control. The previously mentioned work by Zheng et al. [2019] is a notable exception, as their approach deals with controlling dialogue systems for linguistic features, like age, gender, and geographical region. However, Zheng et al. [2019] still suffers from significant computational costs, because control is achieved by fine-tuning a large system for every specific set of attributes. Furthermore, their proposed architectures are RNN-based, as opposed to my Transformer-based approach. My work therefore aims to extend the applicability of plug-and-play controlled generation to more abstract linguistic characteristics than those explored by Dathathri et al. [2020] and Madotto et al. [2020], and without the significant fine-tuning cost of Zheng et al. [2019].

[L: TODO

- Ask for Sandro’s feedback on this last rephrased paragraph.
- Is it worth mentioning that Zheng et al 2019 deals with Chinese Mandarin dialogue systems, and mine with English?

]

Chapter 3

Experiment 1: Detecting Age-Related Linguistic Patterns in Dialogue and Discourse

In this chapter we report experiments aimed at age detection from text, and the components involved. The problem we tackle in this first phase of experiments is automated detection of age-related linguistic patterns in dialogue and discourse, using current text-based NLP models. Being able to detect and investigate these linguistic differences is important for controlled dialogue generation, because it suggests that adapting automated conversational responses to a user’s age is possible. Moreover, it can provide us with insights about which linguistic features are most salient for distinguishing between, and adapting to, different age groups. We expect that the classification models are able to reliably detect age-related differences in both transcribed dialogue and discourse, and the most informative differences to lie at the syntactic-level.

The following section describes the two datasets used for these experiments. There we provide descriptive statistics, examples, and comparisons between the datasets. Section 3.2 covers the problem description in more detail, along with the models used, and our experimental setup. The classification results are presented in Section 3.3. Then for the dialogue classification models, Section 3.4 contains both quantitative and qualitative analyses of the results.

age 19-29	
A: oh that's cool	B: different sights and stuff
A: oh	
age 50+	
A: well quite and I'd have to come back as well	B: that's of course
A: and make up for you know	

Figure 3.1: Example dialogue snippets from speakers of different age groups (19-29 vs. 50+) in the British National Corpus. We conjecture that stylistic and lexical differences between age groups can be detected. In our approach, we experiment at the level of the utterance.

3.1 Data

We use a dataset of dialogue data where information about the age of the speakers involved in the conversation is available (see the dialogue snippets in Figure), i.e., the spoken partition of the British National Corpus [Love et al., 2017]. We henceforth refer to it as our *dialogue* dataset. For comparison with previous work, and to explore commonalities and differences between various types of language data, we also experiment with a dataset of discourse, i.e., the Blog Authorship Corpus used by Schler et al. [2006], that we henceforth refer to as our *discourse* dataset. Below, we briefly describe the two datasets along with the pre-processing steps we took to make the data suitable for our experiments.

dataset	# age groups	# samples	# tokens	mean length (\pm std)	min - max length	# topics
dialogue	2	67,282	787,352	11.7 (\pm 19.0)	1 - 1246	790
discourse	3	677,244	140M	102.2 (\pm 212.9)	1 - 71,580	40

Table 3.1: Descriptive statistics of the datasets used in our experiments. Length is the number of tokens in a sample.

3.1.1 Dialogue Dataset

This partition of the British National Corpus includes spoken informal open-domain conversations between people that were collected between 2012 and 2016 via crowd-sourcing, and then recorded and transcribed by the creators. Dialogues can be between two or more interlocutors, and are annotated along several dimensions including age and gender together with geographic and social indicators. Speaker ages in the original dataset are categorized in the following ten brackets: 0-10, 11-18, 19-29, 30-39, 40-49, 50-59, 60-69, 70-79, 80-89, and 90-99.

We focus on conversations in the British National Corpus that took place between two interlocutors, and only consider dialogues between people of the same age group. We then focus

on dialogues by speakers belonging to two age groups: 19-29 and 50+, in which we group conversations from five original brackets: 50-59, 60-69, 70-79, 80-89, and 90-99. We omit the intermediate age bracket to allow for clearer differentiation.

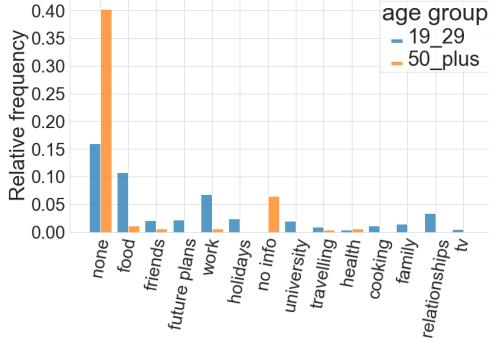
We split the dialogues into their constituent utterances (e.g., from each dialogue snippet in Figure 3.1 we extract three utterances), and further pre-process them by removing non-alphabetical characters. Only samples which were not empty after pre-processing were kept. The resulting dialogue dataset, that we use for our experiments, includes around 67K utterances with an average length of 11.7 tokens. Descriptive statistics of it are reported in Table 3.1.

Each conversation in the British National Corpus is annotated with a list of *topics* provided by the speakers during data collection. To extract a single representative topic from this list, we first compute the frequency of all topic labels in the whole dataset. Then, for each utterance, we take the label in the conversation with the highest frequency in the ranking. In total, our final dataset includes 790 unique topic labels. The distribution of the most frequent ones is reported in Figure 3.2a. As can be seen, frequent topics (besides the frequent *none* label) are *food*, *work*, and *holidays*, which reveals the colloquial and everyday nature of the dialogues in this dataset.

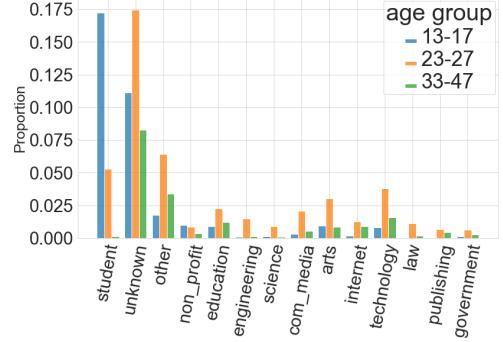
3.1.2 Discourse Dataset

The Blog Authorship Corpus Schler et al. [2006] is a collection of blog posts posted on <https://www.blogger.com>, gathered in or before August 2004. Each blog entry is written by a single user whose age, gender, and astrological sign are reported. The corpus contains almost 700,000 posts by 19,000 unique bloggers (i.e., ~35 posts per blogger on average). For our experiments, similar to Schler et al. [2006], we consider three age groups: 13-17, 23-27, and 33+. We pre-process the data in the same way as described above, namely by removing stopwords and non-alphabetical characters. The resulting dataset, that we use for our experiments, includes slightly more than 677K samples with an average length of 102.2 tokens. Descriptive statistics of it are reported in Table 3.1. [L: TODO - Adjust paragraph to W.S. setting.]

Each sample in the Blog Authorship Corpus is annotated with one topic. In our final discourse dataset, the unique topics present are 40. Figure 3.2b reports the distribution of the most frequent ones. As can be noted, frequent topics are *student*, *arts*, and *technology*, which reveals that this and the dialogue dataset are rather different.



(a) Distribution of most frequent topics (including the *none* and *no info* labels) in the **dialogue dataset**, shown by age group. Best viewed in color.



(b) Distribution of most frequent topics (including the *unknown* label) in the **discourse dataset**, shown by age group. Best viewed in color.

Figure 3.2

3.2 Methodology and Experimental Setup

The current section contains the methodology and experimental details of our automated age-detection experiments. We frame the problem as a N -class classification problem: given a fragment of text X , we seek to predict the age class of its speaker/writer. For the dialogue dataset, $N = 2$, while $N = 3$ for the discourse dataset. We experiment with various models, that we briefly describe here below. Details on the training and evaluation of models are given at the end of the sub-section.

n -gram Our simplest models are based on n -grams, which have the advantage of being highly interpretable. Each data entry (i.e., a dialogue utterance or blog post) is split into chunks of all possible contiguous sequences of n tokens. The resulting vectorized features are used by a logistic regression model to estimate the odds of a text sample belonging to a certain age group. We experiment with unigram, bigram and trigram models. Note that a bigram model uses unigrams and bigrams, and a trigram model unigrams, bigrams, and trigrams.

LSTM and BiLSTM We use a standard Long Short-Term Memory network [LSTM; Hochreiter and Schmidhuber, 1997] with two layers, embedding size 512, and hidden layer size 1024. Batch-wise padding is applied to variable length sequences. The original model’s bidirectional extension, the bidirectional LSTM [BiLSTM; Schuster and Paliwal, 1997], is also used. BiLSTM more thoroughly leverages forward and backward directed information by combining the hidden states from both directions. Padding is similarly applied to this model, and the following optimal

architecture is found: embedding size 64, 2 layers, and hidden layer size 512. Both RNN-model are found to perform optimally for a learning rate of 10^{-3} .

BERT We experiment with a Transformer-based model, i.e., Bidirectional Encoder Representations from Transformers [BERT; Devlin et al., 2019] for text classification. BERT is pre-trained to learn deeply bidirectional language representations from massive amounts of unlabeled textual data. We experiment with the base, uncased version of BERT, in two settings: by using its pre-trained frozen embeddings (BERT_{frozen}) and by fine-tuning the embeddings on our age classification task (BERT_{FT}). The BERT embeddings are followed by a dropout layer with dropout probability 0.1, and a linear layer with input size 768.

Experimental Details Both datasets are randomly split into a training (75%), validation (15%), and test (10%) set. Each model with a given configuration of hyperparameters is run 5 times with different random initializations. All models are trained on an NVIDIA TitanRTX GPU.

The n -gram models are trained in a One-vs-Rest (OvR) fashion, and optimized using the Limited-memory Broyden–Fletcher–Goldfarb–Shanno (L-BFGS) algorithm Liu and Nocedal [1989], with a maximum of 10^6 iterations. The n -gram models are trained until convergence or for the maximum number of iterations.

LSTMs and BERT-based models are optimized using Adam [Kingma and Ba, 2015], and trained for 10 epochs, with an early stopping patience of 3 epochs. The RNN-based models’ embeddings are jointly trained, and optimal hyperparameters (i.e., learning rate, embedding size, hidden layer size, and number of layers) are determined using the validation set and a guided grid-search. BERT_{FT} is fine-tuned on the validation set for 10 epochs, or until the early stopping criterion is met. BERT models have a maximum input length of 512 tokens. Sequences exceeding this length are truncated.

3.3 Detecting Age-Related Linguistic Patterns in Dialogue

We first report results on *discourse* to check whether we replicate previous findings. Then, we focus on *dialogue* to answer our research questions. We report accuracy and F_1 for each age group.

3.3.1 Classification Performance on Discourse

Table A.3 reports the results. As can be seen, all models are well above the baseline in terms of both accuracy and F_1 s. This overall confirms previous evidence Schler et al. [2006] that language features of (written) *discourse* can predict, to some extent, the age group to which the person belongs. At the same time, BERT fine-tuned on the age classification task stands out as our best-performing model by achieving highest accuracy (0.731) and highest F_1 in all age groups. BiLSTM and LSTM rank second (0.720) and third (0.714) in terms of accuracy, respectively, while a somehow more mixed pattern is observed for F_1 scores.

Overall, these results indicate that powerful neural models that are capable of representing the linguistic context have a great advantage on this dataset over simpler n -gram models, which are more than 10 accuracy points behind.

Finally, it should be noted that our best results are slightly lower than those obtained by Schler et al. [2006]. This could be due to two main reasons: First, they experiment with a different (smaller) dataset than ours,¹ which also has a different majority baseline (see Table A.3). Second, while in our approach all models are trained end-to-end on the task, Schler et al. [2006] use hand-crafted features that are specific to the dataset, (as mentioned in the Introduction), which could constitute an advantage.

3.3.2 Classification Performance on Dialogue

Table 3.3 reports the results. As can be seen, BERT fine-tuned on the task is again the best-performing model in terms of accuracy (0.729), which confirms the effectiveness of this model in detecting age-related linguistic differences. At the same time, it can be noted that the model based on trigrams is basically on par with it in terms of accuracy (0.722) and well above both LSTM and BiLSTM (0.693 and 0.691, respectively). A similar pattern is shown for F_1 scores, where BERT fine-tuned and the trigram model achieve comparable performance, with LSTMs being overall behind.

Overall, our results indicate that predicting the age group to which a speaker belongs, using text-based models, is possible also for *dialogue* data, though the task appears to be somehow more challenging compared to when performed on discourse (note that the improvement with respect to the majority/random baseline is lower in dialogue). At the same time, the different ranking of models observed between discourse and dialogue suggests possibly different strategies

¹They are left with roughly 511K datapoints after pre-processing, while we experiment with around 677K.

used by models to solve the task. In particular, the very good performance of the trigram model in *dialogue* suggests that leveraging ‘local’ linguistic features captured by n -grams is extremely effective in this setup. This could indicate that differences among various age groups are at the level of local lexical constructions. This deserves further analysis, which we carry out in the next section.

Model	Accuracy ↑ better	$F_1^{(13-17)}$ ↑ better	$F_1^{(23-27)}$ ↑ better	$F_1^{(33+)}$ ↑ better
Majority class	0.472	*	0.642	*
Schler et al. [2006]	0.762	0.860	0.748	0.504
unigram	0.603 (0.001)	0.760 (0.003)	0.706 (0.001)	0.491 (0.003)
bigram	0.627 (0.001)	0.788 (0.001)	0.715 (0.001)	0.504 (0.002)
trigram	0.625 (0.002)	0.789 (0.001)	0.716 (0.002)	0.485 (0.003)
LSTM	0.714 (0.005)	0.772 (0.007)	0.740 (0.004)	0.501 (0.006)
BiLSTM	0.720 (0.001)	0.778 (0.005)	0.746 (0.001)	0.486 (0.016)
$BERT_{frozen}$	0.604 (0.001)	0.627 (0.011)	0.666 (0.005)	0.198 (0.018)
$BERT_{FT}$	0.731 (0.002)	0.791 (0.003)	0.752 (0.005)	0.521 (0.020)

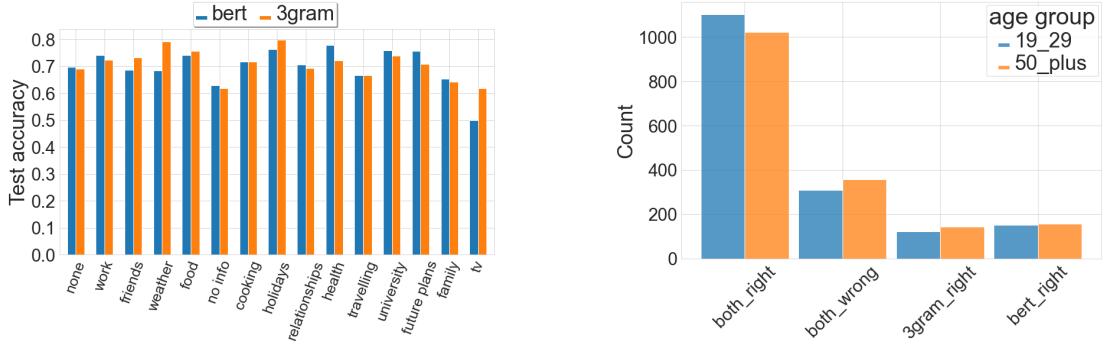
Table 3.2: Discourse dataset. Test set results averaged over 5 random initializations. Format: *average metric (standard error)*. Values in **bold** are the highest in the column; in **blue**, the second highest. *: F_1 is actually 0/0.

Model	Accuracy ↑ better	$F_1^{(19-29)}$ ↑ better	$F_1^{(50+)}$ ↑ better
Random	0.500	0.500	0.500
unigram	0.701 (0.007)	0.708 (0.009)	0.693 (0.004)
bigram	0.719 (0.002)	0.724 (0.003)	0.714 (0.003)
trigram	0.722 (0.001)	0.727 (0.003)	0.717 (0.001)
LSTM	0.693 (0.003)	0.696 (0.005)	0.691 (0.007)
BiLSTM	0.691 (0.009)	0.702 (0.017)	0.679 (0.007)
$BERT_{frozen}$	0.675 (0.003)	0.677 (0.008)	0.673 (0.010)
$BERT_{FT}$	0.729 (0.002)	0.730 (0.011)	0.727 (0.010)

Table 3.3: Dialogue dataset. Test set results averaged over 5 random initializations. Format: *average metric (standard error)*. Values in **bold** are the highest in the column; in **blue**, the second highest.

age	both correct	both wrong	$BERT_{FT}$ correct trigram wrong	trigram correct $BERT_{FT}$ wrong
19-29	oh that's cool	A retrospective exhibition	what even on the green slope?	really?
19-29	a text and then I'll do it	chuck them in those pots	yeah you told me to do you told me to do somebody made the f***ing table	and she like won't eat any carbs and she's like do you not like total greens?
19-29	yeah	mm		
50+	I said no I don't have them	yeah	really?	my under stairs in the kitchen
50+	that's of course	no no that's alright	it's still we we frequently walk that way	in the first place
50+	oh right	what a tragic life	since this this was new this house?	thank you very much

Table 3.4: Examples where both models are correct/wrong or only $BERT_{FT}$ /trigram is correct.



(a) BERT_{FT} and trigram test accuracies per topic for most frequent topics (including none/no info).

(b) Distribution of predicted cases by trigram and BERT_{FT} models for dialogue, split by age groups.

Figure 3.3

3.4 Age Detection Analyses

We focus our analysis on dialogue. In particular, we compare the two best-performing models, namely BERT_{FT} and the one using trigrams, and aim to shed light on what cues they use to solve the task. We first analyze how these models perform with respect to utterances of various topics. Secondly, we compare the prediction patterns of the two models, which allows us to highlight easy and hard examples. Finally, we focus on the trigram model and report the n-grams that turn out to be most informative to distinguish between age groups.

3.4.1 Performance Against Topic

As described in Section 3.1.1, each utterance in the dialogue dataset is annotated with one label which is representative of its topic.² This information is not explicitly available to the models. To explore how the two models deal with utterances in different topical contexts, we compare the accuracy they achieve on the 15 most frequent topics. The results are shown in Figure 3.3a. Two main observations can be made: Firstly, some topics turn out to be generally easier/harder than others, i.e., both models achieve higher/lower performance. To illustrate, both models achieve an accuracy well above 70% on topics like *food*, *holidays* or *university*, while topics such as *tv*, *family* or *no info* appear to be generally more challenging for both models. While this could be due to (a combination of) various factors, one intuitive possibility is that certain topics allow for more discriminative language features, which could be at the level of the lexicon or the style used to talk about them.

²In particular, it represents one of the utterance's topic, i.e., the one most frequently used in the whole data.

Secondly, some topics appear to be easier for one model rather than the other, and *vice versa*. To illustrate, the trigram model outperforms BERT on the topics *weather*, *holidays* and *tv*, while an opposite pattern is observed for *work*, *health*, and *future plans*. We conjecture that these patterns could be indicative of different strategies and cues exploited by various models to make a prediction. We explore this issue more in-depth in the following section, where we compare the predictions by the two models and qualitatively inspect some examples.

3.4.2 Comparing Model Predictions

We split the data for analysis by whether or not both models make the same correct or incorrect prediction, or whether they differ. Table 3.6 shows the breakdown of these results. As can be seen, a quite large fraction of samples are correctly classified by both models (63.17%), while in 19.78% cases neither of the models make a correct prediction. The remaining cases are comparably split between cases where only one of the two is correct, with BERT slightly outperforming Trigram by 1.23 percentage points. As shown in Figure 3.3b, the 19-29 age group appears to be slightly easier compared to the 50+ group, where models are observed to make more errors: the trigram misclassifies 50+ utterances 1.12 times as often as 19-29 utterances, and 1.17 times as often by $BERT_{FT}$.

To qualitatively inspect what the utterances falling into these classes look like, in Table 3.4 we show a few cherry-picked cases for each age group. We notice that, not surprisingly, both models have trouble with backchanneling utterances consisting of a single word, such as *yeah*, *mm*, or *really?*, which are used by both age groups. For example, both models seem to consider *yeah* as a ‘young’ cue, which leads to wrong predictions when *yeah* is used by a speaker in the 50+ group. As for the utterance *really?*, BERT_{FT} assigns it to the 50+ group, while the trigram model makes the opposite prediction. This indicates that certain utterances simply do not contain sufficient distinguishing information, and model predictions that are based on them should therefore not be considered reliable.

This seems to be particularly the case for short utterances. Indeed, through comparing the average length of the utterances incorrectly classified by both models (rightmost column of Table 3.6), we notice that they are much shorter than those belonging to the other cases. This is interesting, and indicates a key challenge in the analysis of dialogue data: on average, shorter utterances contain less signal. On the other hand, short utterances can provide rich conversational signal in dialogue; for example, backchanneling, exclamations, or other acknowledging acts. As a

consequence, using length alone as a filter is not an appropriate approach, as it can remove aspects of language use key to differentiating speaker groups.

3.4.3 Most Informative N-grams

	19-29		50+
coef.	n-gram	coef.	n-gram
-3.20	um	2.37	yes
-2.84	cool	2.12	you know
-2.58	s**t	2.09	wonderful
-2.12	hmm	1.90	how weird
-2.09	like	1.84	chinese
-2.02	was like	1.73	right
-1.96	love	1.71	building
-1.96	as well	1.66	right right
-1.88	as in	1.55	so erm
-1.84	cute	1.43	mm mm
-1.82	uni	1.41	cheers
-1.79	massive	1.39	shed
-1.79	wanna	1.37	pain
-1.79	f***k	1.36	we know
-1.72	tut	1.08	yeah exactly

Table 3.5: [L: Fix scope between table and caption.] For each age group, top 15 most informative *n*-grams used by the trigram model. **coef.** is the coefficient (and sign) of the corresponding *n*-gram for the logistic regression model: the higher its absolute value, the higher the utterance's odds to belong to one age group. * indicates masking of foul language.

Analyzing the most informative *n*-grams used by the trigram model allows us to qualitatively compare the linguistic differences inherent to each age group. In Table A.2 we report the top 15 *n*-grams per group. We find, firstly and intuitively, that colloquial language seems somewhat generational, with unigrams particularly indicative of younger speakers consisting of words such as *cool* and *massive*, and for older speakers, words like *wonderful*. These unigrams are both informative to the model and indicative of differences in both formality and ‘slang’ use across age groups.

These most informative *n*-grams also indicate differences in back-channeling use between age groups; younger speaker’s language is more characterized by the use of *hmm*, *um*, *yeah course*, while the top *n*-grams in the older category will more likely use *yes*, *right*, *right right*. A feature of younger language also apparent from these examples is in their use of more informal language: *yeah course* rather than *yes*. This informal language use also extends to the use of foul language, which make up a percent of the most informative unigrams shown in Table A.2.

Interestingly, while topic words make up many of the most informative n -grams for older speakers in Table A.2, younger speakers are more defined by their use of slang words such as *wanna*, foul language, or adjectives such as *cute*, *cool*, and *massive*. A key finding from Schler et al. [2006] is in the sentiment of language playing an important role, something which some of the most informative n -grams suggest may also be true for the dialogue dataset. As Table A.2 demonstrates, younger speakers use more dramatic language such as negative foul words, and positive *love*, *cute*, *cool*; all words with a strong connotative meaning. This prompts us to hypothesize that further inspection is needed to determine whether the same sentiment pattern will be true of dialogue as it has been reported to be in discourse.

	% cases	avg. length (\pmstd)*
both correct	63.17%	13.51 (\pm 18.98)
both wrong	19.78%	5.82 (\pm 8.33)
only Trigram correct	7.91 %	10.44 (\pm 11.66)
only BERT correct	9.14 %	11.53 (\pm 12.12)

Table 3.6: Percentage (% cases) of (non-)overlapping (in)correctly predicted cases between trigram and BERT_{FT}. *Utterance length measured in tokens.

Chapter 4

Experiment 2: Age-Adaptive Dialogue Generation using PPLM

4.1 Introduction

4.1.1 Recap of previous chapter and Connection to Experiment 1

In Experiment 1, we studied the extent to which text-based NLP models are able to detect speaker and writer age-related linguistic features in dialogue and discourse data. We then investigated which features drive the predictions made by those models. We showed that a fine-tuned version of BERT, $BERT_{FT}$, is capable classifying language by speakers from different age groups based on linguistic features. We also found that much simpler models based on n -grams achieve classification performance comparable to that of $BERT_{FT}$, suggesting that, in dialogue, local features (i.e., at the lexical level) can also be indicative of a speaker’s age. This was shown to be the case, as both lexical and stylistic cues seem to be informative to these models when predicting speaker age from dialogue utterances.

Now, in Experiment 2, we aim to test whether it is possible to generate dialogue responses that possess these age-indicative features, identified in Experiment 1. More specifically, we seek to use Plug-and-Play language models (PPLM) [Dathathri et al., 2020] to develop controlled dialogue generation models that can produce conversation responses contain linguistic elements of the speaking style of a certain age group to the extent that it would be classified as such by Experiment 1’s best classifier. The use of the PPLM method is motivated by its capability to control the style of the text generated by a large pre-trained Transformer-based language

model (e.g., GPT-2 or DialoGPT), without the computationally expensive requirement to re-train or fine-tune it. Namely, the Plug-and-Play approach uses substantially smaller and less costly to train attribute models to make perturbations to the large language model’s activation space, thereby shifting its output distribution towards a desired style. It is important to realize that the PPLM method is fundamentally different from fine-tuning in that it leaves the parameters of underlying language model unchanged.

The analyses of Experiment 1 largely concerned a comparison of performance between simpler n -gram-based models and more sophisticated neural discriminators (i.e., BERT). This comparison is continued throughout Experiment 2, because now we aim to compare the applicability of PPLM for age-adaptation when adaptation is enforced by an aforementioned attribute models, which can be a simple bag-of-words (BoW) model, or a more complex neural discriminator. Using PPLM to control the writing style of generated output text using the former is called *BoW-based control*, and using the latter is referred to as *discriminator-based control*. We also use the insights gained from Experiment 1 (e.g., about the applicability of n -gram-based representations of age-specific speaking style) to inform us about the development of attribute models. Furthermore, in Experiment 1, the classifiers are trained on the spoken component of the BNC, which is a dialogue dataset. Now, in Experiment 2, it is the goal that we generate something similar to a dialogue turn, i.e., a response to a dialogue prompt, in the style of a certain age group. Table 4.1 shows examples of how age-adapted dialogue responses to a prompt might look.

PROMPT	Can we talk?
Age 19-29 style response	Like about what?
Age 50+ style response	What would you like to share about?

Table 4.1: Examples of controlled dialogue responses to a prompt. The responses have been generated by DialoGPT after being controlled to generate younger sounding (red row) and older sounding (blue row) language.

It is also important to realize that the previously seen classifiers from Experiment 1 are not compatible with PPLM, because the method requires a bag-of-words or linear discriminator. Furthermore, one of the goals of Experiment 1 was to develop the best performing age-classifiers, which are ultimately used for (1) evaluation of the attribute adherence of the generated responses in Experiment 2, and (2) informing us about the choice of BoW-based attribute models (i.e., most informative unigrams). It also adds to the generality of our results to use a separate classifier for evaluation than the discriminators used as attribute models.

Lists of unigrams are used as BoWs (as opposed to lists of trigrams from the best performing n -gram-based classifier in Table 3.3), because the PPLM setup is not compatible with lists of n -grams for $n > 1$, as it relies on perturbations at the unigram-level. Making a PPLM-system compatible with, e.g., trigrams, would amount of retraining the entire underlying language model (like GPT-2), thereby defeating the purpose of PPLM, i.e., leveraging large LMs for controllability, without incurring significant re-training costs. However by allowing the best-performing classifiers to inform decisions about generation, we use the best unigram-based classifier’s list of most informative features as an attribute model.

4.1.2 Research Objectives, Hypotheses, and Contributions

Experiment 2 concerns a Plug-and-Play approach to age-adaptive dialogue generation. The aim is to use PPLM for age-adaptive dialogue generation. We hypothesize this to be possible to the extent that a text-based NLP model can detect in the generated dialogue responses linguistic features learned to be associated with certain age groups. Furthermore, we expect discriminator-based control to be more detectable and invasive at the structural level than BoW-based control. We also expect BoW-based control to take a smaller toll on the fluency of generated response than discriminator-based control.

The Plug-and-Play approach has previously been used by Dathathri et al. [2020] and Madotto et al. [2020] for language generation, controlled for very concrete writing styles and topics (e.g., negative/positive sentiment, politics, religion, business, tech). We apply the PPLM method to a novel problem, that is, age-related language generation conditioned on dialogue data. More concretely, we extend the work of Dathathri et al. [2020] in several important ways: (1) we control the generated language for more abstract writing styles, i.e., age-related linguistic style; (2) we use PPLMs to generate dialogue responses; (3) we propose two empirical methods for age-specific BoW development, as opposed to the manually constructed BoWs used in previous work; (4) we carry out in depth analyses about the relationships between evaluation measures of dialogue response quality and style; (5) finally, we address and study the problem of writing style biases induced by the prompt’s style.

It must be noted that the activation perturbation steps necessary for PPLMs to exert influence on the style of generated text make too slow to be suitable for online interactive conversational applications. This problem is addressed by Madotto et al. [2020], who propose Plug-and-Play Conversational Model (PPCM), an extension of PPLM that uses separately trained residual

adapter modules to speed up the PPLM-method, making it suitable for real-time dialogue response generation. Despite their work also being about controlled dialogue generation, our work distinguishes itself in important ways. They control their output for the same styles and topics as Dathathri et al. [2020], and do not experiment with more abstract writing styles, like age-related linguistic characteristics. Overall, this thesis is about establishing whether age-related linguistic features are automatically detectable, and if we can enforce them as a writing style of dialogue responses using PPLMs. The work of Madotto et al. [2020] addresses the problem of interactivity (i.e., speed of response generation) requires solving a separate, very valuable, engineering problem, which is unfortunately beyond the scope of this thesis. Nevertheless, our previously mentioned contributions still hold.

The rest of this chapter is structured as follows. Section 4.2 describes methodology and experimental setup of our controlled dialogue generation experiments. It also covers details about attribute model development, the choice of prompts, and an explanation of the chosen evaluation metrics. The results of the language generation experiments are presented and interpreted in Section 4.3, followed by the outcomes of various quantitative and qualitative analyses in Section 4.4. A separate section about data is omitted in this chapter, because we use the previously mentioned dialogue dataset for these experiments. The reader is directed to Sub-section 3.1.1 for a detailed description of that corpus, and our pre-processing steps.

4.2 Methodology and Experimental Setup

This section describes and motivates the methodology, experimental details of our controlled dialogue generation experiments, and the most important components involved: attribute model development, dialogue response generation, choice of prompt, and evaluation of generated responses. All computations for attribute model development, dialogue response generation, and evaluation are performed on an NVIDIA TitanRTX GPU.

The overall goal of this chapter, Experiment 2, is controlled dialogue generation. Controlled dialogue generation experiments are performed using PPLM-setups that differ with respect to **(1)** pre-trained Transformer-based language model (GPT-2 or DialoGPT), **(2)** type of attribute model (BoW or discriminator), **(3)** style attribute (young, old, or uncontrolled), and **(4)** the style of the prompt on which generation is conditioned (young, neutral, or old). Note that when the style attribute is “uncontrolled”, it implies the unaltered original underlying language model (i.e., GPT-2 or DialoGPT) is recovered.

4.2.1 Attribute Model Development

As mentioned in the introduction to this chapter, in the PPLM-method, controlled generation is achieved by using an attribute model to make activation space perturbations to an underlying language model. This alters the output distribution of the language model to make it more likely to generate text that aligns with a predefined writing style. Such an attribute model is significantly smaller and less costly to train than the underlying large-scale language model (e.g., GPT-2 or DialoGPT). An attribute model can either be a simple bag-of-words, or a more complex neural discriminator. When developing an attribute model, either a discriminator is trained on the dialogue data, or a bag-of-words is statistically constructed from the same the dialogue dataset (See Section 3.1.1).

Bag-of-Words A simple bag-of-words can also serve as an attribute model. We extend the approach of Dathathri et al. [2020] that relies on curated wordlists, by applying two empirical approaches to extract wordlists from the dialogue corpus. An empirical approach has the benefit of being more reproducible, and not requiring a domain expert to manually curate a list. In the first approach, the BoW consists of the 100 most informative unigrams of the unigram classifier used during the text classification experiments (See Table 3.3 for the results). The most informative unigrams per age groups are deemed the most distinguishing features by the unigram classifier. They could therefore be used to make sensible perturbations to a language model’s output, yielding more effective control.

The second method of wordlist extraction is fully frequency-based. The goal of this extraction method is to yield two distinct sets of words that are representative of each age group’s language. The frequency-based wordlists per age group are constructed from the *imbalanced* dialogue dataset. Recall that the 19-29 age group is over-represented in the original imbalanced dialogue dataset, and we discarded utterances from that class to balance the dataset when using it to train classifiers. The use of the original dataset for wordlist extraction is motivated by the fact that the original dataset contains more utterances, and a frequency-based wordlist from that corpus would be a better representation of the younger age group’s speaking style. Wordlist extraction is carried out as follows: Per age group, all unique words are ordered by frequency of occurrence in all utterances. For both ordered lists of word counts, the most frequent words are kept that make up for at least 85% of the cumulative occurrences. Then, the words are removed that appear in both lists (i.e., the overlapping set of words is discarded). Of the resulting two non-overlapping

ordered lists of words and their numbers of occurrences, only the words are kept that account for at least 85% of the respective wordlist’s cumulative occurrences. The resulting lists consist of 56 (young-related), and 92 (old-related) words. The 85-th percentile cutoff points are chosen to yield wordlists of similar lengths as those used by Dathathri et al. [2020]. Both pairs of wordlists are included Section A.2.

Discriminators When training the discriminators, the dialogue dataset is randomly split into a training (90%) and test (10%) set. The frozen embeddings of either GPT-2-medium [Radford et al., 2019] or DialoGPT-medium [Zhang et al., 2020] are fed into trainable linear classifiers, seeking to distinguish between transcribed dialogue utterances from young (ages 19 to 29) and old (ages 50 and over) speakers. The discriminators are trained using Adam [Kingma and Ba, 2015] with a learning rate of $1 \cdot 10^{-4}$ and default values for all other parameters from PyTorch’s implementation of Adam¹, with a maximum sequence length of 512 tokens, for 20 epochs, and a batch size of 64. The discriminator parameters are used of the epoch with the highest test accuracy (67.4% for GPT-2-based discriminator, and 67.6% for DialoGPT-based discriminator). Note that there is no validation set, because the discriminator attribute models have no hyperparameters to be tuned. Namely, they are single-layer linear classifiers, whose input layer size is equal to the output layer size of the underlying language model (GPT-2 or DialoGPT).

4.2.2 Prompts

The prompt is the text on which a language generation model conditions its output. In the specific case of controlled language generation, given a prompt, `prompt`, a predefined style attribute a , and some controlled language generation model parameterized by θ , generating a style-controlled piece of text \mathbf{x} entails modeling $p_\theta(\mathbf{x}|a, \text{prompt})$. The output distribution of controlled generation model p_θ is therefore likely to be affected by `prompt`, which could manifest itself as the prompt’s style influencing the output text’s style. In line with this presumption, Fan et al. [2018] and Lester et al. [2021] show that the content and style of prompts can influence the output of neural text generation models.

Because our research is about the style of generated dialogue responses, we wish to take into account the effect that the style of the prompt might have. We therefore distinguish between three categories of prompt style: young, neutral, and old. We use 3 random initializations of

¹<https://pytorch.org/docs/stable/generated/torch.optim.Adam.html#torch.optim.Adam>

BERT_{FT} classifiers from Experiment 1 (different from the one used for evaluation) to assign target probabilities to prompts written by us. A young prompt is defined as having an average assigned “young” probability (i.e., BERT_{FT} ’s assigned probability of containing features learned to be associated with the younger age group) of at least 80%, with a standard deviation no greater than 0.15. The same holds for an old prompt, but for its assigned “old” probability. A neutral prompt must have an average assigned young (or old, as they are complementary values) probability between 40% and 60%, with a standard deviation no greater than 0.15. We use 5 prompts per style class, and they are shown, along with their average target probabilities and standard deviations, in the Appendix in Table A.1.

The focus of the main quantitative results (covered in Section 4.3) is to see how much the writing style can be adapted towards that of an age group, so it uses neutral prompts, to avoid any unwanted biases induced by the prompts. The effects of using strongly styled prompts are analyzed in Section 4.4.3. It is worth mentioning that it also possible to use PPLM in an unprompted setting, i.e., unconditioned language generation. In practice, this is done by conditioning the generated output on the `<| endoftext |>` token. We also present quantitative results for this unprompted setting in the Appendix, in Section A.7.

4.2.3 Generation

Every PPLM-setup generates 30 sequences per output length 6, 12, 24, 30, 36, 42, 48, 54, and 60 tokens. The lengths are chosen to provide sufficiently distanced sub-samples for later analyses about the effects of generated response length on the quality and control of our PPLM-setups. Sub-sample sizes of 30 are chosen to satisfy the Central Limit Theorem (CLT), making it possible to assume the distribution of sub-sample averages \bar{X}_n approximates the normal distribution with mean μ and variance $\frac{\sigma^2}{n}$ [CLT, 2008], where n is the sub-sample size. Note that perturbations of the base language model’s output can affect the controlled sequence length, so the final sequence length may differ by a few tokens from the given one. All other PPLM-parameters are kept at their default values, as recommended by Dathathri et al. [2020].

Each reported PPLM-configuration represents the best initialization, if the term applies. The pre-trained language models are kept equal across configurations, as using different initializations of these large models is infeasible, and would defeat the purpose of PPLM. A BoW-based setup uses a single list of words as an attribute model, thereby not having random parameters to initialize. And finally, discriminator-based setups use comparatively small linear classifiers (a

few hundred parameters), the initialization effects of which have been found to be negligible on performance.

4.2.4 Evaluation

The generated sequences are evaluated along two main axes: fluency and control. Fluency refers to the degree to which a text passage appears natural, grammatical, and non-repetitive. Control is the extent to which the produced language resembles that of the attribute being controlled for. Evaluation is done using a set of automated metrics. Fluency is measured automatically by perplexity (denoted ppl) with respect to a different language model, GPT-1 [Radford et al., 2018], expressed as

$$ppl(\mathbf{x}) = \exp \left\{ -\frac{1}{t} \sum_i^t \ln p_{\theta}(x_i | x_{<i}) \right\}. \quad (4.1)$$

\mathbf{x} represents a sequence of tokens, t is sequence length, x_i is the i -th token, and θ denotes GPT’s parameters. Perplexity is a measure of a language model’s uncertainty when posed with the task of predicting a succession of words. Assuming a language model to be a reliable representation of relationships within a natural language, low perplexity can serve as a rough proxy for fluency of a text. However, a major caveat of perplexity is that it only measures uncertainty w.r.t. one language model, making it less generalizable. To slightly reduce this effect, we choose to evaluate perplexity with respect to a different language model (GPT-1) than the one used for generation (GPT-2 or DialoGPT).

Furthermore, the normalized number of distinct unigrams (Dist-1), bigrams (Dist-2), and trigrams (Dist-3), are used as measures of text diversity. Experiment 1’s best BERT-classifier’s classification accuracy on a set of sequences generated by a single generation model is used as an automated measure of attribute control. It can be seen as a proxy for control, because it indicates how resemblant of an age group’s speaking style a generated text is deemed to be.

Two types of baselines are used when evaluating text generation performance: a pre-trained model baseline (i.e., a GPT-2-baseline or DialoGPT-baseline), and a corpus-specific baseline. Therefore all controlled generation models that use GPT-2 as their language model, should be compared to the uncontrolled GPT-2 baseline. The same holds for DialoGPT. The second type of baseline combines the underlying language model with a bag-of-words consisting of the 100 most common words in the balanced dialogue corpus, irrespective of age. This setting is included

to give an indication of how biased the balanced BNC’s frequently occurring words might be towards a specific age group.

The deliberate choice to use BERT-based models for classification (and evaluation of generated output), and GPT-based models for generation is motivated by the following reasons. BERT’s encoder-based bidirectional architecture makes it more suitable for sequence classification than for generation [Devlin et al., 2019]. By similar reasoning, GPT’s decoder-based structure makes it a more suitable choice for generation tasks. Furthermore, Experiment 1’s best classifier is a fully fine-tuned BERT model (ca 110M parameters). However, it is computationally highly expensive to fine-tune GPT-2-medium (ca 345M parameters) which makes it infeasible given our computational resources. Finally, using a separate model class (i.e., BERT) for evaluation of GPT-based generation models makes the results more generalizable.

4.3 Controlled Dialogue Generation Results

The quantitative results of generating younger (19 to 29) and older (50 and over) sounding responses to neutral prompts are reported in Tables 4.2 and 4.3, respectively. The metrics per row averaged over $N = 30 \cdot 9 = 270$ samples (See Section 4.2.3). In these tables, the underlying language model being used in a setup (i.e., a row in a table) is indicated by the prefixes G- for GPT-2 and D- for DialoGPT. Additionally, both tables report the results of the unperturbed GPT-2 and DialoGPT baselines (labeled G-baseline and D-baseline, respectively), and those of the 100 most common age-independent bag-of-words setups for both GPT-2 and DialoGPT (labeled G-100MCW and D-100MCW, respectively). The accuracies for these two setups (i.e., baseline and 100 most common words) are omitted because they do not aim to generate responses that resemble any target age group. Moreover, two bag-of-words (BoW) setups are reported per underlying language model (GPT-2 or DialoGPT): the frequency-based BoW setup (indicated by the suffix, -BoW_{FB}), and the 100 most informative unigram setup (indicated by the suffix, -BoW_{100MIU}). Detailed descriptions of how and why these aforementioned wordlists are constructed are provided in Section 4.2.1. Finally, the discriminator-based setups are indicated by the suffix, -Discrim. The aforementioned reporting conventions also hold for the tables containing the results of response generation to younger and older sounding prompts (i.e., Tables 4.4, 4.5, 4.6, and 4.7). The results of those experiments are discussed in Section 4.4.3.

As one would expect, Tables 4.2 and 4.3 show that the GPT-2 baseline consistently scores among the best on perplexity (best perplexity compared to young generation models, and second best compared to old-generation models) and diversity of generation (the Dist- $n|_{n=1,2,3}$ scores are almost consistently in the upper registers). Similarly, the unperturbed DialoGPT baseline also scores best in terms of perplexity, when compared to other DialoGPT-based setups. This means that the responses generated by GPT-2 and DialoGPT are found to be among the least perplexing to GPT-1. This is unsurprising, as GPT-2, and thereby also DialoGPT, are pre-trained in similar fashion to GPT-1 [Radford et al., 2018, 2019, Zhang et al., 2020]. Additionally, the target probabilities ($\bar{P}_Y = 0.62$ in Table 4.2, and $\bar{P}_O = 0.38$ in Table 4.3) indicate that the GPT-2 baseline is biased towards generating young language. That is, given a neutral prompt, GPT-2 is inclined to produce responses that are likely to contain features learned to be young by BERT_{FT}. This could be attributable to GPT-2 being pre-trained on WebText, a corpus of high-quality documents scraped from web pages, which could be over-represented by millennials². Moreover,

²<https://www.statista.com/statistics/272365/age-distribution-of-internet-users-worldwide/>

DialoGPT has very strong bias towards generating younger sounding responses, given a neutral prompt (0.76 average probability to contain detectable young features). This is most likely due to DialoGPT having been fine-tuned on Reddit threads [Zhang et al., 2020], as the majority of Reddit users are between the ages 20 and 29³.

The G-BoW_{100MCW} setup performs on par with baseline w.r.t. perplexity and diversity (second best perplexity, second best Dist-2, best Dist-3), when compared to the young generation models. A similar pattern can be seen for old models, where G-100MCW has the second best Dist-2, and best Dist-3 scores. Additionally, the target probabilities seem virtually unaffected by the 100MCW setups, suggesting that perturbing GPT-2’s output with an age-agnostic bag-of-words of the most frequently used words in the dialogue dataset does not noticeably shift the writing style towards that of the younger or older age group. This is to be expected, as such a wordlist should be an unbiased representation of the dataset’s language style, given similar sample sizes per class.

BoW-based models seem to generate responses that are slightly more likely to contain features of the target age (young GPT2-BoW_{FB} models result in 0.06 target probability improvement over the baseline, and old GPT2-BoW_{FB} model in 0.04 target probability improvement over baseline). However, these differences could also be due to randomness. For the 50-plus style GPT-2-based response models, the BoW_{100MIU} setup does not manage to generate older sounding language than the baseline. Furthermore, BoW-based models seem to barely impact the syntactical structure of generated responses, because changes are more local than those made by discriminator-based setups Dathathri et al. [2020]. This is confirmed by the barely altered perplexity and Dist-scores.

The GPT-2 discriminator-based old-style model manages to generate responses that more convincingly resemble the style of the target group (0.38 \bar{P}_O improvement over the GPT-2 baseline). However, this comes at the cost of perplexity (+19.65 compared to baseline) and diversity (much lower dist-scores and higher corresponding standard deviations). By contrast, the GPT-2 discriminator-based young setup does not generate convincingly more young-sounding responses (only 0.04 average target probability improvement over baseline), despite having noticeably worse and less precise perplexity (4.59 increase over baseline) and diversity.

The unperturbed DialoGPT baseline produces more perplexing and less diverse text than GPT-2 according to GPT-1 perplexity and the Dist- $n|_{n=1,2,3}$ scores. The higher perplexity is to be

³<https://www.statista.com/statistics/1125159/reddit-us-app-users-age/>

expected, as DialoGPT pre-training and fine-tuning method deviates more from GPT-1’s than GPT-2, making it likely to produce more unexpected sentences [Zhang et al., 2020]. When using DialoGPT as an underlying language model, the BoW-based models (including 100MCW) seem to reinforce young-bias for DialoGPT, regardless of age (i.e., for all young-targeted BoW-based models \bar{P}_Y goes up, and \bar{P}_O goes down for all old-targeted BoW-based models). And the discriminator-based DialoGPT models, are superior w.r.t. target probability (+0.10 \bar{P}_Y difference compared to DialoGPT baseline for young-models, and +0.34 \bar{P}_O compared to baseline for old-models). However, this again comes at the cost of much higher and more volatile perplexity. Overall, according to these results it seems to be possible to control dialogue responses for a certain age group. The used underlying language models are biased (in varying degrees) towards generating younger-sounding language. In these tested PPLM-setups, there seems to be a tradeoff between increased control and decreasing perplexity and diversity of generated language. Furthermore, the BoW-based models achieve less detectable levels of control, but preserve the fluency and diversity of generated text. In other words, the discriminator-based models make more invasive changes to the unperturbed sentences, which can result in less fluent and more repetitive text. However, they do produce more detectably age-appropriate passages.

Model	ppl. ↓ better	Dist-1 ↑ better	Dist-2 ↑ better	Dist-3 ↑ better	\bar{P}_Y ↑ better	Acc. ↑ better
G-baseline	27.50 (6.58)	0.87 (0.09)	0.94 (0.04)	0.90 (0.06)	0.62 (0.42)	-
G-100MCW	27.56 (6.60)	0.86 (0.10)	0.93 (0.04)	0.90 (0.05)	0.63 (0.42)	-
G-BoW _{FB}	27.91 (7.18)	0.87 (0.10)	0.93 (0.05)	0.90 (0.06)	0.69 (0.41)	70.4%
G-BoW _{100MIU}	28.37 (7.31)	0.87 (0.09)	0.93 (0.04)	0.90 (0.06)	0.67 (0.41)	67.4%
G-Discrim	32.09 (18.98)	0.77 (0.20)	0.86 (0.13)	0.84 (0.15)	0.66 (0.43)	67.8%
D-baseline	37.52 (12.06)	0.86 (0.13)	0.90 (0.08)	0.85 (0.10)	0.76 (0.37)	-
D-100MCW	37.80 (10.89)	0.85 (0.14)	0.89 (0.10)	0.85 (0.10)	0.82 (0.33)	-
D-BoW _{FB}	38.53 (12.64)	0.87 (0.12)	0.90 (0.08)	0.86 (0.10)	0.82 (0.33)	83.0%
D-BoW _{100MIU}	38.67 (11.70)	0.88 (0.11)	0.91 (0.07)	0.86 (0.10)	0.87 (0.28)	88.5%
D-Discrim	42.01 (16.94)	0.90 (0.12)	0.86 (0.14)	0.77 (0.22)	0.86 (0.29)	85.9%

Table 4.2: Results of age-controlled dialogue generation: **young**-targeted models, conditioned on **neutral prompts**. Format: *average metric (standard error)*. ppl. is perplexity w.r.t. GPT-1. Dist- n (for $n = 1, 2, 3$) is number of distinct n -grams normalized by text length, as a measure of diversity. \bar{P}_Y is the sample’s average probability to contain features learned to be young by BERT_{FT}. Acc. is BERT_{FT}’s accuracy when classifying the row’s samples.

Model	ppl. ↓ better	Dist-1 ↑ better	Dist-2 ↑ better	Dist-3 ↑ better	\bar{P}_O ↑ better	Acc. ↑ better
G-baseline	27.50 (6.58)	0.87 (0.09)	0.94 (0.04)	0.90 (0.06)	0.38 (0.42)	-
G-100MCW	27.56 (6.60)	0.86 (0.10)	0.93 (0.04)	0.90 (0.05)	0.37 (0.42)	-
G-BoW _{FB}	27.58 (7.07)	0.86 (0.10)	0.93 (0.04)	0.90 (0.06)	0.42 (0.42)	43.0%
G-BoW _{100MIU}	27.25 (6.15)	0.87 (0.09)	0.93 (0.04)	0.90 (0.06)	0.38 (0.42)	37.4%
G-Discrim	47.15 (47.56)	0.73 (0.24)	0.75 (0.28)	0.75 (0.27)	0.76 (0.36)	74.3%
D-baseline	37.52 (12.06)	0.86 (0.13)	0.90 (0.08)	0.85 (0.10)	0.24 (0.37)	-
D-100MCW	37.80 (10.89)	0.85 (0.14)	0.89 (0.10)	0.85 (0.10)	0.18 (0.33)	-
D-BoW _{FB}	37.85 (11.17)	0.87 (0.12)	0.90 (0.08)	0.86 (0.09)	0.22 (0.35)	21.5%
D-BoW _{100MIU}	37.91 (12.27)	0.87 (0.11)	0.90 (0.07)	0.85 (0.10)	0.22 (0.34)	21.9%
D-Discrim	41.17 (20.72)	0.87 (0.12)	0.89 (0.13)	0.83 (0.16)	0.57 (0.41)	56.7%

Table 4.3: Results of age-controlled dialogue generation: **old**-targeted models, conditioned on **neutral prompts**. Format: *average metric (standard error)*. ppl. is perplexity w.r.t. GPT-1. Dist-*n* (for $n = 1, 2, 3$) is number of distinct *n*-grams normalized by text length, as a measure of diversity. \bar{P}_O is the sample’s average probability to contain features learned to be young by BERT_{FT}. Acc. is BERT_{FT}’s accuracy when classifying the row’s samples.

4.4 Controlled Dialogue Generation Analyses

By means of quantitative and qualitative analyses, we seek to study which relationships affect the grammatical quality and attribute relevance of the generated responses. The discriminator-based setups, and the BoW-models with the highest average target probabilities (i.e., BoW_{FB} for GPT-2, and BoW_{100MIU} for DialoGPT) are considered for the analyses. In the following sections, we presents a series of analyses about the relationship between perplexity and target probability (Section 4.4.1), the effects of response length on generation quality (Section 4.4.2), the impact of the prompt’s style on generation style and quality (Section 4.4.3), and qualitatively observable patterns in generated samples 4.4.5.

4.4.1 The Relationship between Perplexity and Target Probability

Figures 4.1 and 4.2 show bar charts depicting the relationship between the average target probability (*y*-axes) and perplexity (*x*-axes) assigned to the samples generated by various PPLM-model setups. The error bars around the average target probabilities are 95% confidence intervals. Based on the distribution of perplexity observed over all responses generated conditioned on neutral prompts, perplexity is binned into three consecutive intervals, bordered by the tertiles of distribution (i.e., the two points that divide the distribution of perplexities into three parts, each containing a third of the distribution). These intervals are named low perplexity ($\text{ppl} < 27.52$), medium perplexity ($27.52 \leq \text{ppl} < 35.63$), and high perplexity ($\text{ppl} \geq 35.63$).

The GPT2-based old models (Figures 4.1b and 4.1d) show a pattern of increasing perplexity coinciding with higher and more precise assigned target probabilities. Especially so for the old-targeted discriminator-based GPT-2 setup, whose responses with relatively high perplexity (50+) are (at a 5% level) significantly more likely than those with low perplexity to contain features learned to be old by $BERT_{FT}$. High-perplexity responses of this model are also assigned probabilities with more precision, as indicated by the narrower confidence regions. These observations suggest that, controlling the output of GPT-2 to contain more old-related features, the responses become more perplexing (and arguably less fluent). This is relationship more salient for discriminator-based control, because the perturbations are more invasive to the sentence structure, which is in line with earlier observations. It also suggests that PPLM prioritizes enforcing attribute control over promoting fluency, when making perturbations to the underlying language model’s output distribution. By contrast, we do not observe a clear pattern for the young-targeted GPT-2-based models’ responses (Figures 4.1a and 4.1c). Higher perplexity responses generated by young-targeted BoW-based GPT-2’s seem to coincide being less detectably young-sounding (with lower precision) than lower perplexity ones. And the responses generated the young-targeted discriminator-based GPT-2 PPLM setups appear to receive roughly equal average target probabilities for increasing perplexity, with slightly lower precision. It must be noted that there are no significant differences at the 5% level between the average \bar{P}_Y , so conclusions about the relationship between perplexity and target probability should be taken tentatively.

DialoGPT’s strong proclivity to generate younger sounding responses is noticeable in Figures 4.2a and 4.2c, as depicted by the high average target probabilities and relatively narrow confidence intervals. Furthermore, Figure 4.2a shows a similar pattern as the ones observed for the old-targeted GPT-2-based setups, that is, increasing detectability of target age features for increasing perplexity, with an almost significant difference in average assigned target probabilities between the low and high perplexity responses. Additionally, one could argue that the BoW-based DGPT young-targeted setup (Figure 4.2c) shows a similar, yet much less pronounced, pattern to its GPT-2-based counterpart (Figure 4.1c), where the average assigned target probability decreases for increasing perplexity. However, the differences between the average target probabilities are too small and statistically insignificant to draw a conclusion from this result alone.

The DGPT-based models targeted towards generating old-related features (Figures 4.2b and 4.2d) display patterns of decreasing target probabilities for increasing perplexity. However, it must

be emphasized that responses generated by both these models are rarely seen by BERT_{FT} as likely to contain linguistic features associated with their target age (overall \bar{P}_O of 0.22 and 0.57 for DGPT-BoW Old and DPGT-Discrim Old, respectively). This makes it less reliable to draw conclusions about the relationship between the style-adherence and perplexity of their samples, as they are often not style-adherent to begin with.

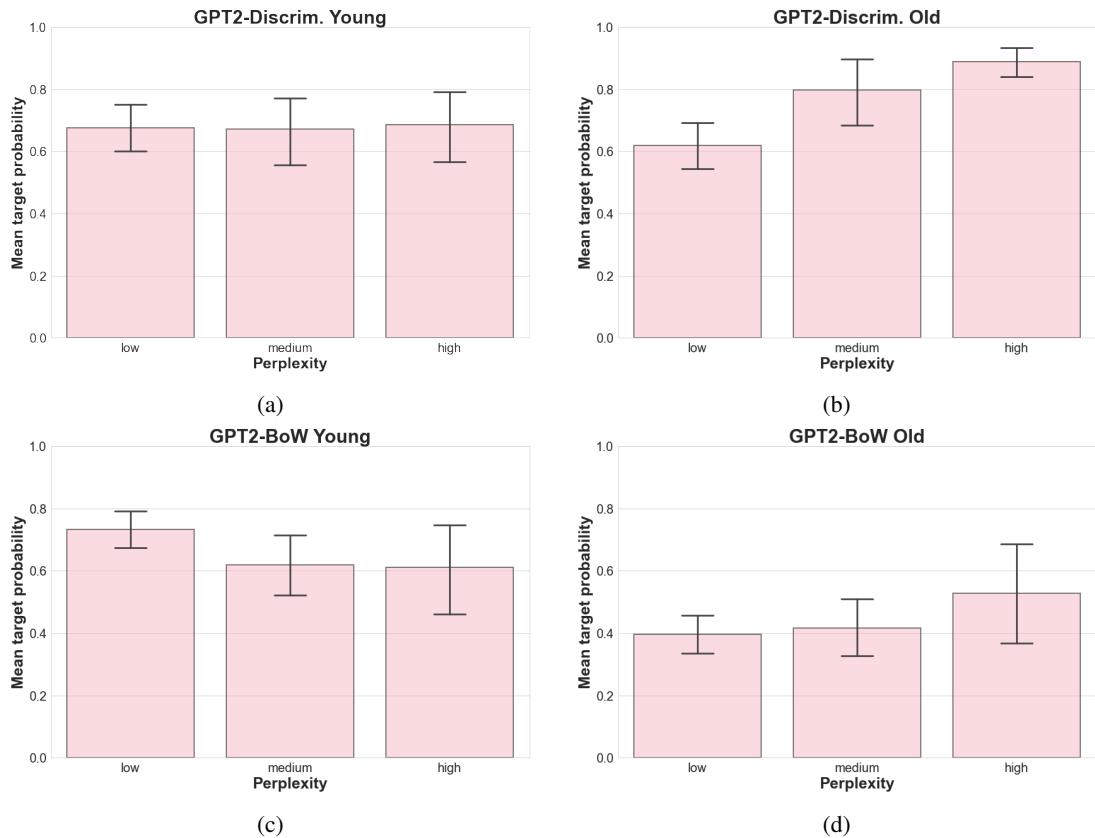


Figure 4.1: Mean target probability (x -axes) assigned to GPT-2-based models' samples by BERT_{FT} for increasing ranges of GPT-1 perplexity (y -axes). Error bars are 95% confidence intervals.

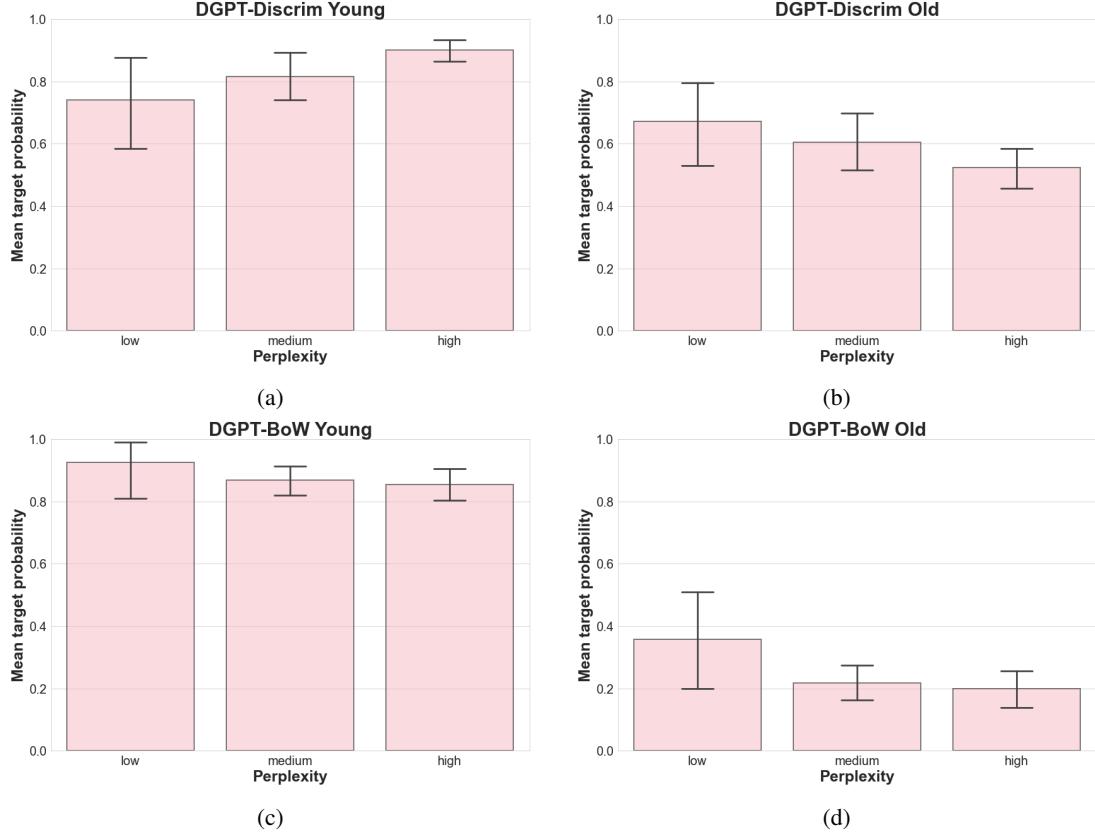


Figure 4.2: Mean target probability (y -axes) assigned to DialoGPT-based models’ samples by $BERT_{FT}$ for increasing ranges of GPT-1 perplexity (x -axes). Error bars are 95% confidence intervals.

4.4.2 The Relationship between Response Length and Various Evaluation Metrics

The number of tokens in a generated response coincides with noticeable differences in our automated evaluation metrics. It is therefore important to get a clearer picture of how the various measures for fluency and control change for different sequence lengths. Moreover, properly understanding these relationships can inform developers of adaptive dialogue systems about preserving output quality and adaptation of responses of arbitrary lengths.

Response length (on the x -axes) is plotted against various evaluation metrics in Figures 4.3 (average $BERT_{FT}$ accuracy), 4.4 (GPT-1 perplexity), 4.5 (normalized number of distinct unigrams), 4.6 (normalized number of distinct bigrams), and 4.7 (normalized number of distinct trigrams).

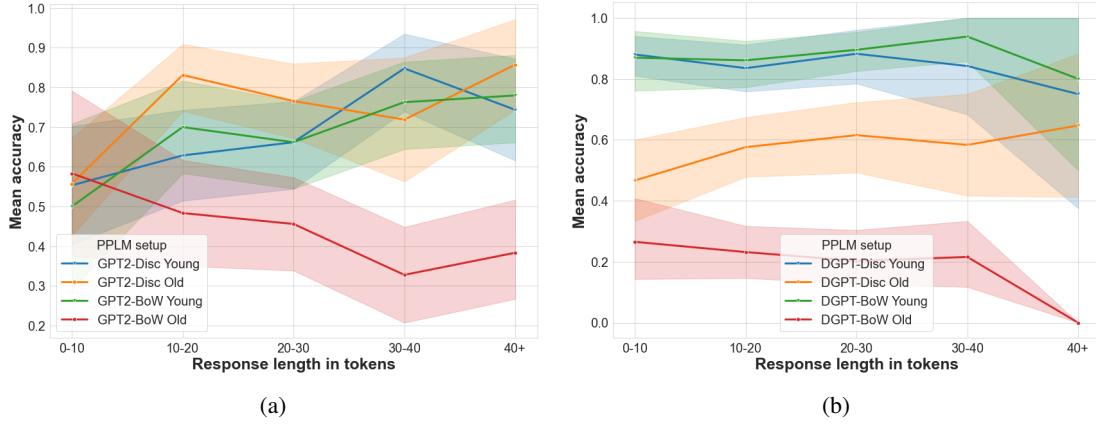


Figure 4.3: Mean BERT_{FT} accuracy. GPT-2-based models (left), DialoGPT-based models (right). Translucent error bands represent 95% confidence intervals. Plots best viewed in color.

Figure 4.3a shows a slight upward trend in average accuracy with greater uncertainty for increasing response length for all GPT-2-based models, except for the BoW-based old generation model. That is, longer sequences are, on average, slightly easier to classify, though with less precision. This is probably due to the fact that longer sentences contains more information to base predictions on. By contrast, the DialoGPT-based models in Figure 4.3 do not seem to show a clear general trend that mean accuracy follows for increasing response length. However, it does seem that DialoGPT’s strong bias towards generating younger sounding responses causes output from DialoGPT-based young generation models to be much easier to classify than that from the old generation models. Overall, it can be seen that the BoW-based old models (GPT-2 and DialoGPT) are significantly more challenging to classify at almost every length bracket.

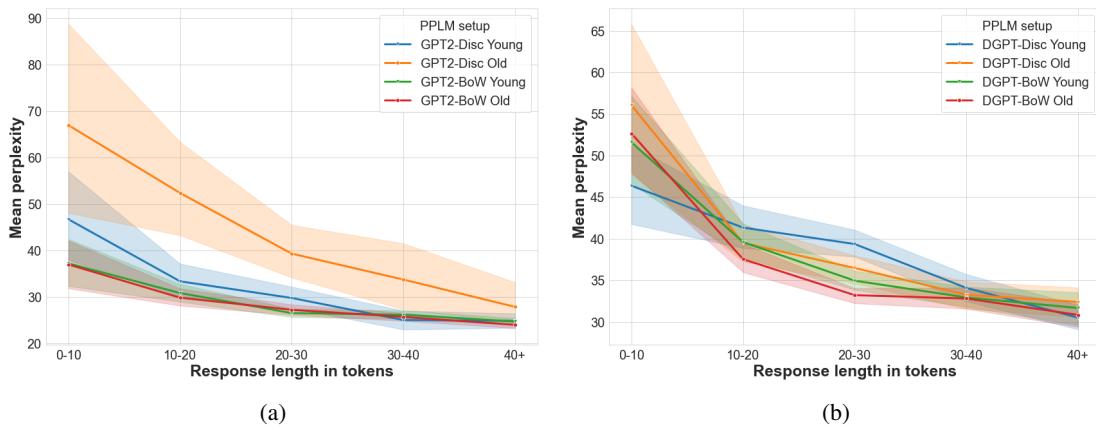


Figure 4.4: Perplexity. GPT-2-based models (left), DialoGPT-based models (right). Translucent error bands represent 95% confidence intervals. Plots best viewed in color.

Figure 4.4 shows a clear downward trend for both sets of PPLM-setups. Irrespective of the underlying language model being used, longer responses are deemed less perplexing, with more certainty, by GPT-1 than shorter ones. It is worth emphasizing that the model with the highest target probability improvement over its relevant baseline, G-Discrim_{Old}, is found to produce significantly more perplexing responses than its GPT-2-based counterparts at most length brackets (See Figure 4.4a). This finding resonates with Figure 4.1b and the idea, that especially for old-generation models, increased levels of attribute relevance coincide with worse perplexity.

However, it must be noted that the downward slope of perplexity for increasing response length could be attributable to the nature of calculating perplexity, rather than generation properties of the models. Namely, perplexity essentially averages the sum of the negative exponentiated probabilities $p(\text{word}|\text{context})$, for every word in a sentence. Because the context increases with every successive word, and larger contexts typically result in less uncertainty, shorter sequences are often given unfairly high perplexities.

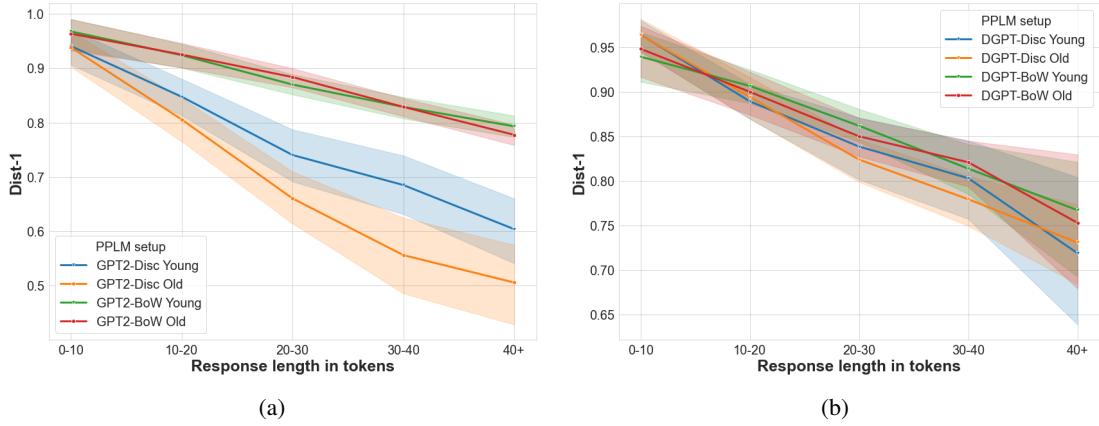


Figure 4.5: Dist-1. GPT-2-based models (left), DialoGPT-based models (right). Translucent error bands represent 95% confidence intervals. Plots best viewed in color.

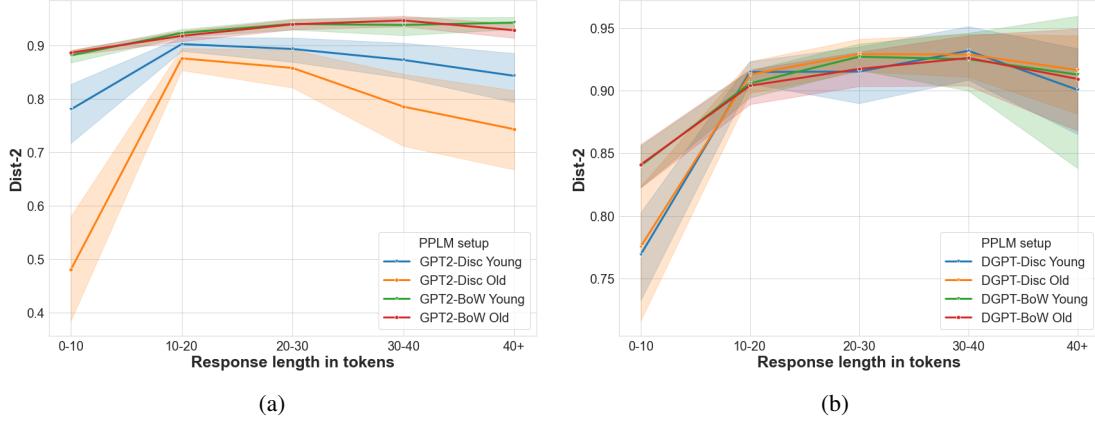


Figure 4.6: Dist-2. GPT-2-based models (left), DialoGPT-based models (right). Translucent error bands represent 95% confidence intervals. Plots best viewed in color.

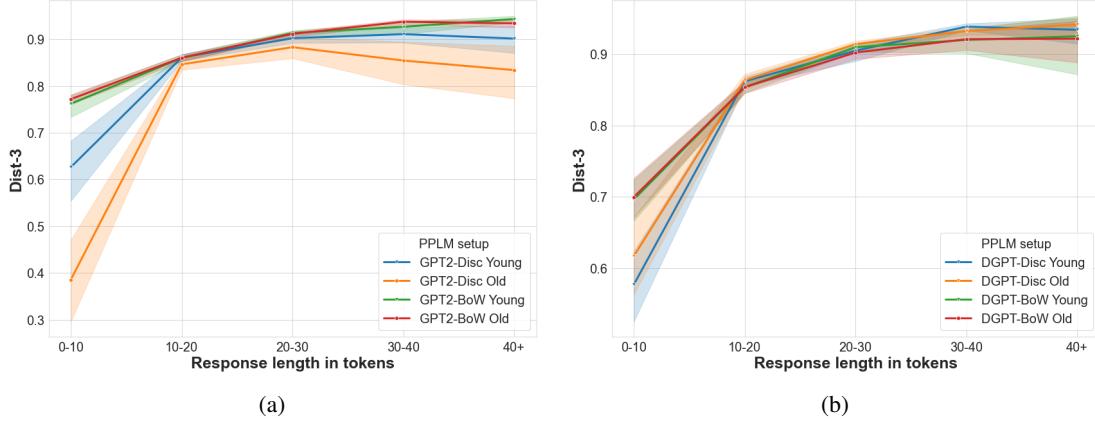


Figure 4.7: Dist-3. GPT-2-based models (left), DialoGPT-based models (right). Translucent error bands represent 95% confidence intervals. Plots best viewed in color.

Figure 4.5 shows that diversity w.r.t. unigrams decreases for longer responses. This is most likely due to the fact that longer sentences have an *a priori* higher probability of containing repeated words. E.g., stopwords like "the" and "of" are likely to appear multiple times in longer sentences. The same figure shows that GPT-2-BoW models generate significantly more diverse responses w.r.t. unigrams for almost every bracket of response length. This could be attributable to BoW-based control altering base-GPT2's generated sentences at the token-level, thus being more likely to preserve the unigram diversity of the unperturbed baseline (the GPT-2 baseline's Dist-1 is always in the upper register in Tables 4.2 and 4.3).

Figure 4.6 and shows that variety w.r.t. bigrams makes an initial upward jump between response length of 1-10 and 10-20. Dist-2 then follows a mild downward trend for both GPT2- and DialoGPT-based models. However, detailed inspection of Figure 4.5a shows that only the

discriminator-based setups have a negative slope, whereas the BoW-based setups follow a very slight upward trend. Thus, the BoW-based GPT-2 models produce significantly more diverse language w.r.t. bigrams for most response lengths, attributable to the same reason mentioned above.

Figure 4.7 shows similar patterns: an initial upward jump in trigram diversity between shortest and second-to-shortest length brackets. GPT-2 models then show a slight downward trend for the discriminator-based setups from the 20-30 lengths onward, while the BoW-based models become slightly more diverse w.r.t. trigrams.

Overall, it can again be seen that BoW-based models generate detectably more diverse responses (with greater precision), and remain to do so as response length increases. The decreasing diversity of discriminator-based generated responses further confirms that more invasive control during generation impedes textual variety.

4.4.3 The Effects of Prompt Style on PPLM Control

Recall that given a conditioning prompt prompt , a predefined style attribute a , and some controlled dialogue generation model parameterized by θ , generating a style-controlled piece of text \mathbf{x} entails modeling $p_\theta(\mathbf{x}|a, \text{prompt})$. It is therefore reasonable to expect the output distribution of controlled generation model p_θ to depend (to some extent) on the content and style of the conditioning text, prompt . Indeed, the content and style of prompts are found to strongly influence the output of neural text generation models [Fan et al., 2018, Lester et al., 2021]. Thus, studying the effects of a prompt’s age-style (i.e., whether a prompt is considered young, old, or neutral by BERT_{FT}) on the style and grammatical quality of PPLM-setups is of great importance, as it could inform developers of adaptive dialogue systems about mitigation of prompt-induced biases.

It is worth mentioning that the effects of prompt-style on PPLM-generation are not considered by Dathathri et al. [2020] or Madotto et al. [2020]. Studying these effects is an important extension of their methods, as not quantitatively taking into account the effects of prompt-style obfuscates the degree to which one can conclude whether attribute-relevance is the result of controlled generation or prompt-induced bias.

Figures 4.8 and 4.9 depict the average target probability and perplexity over responses generated by the baseline, best BoW-based model, and discriminator-based model, when prompted with a prompt of either young, neutral, or old style. More specifically, each bar represents a metric

(target probability or perplexity) averaged over $N = 270$ samples generated by a single model, when presented with five prompts of the same age-style. E.g., the blue bar in Figure 4.8a represents the average probability of samples generated by GPT-2 + frequency-based BoW to contain features learned to be young by BERT_{FT} , when the model was presented young-sounding prompts. The explicit numerical values of Figures 4.8 and 4.9 are found in Tables 4.4, 4.5, 4.6, and 4.7.

Ideally, the neutrally prompted baseline’s assigned probability of generating age-specific responses should be around 0.50. Furthermore, a prompt should shift the target probability in the direction of the prompt class, e.g., a young prompt should shift a young-model’s target prob upwards, and an old-model’s target prob downwards. We know from previous results in Tables 4.2 and 4.3 that the language model baselines are (in varying degrees) biased towards generating younger sounding responses to neutral prompts. Nevertheless, we should expect the impact of prompt-style to persist, albeit to differing degrees based on the (dis)similarity in style between the prompt and response, and the type of attribute model being used (BoW or discriminator).

Figures 4.8 and 4.9, show that the models’ target probabilities indeed move accordingly with the prompts’ styles. E.g., young-prompted young-model achieves highest young target prob, then neutral prompted, and then old prompted (Figures 4.8a and 4.8c). The same pattern holds the other way around: an old-prompted old-model has the highest (old) target probability, then neutrally prompted, and then the young- prompted ones (Figures 4.8b and 4.8d). In these last two sub-figures, we also clearly see that the discriminator-based old generation models achieve substantial target probability improvements over the baseline (and BoW-based models), for every style of prompt. By contrast, the young-generation models (Figures 4.8a and 4.8c) do not show the same pattern: discriminator-based models achieve similarly subtle improvements in target probability over their baselines as the BoW-based models do. Figure 4.8a even shows the discriminator-based models to perform worse than the baseline and BoW-based models.

So the class of the prompt strongly influences the style of the generated response.

Overall, Figures 4.8 and 4.9 and Tables 4.4, 4.5, 4.6, and 4.7 show that the style of the prompt clearly nudges the assigned probability of containing age-related features in the direction of the prompt’s style. Moreover, using strongly young-prompt results in heavily reinforced young-bias for GPT2 and DialoGPT. Similarly, using an older sounding prompt results in a slightly neutralized young-bias for both language models. Additionally, the discriminator-based old models (GPT-2 and DialoGPT) always yield the highest relative improvements over the baselines.

However, they fail to do so when attempting to generate younger sounding responses, which could suggest that the stylistic features learned to be young and old by BERT_{FT} lie at different syntactical levels, and are not equally challenging to control for. Finally, the aforementioned tables show that the effects of prompt-style are largely limited to the target probabilities. For perplexity and distinctiveness, the young- and old-prompted results show similar patterns to the neutral-prompt setting: BoW-based models achieve smaller increases in control, but maintain relatively desirable perplexity and diversity. Whereas, the discriminator-based models achieve higher levels of control, at the cost of worse perplexity and Dist-scores.

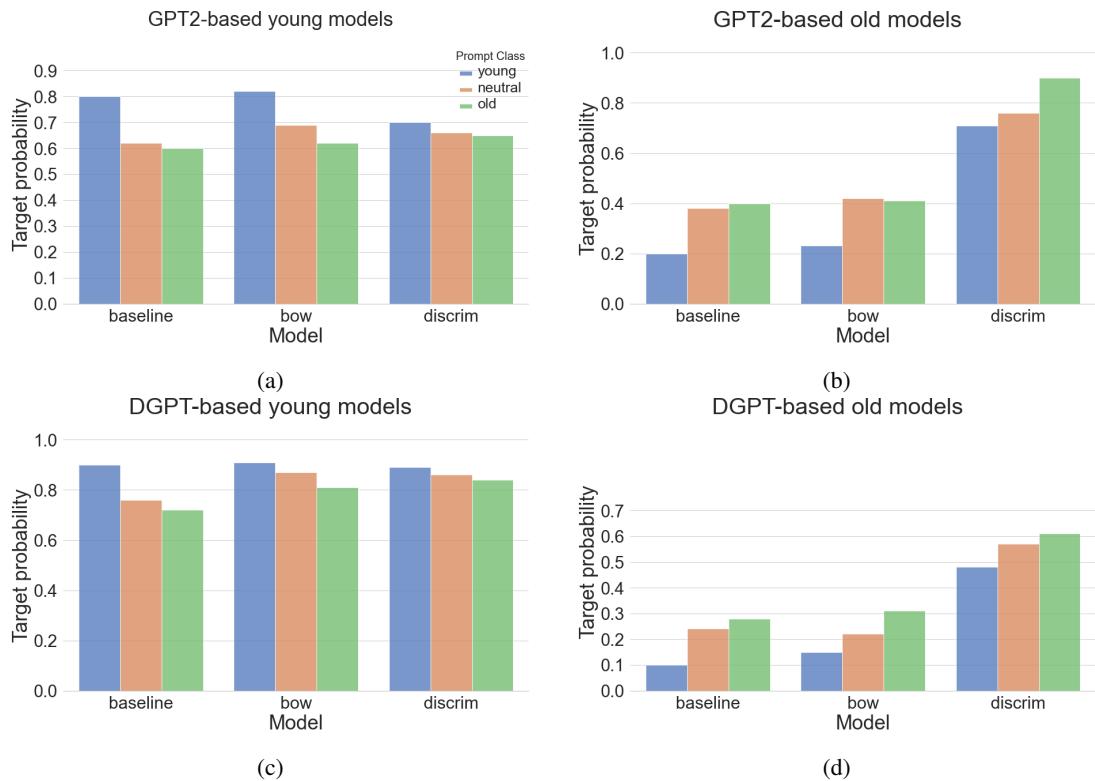


Figure 4.8: The average target probabilities assigned to responses generated by various PPLM-setups, conditioned on young, neutral, or old prompts. The plots are best viewed in color.

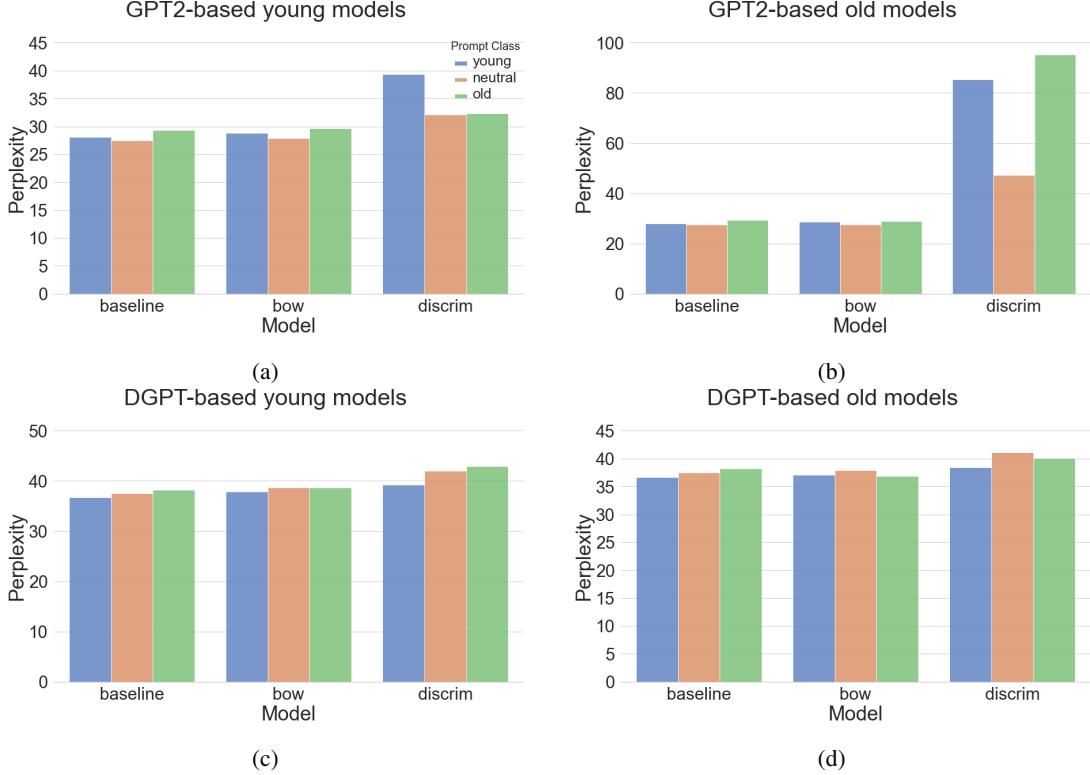


Figure 4.9: The average perplexities of responses generated by various PPLM-setups, conditioned on young, neutral, or old prompts. The plots are best viewed in color. Mind the differences in scale between the y -axes. The plots are best viewed in color.

Model	ppl. ↓ better	Dist-1 ↑ better	Dist-2 ↑ better	Dist-3 ↑ better	\bar{P}_Y ↑ better	Acc. ↑ better
G-baseline	28.05 (± 6.12)	0.85 (± 0.13)	0.91 (± 0.08)	0.88 (± 0.08)	0.80 (± 0.33)	-
G-100MCW	27.71 (± 6.20)	0.85 (± 0.12)	0.91 (± 0.09)	0.88 (± 0.09)	0.75 (± 0.37)	-
G-B _{FB,Y}	28.81 (± 7.09)	0.86 (± 0.12)	0.92 (± 0.08)	0.89 (± 0.08)	0.82 (± 0.32)	83.3%
G-B _{100MIU,Y}	28.49 (± 6.49)	0.86 (± 0.12)	0.91 (± 0.08)	0.88 (± 0.08)	0.83 (± 0.32)	83.0%
G-D _Y	39.32 (± 37.49)	0.84 (± 0.21)	0.61 (± 0.40)	0.57 (± 0.40)	0.70 (± 0.40)	70.7%
D-baseline	36.69 (± 9.11)	0.87 (± 0.10)	0.91 (± 0.06)	0.87 (± 0.08)	0.90 (± 0.24)	-
D-100MCW	36.93 (± 9.18)	0.86 (± 0.11)	0.91 (± 0.06)	0.88 (± 0.07)	0.90 (± 0.25)	-
D-B _{FB,Y}	37.35 (± 8.60)	0.88 (± 0.10)	0.91 (± 0.06)	0.87 (± 0.08)	0.90 (± 0.26)	90.0%
D-B _{100MIU,Y}	37.87 (± 8.32)	0.88 (± 0.10)	0.91 (± 0.07)	0.87 (± 0.09)	0.91 (± 0.24)	92.6%
D-D _Y	39.22 (± 14.96)	0.89 (± 0.12)	0.86 (± 0.19)	0.79 (± 0.23)	0.89 (± 0.25)	91.1%

Table 4.4: Results of age-controlled dialogue generation: **young**-targeted models, conditioned on **young prompts**. Format: *average metric (standard error)*. ppl. is perplexity w.r.t. GPT-1. Dist- n (for $n = 1, 2, 3$) is number of distinct n -grams normalized by text length, as a measure of diversity. \bar{P}_Y is the sample's average probability to contain features learned to be young by BERT_{FT}. Acc. is BERT_{FT}'s accuracy when classifying the row's samples.

Model	ppl. ↓ better	Dist-1 ↑ better	Dist-2 ↑ better	Dist-3 ↑ better	\bar{P}_O ↑ better	Acc. ↑ better
G-baseline	28.05 (± 6.12)	0.85 (± 0.13)	0.91 (± 0.08)	0.88 (± 0.08)	0.20 (± 0.33)	-
G-100MCW	27.71 (± 6.20)	0.85 (± 0.12)	0.91 (± 0.09)	0.88 (± 0.09)	0.25 (± 0.37)	-
G-B _{FB,O}	28.54 (± 6.45)	0.86 (± 0.12)	0.92 (± 0.08)	0.89 (± 0.08)	0.23 (± 0.36)	22.6%
G-B _{100MIU,O}	28.18 (± 5.70)	0.87 (± 0.11)	0.92 (± 0.08)	0.89 (± 0.09)	0.21 (± 0.34)	21.5%
G-D _O	85.40 (± 150.28)	0.67 (± 0.30)	0.62 (± 0.31)	0.62 (± 0.32)	0.71 (± 0.40)	70.5%
D-baseline	36.69 (± 9.11)	0.87 (± 0.10)	0.91 (± 0.06)	0.87 (± 0.08)	0.10 (± 0.24)	-
D-100MCW	36.93 (± 9.18)	0.86 (± 0.11)	0.91 (± 0.06)	0.88 (± 0.07)	0.10 (± 0.25)	-
D-B _{FB,O}	37.25 (± 9.45)	0.87 (± 0.11)	0.91 (± 0.06)	0.87 (± 0.08)	0.12 (± 0.29)	11.1%
D-B _{100MIU,O}	37.04 (± 8.78)	0.88 (± 0.10)	0.91 (± 0.05)	0.88 (± 0.07)	0.15 (± 0.32)	15.2%
D-D _O	38.46 (± 14.91)	0.82 (± 0.15)	0.87 (± 0.15)	0.83 (± 0.17)	0.48 (± 0.44)	47.4%

Table 4.5: Results of age-controlled dialogue generation: **old**-targeted models, conditioned on **young prompts**. Format: *average metric (standard error)*. ppl. is perplexity w.r.t. GPT-1. Dist- n (for $n = 1, 2, 3$) is number of distinct n -grams normalized by text length, as a measure of diversity. \bar{P}_O is the sample’s average probability to contain features learned to be young by BERT_{FT}. Acc. is BERT_{FT}’s accuracy when classifying the row’s samples.

Model	ppl. ↓ better	Dist-1 ↑ better	Dist-2 ↑ better	Dist-3 ↑ better	\bar{P}_Y ↑ better	Acc. ↑ better
G-baseline	29.34 (± 10.30)	0.86 (± 0.09)	0.94 (± 0.04)	0.90 (± 0.06)	0.60 (± 0.43)	-
G-100MCW	29.14 (± 10.11)	0.86 (± 0.10)	0.93 (± 0.04)	0.90 (± 0.06)	0.60 (± 0.44)	-
G-B _{Y,FB}	29.61 (± 10.28)	0.86 (± 0.10)	0.93 (± 0.04)	0.91 (± 0.06)	0.62 (± 0.43)	61.1%
G-B _{Y,100MIU}	29.51 (± 0.09)	0.87 (± 0.09)	0.93 (± 0.05)	0.90 (± 0.06)	0.68 (± 0.42)	68.5%
G-D _Y	32.34 (± 19.88)	0.77 (± 0.20)	0.84 (± 0.19)	0.80 (± 0.23)	0.65 (± 0.43)	65.4%
D-baseline	38.18 (± 12.03)	0.86 (± 0.12)	0.90 (± 0.08)	0.86 (± 0.09)	0.72 (± 0.38)	-
D-100MCW	37.73 (± 11.88)	0.85 (± 0.13)	0.90 (± 0.08)	0.86 (± 0.09)	0.73 (± 0.39)	-
D-B _{Y,FB}	38.24 (± 11.53)	0.86 (± 0.12)	0.90 (± 0.08)	0.86 (± 0.10)	0.81 (± 0.34)	82.6%
D-B _{Y,100MIU}	38.66 (± 11.57)	0.85 (± 0.12)	0.90 (± 0.07)	0.86 (± 0.09)	0.81 (± 0.33)	80.7%
D-D _Y	42.93 (± 20.18)	0.90 (± 0.14)	0.79 (± 0.22)	0.68 (± 0.28)	0.84 (± 0.30)	85.2%

Table 4.6: Results of age-controlled dialogue generation: **young**-targeted models, conditioned on **old prompts**. Format: *average metric (standard error)*. ppl. is perplexity w.r.t. GPT-1. Dist- n (for $n = 1, 2, 3$) is number of distinct n -grams normalized by text length, as a measure of diversity. \bar{P}_Y is the sample’s average probability to contain features learned to be young by BERT_{FT}. Acc. is BERT_{FT}’s accuracy when classifying the row’s samples.

4.4.4 BERT_{FT} Attention Patterns

We use visualizations of attention mechanisms [Vig, 2019] in BERT_{FT}’s attention heads to analyze recurring patterns when assigning target probabilities to generated prompt-response pairs. An attention weight can be interpreted as an indication of how important particular token is when producing the next representation of the current token [Bahdanau et al., 2015, Clark et al., 2019]. Please note that the PPLM-method does not change the attention weights of the underlying language models, so visualizing the attention weights of our generation models would not reveal

Model	ppl. ↓ better	Dist-1 ↑ better	Dist-2 ↑ better	Dist-3 ↑ better	\bar{P}_O ↑ better	Acc. ↑ better
G-baseline	29.34 (± 10.30)	0.86 (± 0.09)	0.94 (± 0.04)	0.90 (± 0.06)	0.40 (± 0.43)	-
G-100MCW	29.14 (± 10.11)	0.86 (± 0.10)	0.93 (± 0.04)	0.90 (± 0.06)	0.40 (± 0.44)	-
G-B _{O,FB}	28.81 (± 10.10)	0.86 (± 0.10)	0.93 (± 0.05)	0.90 (± 0.06)	0.41 (± 0.43)	41.1%
G-B _{O,100MIU}	29.05 (± 9.80)	0.86 (± 0.09)	0.93 (± 0.04)	0.90 (± 0.06)	0.40 (± 0.43)	39.6%
G-D _O	95.21 (± 174.42)	0.65 (± 0.27)	0.78 (± 0.18)	0.78 (± 0.18)	0.90 (± 0.25)	90.3%
D-baseline	38.18 (± 12.03)	0.86 (± 0.12)	0.90 (± 0.08)	0.86 (± 0.09)	0.28 (± 0.38)	-
D-100MCW	37.73 (± 11.88)	0.85 (± 0.13)	0.90 (± 0.08)	0.86 (± 0.09)	0.27 (± 0.39)	-
D-B _{O,FB}	37.80 (± 11.74)	0.86 (± 0.12)	0.90 (± 0.07)	0.87 (± 0.08)	0.28 (± 0.39)	29.3%
D-B _{O,100MIU}	36.93 (± 11.68)	0.87 (± 0.12)	0.90 (± 0.09)	0.86 (± 0.09)	0.31 (± 0.41)	29.6%
D-D _O	40.08 (± 16.77)	0.85 (± 0.14)	0.88 (± 0.10)	0.83 (± 0.14)	0.61 (± 0.42)	61.1%

Table 4.7: Results of age-controlled dialogue generation: **old**-targeted models, conditioned on **old prompts**. Format: *average metric (standard error)*. ppl. is perplexity w.r.t. GPT-1. Dist- n (for $n = 1, 2, 3$) is number of distinct n -grams normalized by text length, as a measure of diversity. \bar{P}_O is the sample’s average probability to contain features learned to be young by BERT_{FT}. Acc. is BERT_{FT}’s accuracy when classifying the row’s samples.

any age-related attention patterns. By contrast, BERT_{FT} is fine-tuned to detect age-related patterns, so visualizing its attention maps can inform us about which features are important when assigning a target-age probability to a generated sentence. Despite this seemingly natural interpretation, there is much debate about the validity of using attention mechanisms (as opposed to, e.g., saliency methods) as explanations for model output [Jain and Wallace, 2019, Wiegreffe and Pinter, 2019, Bastings and Filippova, 2020]. However, as Vig [2019] and Clark et al. [2019] suggest, attention weight visualization can be used tentatively as a complementary analysis tool to add to sets of different analysis methods to inform researchers about, e.g., possible linguistic patterns that may be attended to by attention-based models.

To provide a clear reading experience of the analyses presented below, we revise a few important concepts about BERT and how to read attention weight visualizations. During pre-processing, BERT tokenizes the input text and adds a special token [CLS] to the beginning of the text, and another special token [SEP] is appended to the text. If an input sequence consists of multiple sentences (e.g., a question-answer, or prompt-response input), [SEP] tokens are also used to separate the sentences. BERT_{FT} is a fine-tuned version of BERT-base-uncased [Devlin et al., 2019], which consists of 12 layers of 12 attention heads. We refer to a specific attention head as Head `layer_number-head_number`, where `layer_number` and `head_number` range from 0 to 11. E.g., Head 3-1 refers to Head 1 in Layer 3. Furthermore, attention weights are visualized as colored lines between tokens of the input sequence, where a thicker line corresponds to a greater attention weight, and the color represents the layer in which the head is present. The input text is

displayed twice in parallel columns, to make visualizations of self-attention possible (visualized as lines between identical tokens in the same positions). Because BERT is designed to be deeply bidirectional, tokens can also attend to tokens in previous positions in the input text.

We show attention visualizations of BERT_{FT} when processing prompt-response pairs cherry-picked from the age-targeted prompted results, presented in Tables 4.4 and 4.7. The generated sequences are chosen to display more pronounced examples of recurring attention patterns. We show two young-targeted generated responses to younger sounding prompts in Figures 4.10 and 4.11, and two old-targeted generated responses to older sounding prompts in Figures 4.12 and 4.13. All prompts and responses received target probabilities from BERT_{FT} of at least 95%.

Heads in the same layer have the tendency to attend to similar patterns and linguistic phenomena Clark et al. [2019]. Our analyses seem to confirm this behavior, as recurring patterns are observed for heads in the same layers. For instance, we see (in Figures 4.10a, 4.11a, 4.12a, and 4.13a) that heads in the last layer tend to broadly disperse attention among all tokens. Recurring patterns are also observed in earlier layers, such as Head 2-9 attending to the next token in the sequence (Figures 4.10b, 4.11b, 4.12b, and 4.13b), and Head 4-4 attending to the special tokens (Figures 4.10c, 4.11c, 4.12c, and 4.13c). Certain heads also seem to pay special attention to age-related linguistic features, specifically certain tokens associated with an age group (mentioned in Section 4.4.5). This is most noticeable in Head 9-0 (Figures 4.10d, 4.11d, 4.12d, and 4.13d), which seems to consistently devote the majority of its attention to tokens that are found to be indicative of age. For instance, Head 9-0 attends strongly to younger sounding tokens like *facebook*, *awesome*, *cool*, and slang and swear words in Figures 4.10d and 4.11d. And the same attention head then focuses strongly on tokens associated with older age in Figures 4.12d and 4.13d: e.g., *workers*, *union*, *greetings*, or *fellow*.

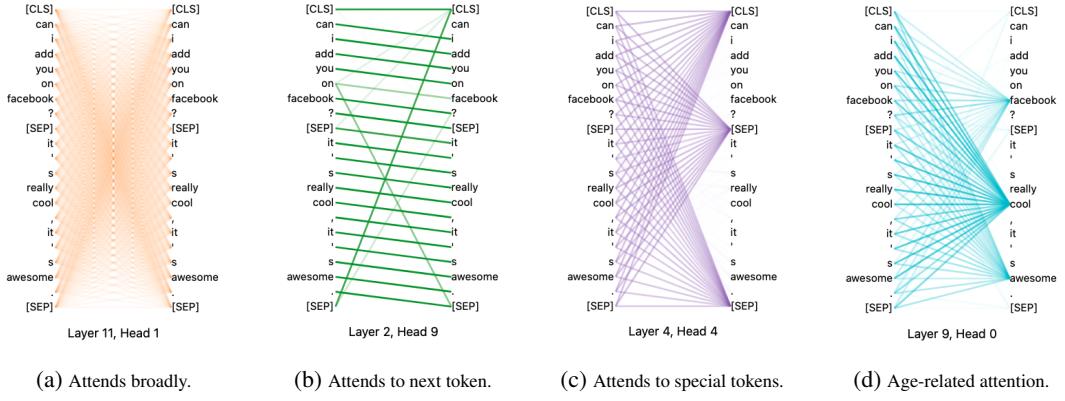


Figure 4.10: Attention weights visualizations of four of BERT_{FT}'s attentions heads and the patterns to which they presumably attend when processing representations for a cherry-picked prompt-response pair generated by **young**-targeted GPT2-Discrim.

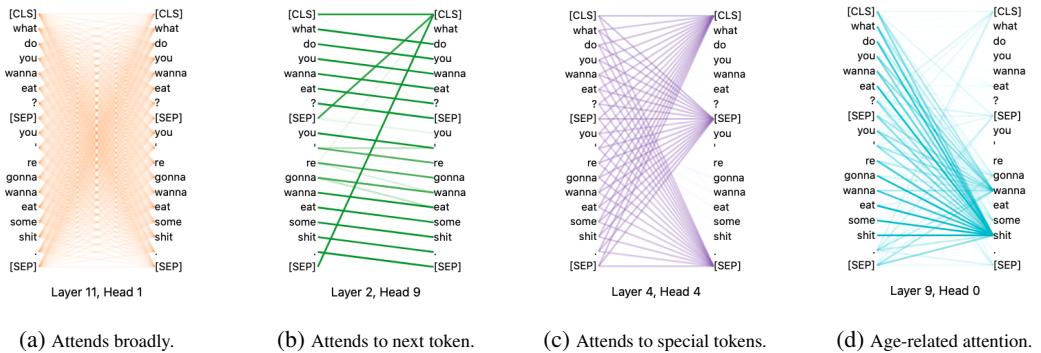


Figure 4.11: Attention weights visualizations of four of BERT_{FT}'s attentions heads and the patterns to which they presumably attend when processing representations for a cherry-picked prompt-response pair generated by **young**-targeted GPT2-BoW_{100MIU}. [L: **@Sandro, Raquel, or Janie - Is it OK to have a swear word in this figure? Unfortunately, this is one of the prompt-response pairs that most clearly displays the pattern in subfigure (d).]**

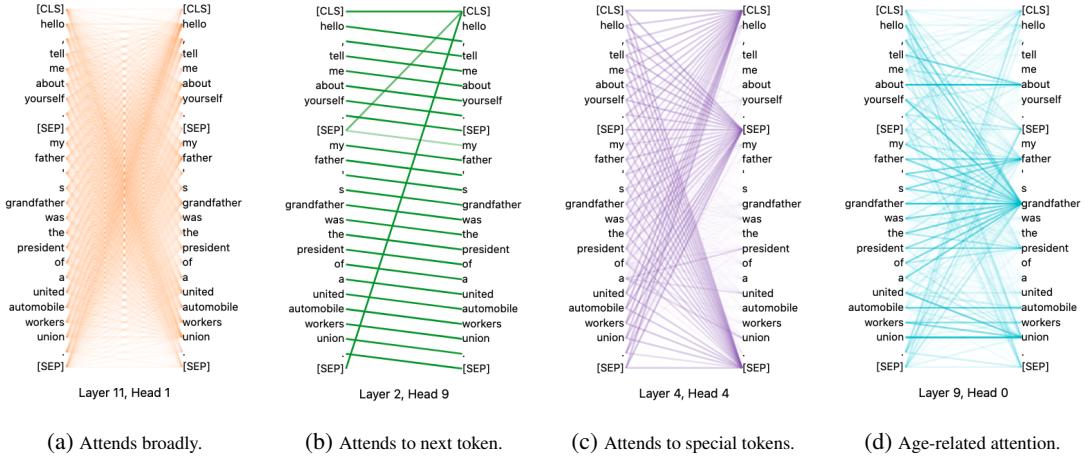


Figure 4.12: Attention weights visualizations of four of BERT_{FT}'s attentions heads and the patterns to which they presumably attend when processing representations for a cherry-picked prompt-response pair generated by **old**-targeted GPT2-Discrim.

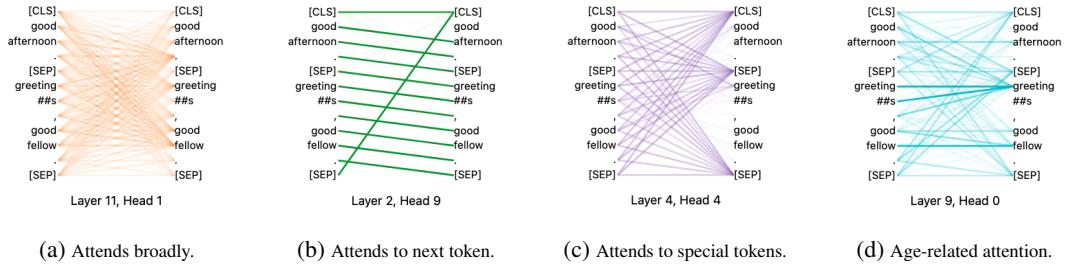


Figure 4.13: Attention weights visualizations of four of BERT_{FT}'s attentions heads and the patterns to which they presumably attend when processing representations for a cherry-picked prompt-response pair generated by **old**-targeted DGPT-Discrim.

4.4.5 Qualitative Analyses

The generated responses to neutral prompts are split by (1) whether or not BERT_{FT} correctly classified a response as its target class, and (2) the level of perplexity: low (ppl. < 27.52), medium ($27.52 \leq \text{ppl.} < 35.63$), or high ($\text{ppl.} \geq 35.63$). See Section 4.4.1 for an explanation of the rationale behind these intervals. Table 4.9 shows how the generated responses are distributed among these cases for the best performing BoW-based models, and discriminator-based models for both underlying language models and target age groups. As can be seen, the majority of the responses generated by GPT-2 BoW-based models lie in the low perplexity range, whereas the discriminator-based GPT-2 models also have percentage peaks in the higher perplexity registers (e.g., 37.2% high-perplexity correctly classified for G-Discrim_{Old}). By contrast, the DialoGPT-

based models all show the majority of their distributions lying in the medium-to-high perplexity range.

To narrow down the comparison, the qualitative inspection of samples is limited to responses generated by the discriminator-based models of both language models and age groups (i.e., samples generated by $G\text{-Discrim}_{Young}$, $G\text{-Discrim}_{Old}$, $D\text{-Discrim}_{Young}$, and $D\text{-Discrim}_{Old}$). Moreover, these models are all among those with the highest target probability improvements over their respective baselines in Tables 4.2 and 4.3, so the differences in language style should be most pronounced between these setups. Furthermore, to emphasize the differences between perplexity, we only consider low versus high perplexity samples. Finally, the samples used for qualitative inspection are also split by whether or not they were correctly classified by BERT_{FT}. To summarize, qualitative inspection are performed on a total of 16 splits: by model ($G\text{-Discrim}_{Young}$, $G\text{-Discrim}_{Old}$, $D\text{-Discrim}_{Young}$, or $D\text{-Discrim}_{Old}$), perplexity (low or high), and classification outcome (correct or incorrect). Table 4.8 shows examples of generated responses containing the patterns and observations discussed in the remainder of this section.

PROMPT	Good weather we're having.
GPT-2_{Young}	It's awesome weather.
GPT-2_{Old}	It's not been quite a good weather for long.
DialoGPT_{Young}	Gonna make the weather better.
DialoGPT_{Old}	It's not the best weather in the country.
PROMPT	Hi, how's it going?
GPT-2_{Young}	It's been almost 2 weeks now since my boyfriend has been gone.
GPT-2_{Old}	I've had my first surgery recently.
DialoGPT_{Old}	You have you're first born daughter going on with the surgery????
PROMPT	Can we talk?
GPT-2_{Young}	We have to stop talking about this is all going so stupid lol lol, sh*t just got weird lol LOL.
GPT-2_{Old}	The question was raised in a response sent by an official.
DialoGPT_{Young}	You guys are doing a lot.
DialoGPT_{Old}	We have talked about your past your previous history.
PROMPT	Good weather we're having.
GPT-2_{Young}	The Weather is nice, the temperature cool, the traffic good.
GPT-2_{Old}	The sun was setting when we were getting up with a huge rain.
DialoGPT_{Young}	Cold weather we can't handle it
PROMPT	Hello, tell me about your latest holiday.
GPT-2_{Young}	I'm trying hard to stay away.
GPT-2_{Old}	It's not my favorite holiday.
DialoGPT_{Old}	It was a great weekend!
PROMPT	Can we talk?
GPT-2_{Young}	Can we talk with some sense?
GPT-2_{Old}	No, not about your favourite club.
DialoGPT_{Young}	We have a meeting in a bit.
DialoGPT_{Old}	What I can't.
DialoGPT_{Old}	We have a couple more people interested!

(b) Low perplexity & incorrectly classified.

(a) Low perplexity & correctly classified.

Table 4.8: Cherry-picked examples of low-perplexity young-targeted (red rows) and old-targeted (blue rows) responses to neutral prompts (yellow rows) generated by various PPLM-setups, and grouped by whether or not they were correctly classified by BERT_{FT}. High-perplexity samples are omitted due to excessive amounts of gibberish and nonsensical strings present in the samples.

Manual inspection of the sub-samples show that correctly classified low-perplexity samples from models targeted towards the older age group use more formal words than their young-targeted counterparts, e.g., *quite*, *significant*, *powerful*, or *institutions*. Samples generated by these models also show recurring topics that are typically associated with older age, e.g., children (*my son* or *your daughter*), history, or politics. Responses about healthcare related subjects are also more common among the samples generated by these models, as indicated by the use of words

like *surgery*. When viewing high-perplexity samples from the same set of models (i.e., those target towards the older age group and correctly classified), we see a substantial increase in the amount of gibberish and nonsensical sequences (space-less sequences of words, repetitions of the same words, or sequences of punctuation marks). When we inspect the incorrectly classified samples from models targeted towards the older class, we also see a considerable increase in the amount of nonsensical strings. Remarkably, we also observe more linguistic patterns associated with younger age in this sub-sample, e.g., words of excitement like *favourite* or *best*, informal vocabulary (*pretty much*).

The low-perplexity correctly classified samples generated by models targeted towards the younger age group also contain clear indications of their target age. We observe more usage of slang words, swear words, and informal language, like *yeah*, *dude*, *cool*, *kinda*, and *lol*. Furthermore, these sub-samples are also characterized by increased use of words of excitement and exclamation: e.g., *awesome*, *really*, *love*, *fun*, or *amazing*. Especially in the correctly classified low-perplexity responses generated by G-Discrim_{Young}, we see a strong presence of topics such as dating (indicated by words such as *girlfriend* or *boyfriend*), having parents (*my dad*), parties, and student life (*roommate*). An interesting pattern observed in this sub-sample is one often associated with millennial or social networking language: the tendency to end a (serious-sounding) statement with *lol* or *haha* as a means of softening the perceived severity of the statement or as a signal of interlocutor involvement [Newitz, 2019, Tagliamonte and Denis, 2008]. For example, *We have to stop talking about this all going so stupid lol lol*. When inspecting the high-perplexity and/or incorrectly classified samples generated by the younger-targeted models, we observe similar patterns. Namely, a substantial increase in non-alphabetical strings, and gibberish.

When comparing the samples from GPT-2 and DialoGPT-based models, we see that the responses generated by DialoGPT are more dialogic, whereas GPT-2 sometimes generates sequences that look more like sentence completions. This is to be expected from the differences in pre-training methods between the two language models [Zhang et al., 2020]. Furthermore, we see a lot more nonsensical low-perplexity responses generated by DialoGPT-based models. Perplexity remains a rough proxy for fluency, and observations like these confirm this problem. That is, low perplexity often does not imply lack of gibberish or nonsense. DialoGPT’s strong bias towards generating younger-sounding language is also noticeable in its generated samples. For example, DialoGPT-based models targeted towards the older group still produce lots of word of excitement.

However, this older-targeted DialoGPT-based model does succeed in producing significantly fewer usages of slang or swear words, when compared to its younger-target counterpart.

model	low ppl. ✓	low ppl. ✗	med ppl. ✓	med ppl. ✗	high ppl. ✓	high ppl. ✗
G-BoW _{Young}	48.0%	15.6%	15.6%	10.4%	6.7%	3.7%
G-BoW _{Old}	25.6%	37.7%	11.9%	14.8%	5.6%	4.4%
G-Discrim _{Young}	33.3%	15.9%	17.4%	8.5%	17.0%	7.9%
G-Discrim _{Old}	23.8%	18.2%	13.4%	3.3%	37.2%	4.1%
D-BoW _{Young}	5.2%	0.4%	40.0%	5.2%	43.3%	5.9%
D-BoW _{Old}	3.0%	5.9%	10.4%	35.2%	8.5%	37.0%
D-Discrim _{Young}	5.9%	2.7%	23.0%	5.6%	56.9%	5.9%
D-Discrim _{Old}	7.0%	4.1%	19.3%	11.5%	30.3%	27.8%

Table 4.9: Distribution of dialogue responses generated by various PPLM-setups (under the **model** column) into categories of low, medium, or high perplexity, and whether or not they are correctly classified as their target age group by BERT_{FT} (✓: correctly classified, ✗: incorrectly classified). Alternating row colors are for readability and bear no meaning.

Chapter 5

Discussion

5.1 Summary of Key Findings and Interpretations

In this thesis, the problems of automated age detection from dialogue and age-adaptive controlled dialogue generation are investigated. We first studied the extent to which purely text-based NLP models can detect age-related linguistic features in dialogue data, and which features drive their predictions. Subsequently, we studied the extent to which age-adaptive dialogue generation is possible using Plug-and-Play language models (PPLM) [Dathathri et al., 2020].

The results of the age detection experiments indicated that our fine-tuned version of BERT, BERT_{FT}, is capable of detecting age-related linguistic features in dialogue utterances with reasonable accuracy. BERT_{FT} was particularly useful for age detection when the dialogue fragment is long enough to contain discriminative signal. At the same time, we found that much simpler models based on n -grams achieve comparable performance, which suggests that, in dialogue, ‘local’ features can be indicative of the language of speakers from different age groups. We showed this to be the case, with both lexical and stylistic cues being informative to these models in this task.

Furthermore, the age detection results informed us about the development of controllable dialogue generation systems using PPLM. Namely, the presumed locality of age-indicative features, as suggested by the comparable performance achieved by the n -gram models, motivates the use of unigram-based bag-of-words attribute models in PPLM setups for, in comparison to more sophisticated discriminator-based PPLM setups. Additionally, we used as an empirically generated attribute model for a BoW-based PPLM setup, the list of unigrams deemed most

informative by the best-performing unigram classifier. Overall, the analyses of experiment 1 predominantly compare the classification performances and feature importance of n -gram and Transformer-based neural models. This comparison is continued in the controlled dialogue generation experiments, in the form of comparing BoWs versus neural discriminators as attribute models in PPLM setups. Also, the best performing classifier from the age detection experiments, BERT_{FT} , is used in the second phase of experiments to evaluate the prompts and generated responses for their resemblance to the linguistic style of the younger or older age group.

The controlled dialogue generation results indicate that it is possible to use PPLM to generate dialogue responses that have been detectably adapted to the styles of different age groups. We observed that discriminator-based PPLM setups typically achieve higher levels of detectable control (i.e., statistical resemblance to a specific writing style) than BoW-based setups, but generate significantly more perplexing and repetitive responses. This could be attributable to the fact that BoW-based control is more local (i.e., token-level) and less invasive than discriminator-based control, which can operate at the structural level. The results also indicate that the underlying language models used in our PPLM setups, GPT-2 and DialoGPT, are both biased towards generating younger sounding language. This is most likely due to both models having been pre-trained on texts scraped from web pages, which are typically over-represented by millennials [Radford et al., 2019, Zhang et al., 2020]. Quantitative analyses of the dialogue generation results revealed that the outputs of some PPLM setups display a correlation between improved control and worsened perplexity and diversity. Also, the style of the prompt (i.e., whether it's classified by BERT_{FT} as young, neutral, or old) steers the style of the generated response in the direction of the prompt's style. This observation indicates the importance of taking prompt-induced biases into account when developing controlled generation models.

We also find that BoW-based models are significantly worse at controlling for older sounding language than they are for younger sounding language. It could be that there are some detectable, very young-sounding, tokens, that make the token-level perturbations made by BoW-based models more likely to contain detectably young-features, whereas old-sounding features are more salient at the structural syntactical level and thereby harder to control for. BERT_{FT} 's attention weight visualizations could also be evidence of this phenomena, because the attention weights of linguistic-feature attending heads are much more focused for strongly young-input, than the more dispersed attention weights of old sentences. Finally, qualitative inspection of

generated responses reveal that differences in style are noticeable by use of age-indicative words, formality of language, and the prevalence of certain topics.

5.2 Contributions and Implications

The age detection experiments build on the work of Schler et al. [2006], who focus on age detection in written discourse using handcrafted features. We extend their work by: (1) eliminating the need for handcrafted features by learning end-to-end representations using state-of-the-art NLP models; (2) applying this approach to dialogue data, using a dataset of transcribed spontaneous open-domain dialogues; (3) showing that text-based models can indeed detect age-related differences, even in the case of very sparse signals at the level of dialogue utterances; (4) carrying out an in-depth analysis of the models’ predictions to gain insight about which elements of language use are most informative. Furthermore, the age detection analyses motivate the use of local features (i.e., BoW-based attribute models) for controlled generation as a viable alternative to neural discriminator-based attribute models. Our work on age detection from dialogue can be considered a preliminary step to the modeling of age-related linguistic adaptation by AI conversational systems. In particular, these results informed our work on controlled dialogue generation using PPLM.

Our work on controlled dialogue generation using PPLM builds on previous work on controlled language generation using PPLM by Dathathri et al. [2020], who focus on controlled story writing for concrete styles (e.g., sentiment, or topic). We extend their work in several important ways: (1) we control language generation for more abstract writing styles, i.e., age group specific linguistic style; (2) we use PPLM for dialogue response generation; (3) we propose methods for empirical development of BoW attribute models (as opposed to the manually curated BoWs used by Dathathri et al. [2020]) and demonstrate their applicability for controlled dialogue generation; (4) we thoroughly study the relationships between dialogue response quality, response style, and response length; and finally (5) we carry out an extensive analysis about the effects of prompt-induced biases on the quality and style-attribute adherence of generated language, which has been overlooked by previous work about Plug-and-Play generation. Despite previous work by Madotto et al. [2020] also focusing on Plug-and-Play conversational models, our work demonstrates the achievability of a Plug-and-Play approach to controlled dialogue generation, without the necessity to generate attribute-specific dialogue datasets, or separately optimize residual adapter modules.

Overall, this research is a promising step towards the development of adaptive conversational systems. In particular, the development of age-adaptive conversational systems can benefit from these results. Since consistent language style differences were found between age groups, systems whose language generation capabilities aim to be consistent with a given age group should therefore reproduce these patterns. This could be achieved, for example, by embedding Plug-and-Play modules that control the generation of a system’s output, which could lead to better, more natural interactions between human speakers and a conversational system.

5.3 Limitations

There are several limitations to take into account when interpreting the results presented in this thesis. The generalizability of the age detection results are limited by coarse granularity with which the age groups are defined. One could imagine that if we aim to leverage the predicted age signal in real-world conversational systems, a fine-grained grouping is more reasonable. Furthermore, age group information remains difficult to identify in short utterances. A non-negligible proportion of utterances seem too short for reliable classification, and it remains to be shown investigated if NLP models or even humans could recognize the targets. However, for future applications, the combination of speech analysis could improve performance. Acoustic signals are not considered in this work, but they are important characteristics of speaking style, and should also be taken into account when developing audio-based conversational systems.

Due to lack of information about the *none* and *no info* topics, the age detection results cannot rule out confounding factors. This is especially important for the 50-plus age group, where the *none* and *no info* topics apply to nearly 50% of the utterances. Nevertheless, our analysis of classification performance per topic do not seem to indicate severe differences in the identifiability of age-related signals between topic categories.

The reliability of the controlled dialogue generation results is impacted by perplexity (measured by GPT-1) being a crude proxy for fluency. Perplexity also lacks generalizability as an evaluation measure, because it only measures the uncertainty assigned by one language model. Automated evaluation of fluency could be improved by considering an aggregated perplexity measure (i.e., perplexity averaged over an ensemble of language models). Alternatively, indications of a generated response’s fluency could be made more reliable by consulting human opinion.

It is also important to realize that the learned representations for younger and older style used in this research are specific to the BNC [Love et al., 2017], and should not be interpreted as

general representations of the speaking style of people of ages 19-29 and 50-plus. Even within the context of the BNC, the representations used for classification and generation are indicative of textual features learned to coincide with utterances from certain age groups, and are therefore approximations of the age groups’ speaking styles.

Finally, due to the lack of age-labeled transcribed spontaneous dialogue data, we were constrained to using the same dialogue dataset to develop the attribute models (i.e., the BoWs or neural discriminators used in PPLM setups), and to train evaluation model (i.e., BERT_{FT}). It must be noted that this is still sound practice, because the attribute and evaluation models are not trained jointly, and do not share representations. Moreover, the approach of using an external classifier for evaluation is based on the work of Dathathri et al. [2020] and Madotto et al. [2020]. However, the assigned target probabilities indicating attribute adherence would be more generalizable if the evaluation model were trained on a different dataset of spontaneous dialogue utterances between speakers of the same age.

5.4 Future research directions

While we performed the classification task at the level of single dialogue utterances, future work may take into account larger dialogue fragments, such as the entire dialogue or a fixed number of turns. This would make the setup more comparable to discourse, but would require making experimental choices and dealing with extra computational challenges. Moreover, it could be tested whether the language used by a speaker is equally discriminative when talking to a same-age (this work) or a different-age interlocutor. Also, future work might involve leveraging acoustic signals for automated age-detection and controlled generation. Acoustic signals can contain important features indicative of speaking style, and could therefore substantially improve classification and controlled generation performance.

Developing interactive age-adaptive conversational systems was beyond the scope of this thesis. Therefore, an interesting future research objective would be to use PPCM for age-adaptive dialogue generation. This would require also developing and training residual adapter modules, and generating style-specific dialogue datasets. Our age detection results suggest the importance of n -gram features for age-identification. Future work could therefore entail adapting the PPLM and PPCM to work with n -gram lists (for $n > 1$) as attribute models. This would require finding a way to bypass the need to re-train large underlying language models like GPT-2 and DialoGPT for arbitrary n -grams. Finally, the current set of tools to analyze which features drive controlled

generation in PPLM-setups is limited. It would be valuable to develop probing methods for PPLM-setups. Namely, visualizations or saliency methods targeted towards understanding the activation space perturbations made by PPLM could provide important insights about the effects of attribute models on controlled generated output.

5.5 Ethical and environmental considerations

There is growing concern around the ethics, environmental costs, and societal dangers associated with powerful large-scale language models [Brown et al., 2020, Bender et al., 2021]. There is a non-negligible societal risk accompanying the recent trend of deploying ever larger language models that are more capable of producing texts that are indistinguishable from human-written ones. Namely, such language models can be used to spread misinformation (e.g., fake news), or to perpetuate harmful social biases. Plug-and-Play language models carry the same risks, if not to a greater extent, as they can be used to produce personalized, and ultimately more convincing text. However, Dathathri et al. [2020] also demonstrates the benevolent applicability of PPLM for language detoxification. Additionally, the same methods that generate convincing texts can be used to build more sophisticated tools for detecting artificially generated texts Gehrmann et al. [2019]. Moreover, I believe that the positive applications of PPLM for, e.g, personalized dialogue generation, outweigh the potential risks of misuse.

Chapter 6

Conclusion

In this thesis, I first investigate the extent to which text-based NLP-models can detect age-related linguistic features in dialogue data, and which features drive their predictions. Then, I study the extent to which age-adaptive dialogue generation is possible using Plug-and-Play language models (PPLM) [Dathathri et al., 2020].

The results of the age detection experiments show that a fine-tuned BERT model is capable of detecting age-related linguistic features in dialogue utterances with reasonable accuracy, especially when the dialogue fragment is long enough to contain discriminative signal. However, simpler n -gram-based models achieve comparable performance, suggesting that, in dialogue, ‘local’ features can be indicative of the language of speakers from different age groups. This is shown to be the case, with both lexical and stylistic cues being informative to these models in this task. Furthermore, the age detection results informed the subsequent experiments about controlled dialogue generation using PPLM.

The controlled dialogue generation results show that it is possible to use PPLM to generate dialogue responses that possess detectable linguistic features associated with specific age groups. Discriminator-based PPLM-setups typically achieve higher levels of detectable control than bag-of-words (BoW) based setups, but generate significantly more perplexing and repetitive responses. This could be attributable to the fact that BoW-based control is more local (i.e., at the token-level) and less invasive than discriminator-based control, which can operate at the structural level.

Overall, I believe this thesis is a promising step towards the development of age-adaptive conversational systems.

Bibliography

- Central Limit Theorem*, pages 66–68. Springer New York, New York, NY, 2008. ISBN 978-0-387-32833-1. doi: 10.1007/978-0-387-32833-1_50. URL https://doi.org/10.1007/978-0-387-32833-1_50.
- E. E. Abdallah, J. R. Alzghoul, and M. Alzghoul. Age and gender prediction in open domain text. *Procedia Computer Science*, 170:563–570, 2020.
- L. J. Ba, J. R. Kiros, and G. E. Hinton. Layer normalization. *CoRR*, abs/1607.06450, 2016. URL <http://arxiv.org/abs/1607.06450>.
- D. Bahdanau, K. Cho, and Y. Bengio. Neural machine translation by jointly learning to align and translate. In Y. Bengio and Y. LeCun, editors, *3rd International Conference on Learning Representations, ICLR 2015, San Diego, CA, USA, May 7-9, 2015, Conference Track Proceedings*, 2015. URL <http://arxiv.org/abs/1409.0473>.
- A. Bapna and O. Firat. Simple, scalable adaptation for neural machine translation. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 1538–1548, Hong Kong, China, Nov. 2019. Association for Computational Linguistics. doi: 10.18653/v1/D19-1165. URL <https://www.aclweb.org/anthology/D19-1165>.
- J. Bastings and K. Filippova. The elephant in the interpretability room: Why use attention as explanation when we have saliency methods? In *Proceedings of the Third BlackboxNLP Workshop on Analyzing and Interpreting Neural Networks for NLP*, pages 149–155, Online, Nov. 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.blackboxnlp-1.14. URL <https://aclanthology.org/2020.blackboxnlp-1.14>.
- E. M. Bender, T. Gebru, A. McMillan-Major, and S. Shmitchell. On the dangers of stochastic parrots: Can language models be too big? In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, pages 610–623, 2021.
- C. M. Bishop. *Pattern recognition and machine learning*. Springer, 2006.
- D. Bohm and L. Nichol. *On dialogue*. Routledge, 2013.
- T. Brown, B. Mann, N. Ryder, M. Subbiah, J. D. Kaplan, P. Dhariwal, A. Neelakantan, P. Shyam, G. Sastry, A. Askell, S. Agarwal, A. Herbert-Voss, G. Krueger, T. Henighan, R. Child, A. Ramesh, D. Ziegler, J. Wu, C. Winter, C. Hesse, M. Chen, E. Sigler, M. Litwin, S. Gray, B. Chess, J. Clark, C. Berner, S. McCandlish, A. Radford, I. Sutskever, and D. Amodei. Language models are few-shot learners. In H. Larochelle, M. Ranzato, R. Hadsell, M. F. Balcan, and H. Lin, editors, *Advances in Neural Information Processing Systems*, volume 33, pages 1877–1901. Curran Associates, Inc., 2020. URL <https://proceedings.neurips.cc/paper/2020/file/1457c0d6bfcb4967418bfb8ac142f64a-Paper.pdf>.
- K. Clark, U. Khandelwal, O. Levy, and C. D. Manning. What does BERT look at? an analysis of BERT’s attention. In *Proceedings of the 2019 ACL Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, pages 276–286, Florence, Italy, Aug. 2019. Association for Computational Linguistics. doi: 10.18653/v1/W19-4828. URL <https://aclanthology.org/W19-4828>.

- N. Dai, J. Liang, X. Qiu, and X. Huang. Style transformer: Unpaired text style transfer without disentangled latent representation. *arXiv preprint arXiv:1905.05621*, 2019.
- S. Dathathri, A. Madotto, J. Lan, J. Hung, E. Frank, P. Molino, J. Yosinski, and R. Liu. Plug and play language models: A simple approach to controlled text generation. In *International Conference on Learning Representations*, 2020. URL <https://openreview.net/forum?id=H1edEyBKDS>.
- J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186, Minneapolis, Minnesota, June 2019. Association for Computational Linguistics. doi: 10.18653/v1/N19-1423. URL <https://aclanthology.org/N19-1423>.
- J. Edlund, J. Gustafson, M. Heldner, and A. Hjalmarsson. Towards human-like spoken dialogue systems. *Speech communication*, 50(8-9):630–645, 2008.
- A. Fan, M. Lewis, and Y. Dauphin. Hierarchical neural story generation. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 889–898, Melbourne, Australia, July 2018. Association for Computational Linguistics. doi: 10.18653/v1/P18-1082. URL <https://aclanthology.org/P18-1082>.
- C. Gallois and H. Giles. Communication accommodation theory. *The international encyclopedia of language and social interaction*, pages 1–18, 2015.
- S. Gehrmann, H. Strobelt, and A. Rush. GLTR: Statistical detection and visualization of generated text. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics: System Demonstrations*, pages 111–116, Florence, Italy, July 2019. Association for Computational Linguistics. doi: 10.18653/v1/P19-3019. URL <https://aclanthology.org/P19-3019>.
- K. He, X. Zhang, S. Ren, and J. Sun. Deep residual learning for image recognition. In *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 770–778, 2016. doi: 10.1109/CVPR.2016.90.
- S. Hochreiter and J. Schmidhuber. Long short-term memory. *Neural computation*, 9(8):1735–1780, 1997.
- S. Jain and B. C. Wallace. Attention is not Explanation. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 3543–3556, Minneapolis, Minnesota, June 2019. Association for Computational Linguistics. doi: 10.18653/v1/N19-1357. URL <https://aclanthology.org/N19-1357>.
- N. S. Keskar, B. McCann, L. Varshney, C. Xiong, and R. Socher. CTRL - A Conditional Transformer Language Model for Controllable Generation. *arXiv preprint arXiv:1909.05858*, 2019.
- D. P. Kingma and J. Ba. Adam: A method for stochastic optimization. In *ICLR (Poster)*, 2015. URL <http://arxiv.org/abs/1412.6980>.
- D. P. Kingma and M. Welling. Auto-Encoding Variational Bayes. In *2nd International Conference on Learning Representations, ICLR 2014, Banff, AB, Canada, April 14-16, 2014, Conference Track Proceedings*, 2014.
- S. Kottur, X. Wang, and V. Carvalho. Exploring personalized neural conversational models. In *Proceedings of the Twenty-Sixth International Joint Conference on Artificial Intelligence, IJCAI-17*, pages 3728–3734, 2017. doi: 10.24963/ijcai.2017/521. URL <https://doi.org/10.24963/ijcai.2017/521>.
- P. Kushneryk, Y. P. Kondratenko, and I. V. Sidenko. Intelligent dialogue system based on deep learning technology. In *ICTERI PhD Symposium*, 2019.

- G. Lample, S. Subramanian, E. Smith, L. Denoyer, M. Ranzato, and Y.-L. Boureau. Multiple-attribute text rewriting. In *International Conference on Learning Representations*, 2019. URL <https://openreview.net/forum?id=H1g2NhC5KQ>.
- B. Lester, R. Al-Rfou, and N. Constant. The power of scale for parameter-efficient prompt tuning, 2021.
- C. Li, X. Gao, Y. Li, B. Peng, X. Li, Y. Zhang, and J. Gao. Optimus: Organizing sentences via pre-trained modeling of a latent space. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 4678–4699, Online, Nov. 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.emnlp-main.378. URL <https://www.aclweb.org/anthology/2020.emnlp-main.378>.
- J. Li, M. Galley, C. Brockett, G. Spithourakis, J. Gao, and B. Dolan. A persona-based neural conversation model. In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 994–1003, Berlin, Germany, Aug. 2016. Association for Computational Linguistics. doi: 10.18653/v1/P16-1094. URL <https://aclanthology.org/P16-1094>.
- M. Li, K. J. Han, and S. Narayanan. Automatic speaker age and gender recognition using acoustic and prosodic level information fusion. *Computer Speech & Language*, 27(1):151–167, 2013.
- D. C. Liu and J. Nocedal. On the limited memory bfgs method for large scale optimization. *Mathematical programming*, 45(1):503–528, 1989.
- R. Love, C. Dembry, A. Hardie, V. Brezina, and T. McEnery. The spoken bnc2014: designing and building a spoken corpus of everyday conversations. In *International Journal of Corpus Linguistics*, 22(3):319–344, 2017.
- A. Madotto, E. Ishii, Z. Lin, S. Dathathri, and P. Fung. Plug-and-play conversational models. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 2422–2433, Online, Nov. 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.findings-emnlp.219. URL <https://www.aclweb.org/anthology/2020.findings-emnlp.219>.
- M. McTear. Conversational ai: Dialogue systems, conversational agents, and chatbots. *Synthesis Lectures on Human Language Technologies*, 13(3):1–251, 2020.
- A. Newitz. You've been monetized lol. *New Scientist*, 243(3240):24, 2019.
- A. Nguyen, J. Yosinski, and J. Clune. Deep neural networks are easily fooled: High confidence predictions for unrecognizable images. In *Computer Vision and Pattern Recognition (CVPR), 2015 IEEE Conference on*. IEEE, 2015.
- A. Nguyen, J. Clune, Y. Bengio, A. Dosovitskiy, and J. Yosinski. Plug & play generative networks: Conditional iterative generation of images in latent space. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*. IEEE, 2017.
- D. Nguyen, N. A. Smith, and C. P. Rosé. Author age prediction from text using linear regression. In *Proceedings of the 5th ACL-HLT Workshop on Language Technology for Cultural Heritage, Social Sciences, and Humanities*, pages 115–123, Portland, OR, USA, June 2011. Association for Computational Linguistics. URL <https://www.aclweb.org/anthology/W11-1515>.
- D. Nguyen, D. Trieschnigg, A. S. Doğruöz, R. Gravel, M. Theune, T. Meder, and F. De Jong. Why gender and age prediction from tweets is hard: Lessons from a crowdsourcing experiment. In *Proceedings of COLING 2014, the 25th International Conference on Computational Linguistics: Technical Papers*, pages 1950–1961, 2014.
- J. W. Pennebaker and L. D. Stone. Words of wisdom: Language use over the life span. *Journal of Personality and Social Psychology*, 85(2):291–301, 2003. URL <https://doi.org/10.1037/0022-3514.85.2.291>.
- S. Prabhumoye, A. W. Black, and R. Salakhutdinov. Exploring controllable text generation techniques. In *Proceedings of the 28th International Conference on Computational Linguistics*,

- pages 1–14, Barcelona, Spain (Online), Dec. 2020. International Committee on Computational Linguistics. doi: 10.18653/v1/2020.coling-main.1. URL <https://aclanthology.org/2020.coling-main.1>.
- Q. Qian, M. Huang, H. Zhao, J. Xu, and X. Zhu. Assigning personality/profile to a chatting machine for coherent conversation generation. In *Proceedings of the Twenty-Seventh International Joint Conference on Artificial Intelligence, IJCAI-18*, pages 4279–4285. International Joint Conferences on Artificial Intelligence Organization, 7 2018. doi: 10.24963/ijcai.2018/595. URL <https://doi.org/10.24963/ijcai.2018/595>.
- A. Radford, K. Narasimhan, T. Salimans, and I. Sutskever. Improving language understanding by generative pre-training. 2018.
- A. Radford, J. Wu, R. Child, D. Luan, D. Amodei, and I. Sutskever. Language models are unsupervised multitask learners. *OpenAI blog*, 1(8):9, 2019.
- M. Scheutz, R. Cantrell, and P. Schermerhorn. Toward humanlike task-based dialogue processing for human robot interaction. *Ai Magazine*, 32(4):77–84, 2011.
- J. Schler, M. Koppel, S. Argamon, and J. W. Pennebaker. Effects of age and gender on blogging. In *AAAI spring symposium: Computational approaches to analyzing weblogs*, volume 6, pages 199–205, 2006.
- M. Schuster and K. K. Paliwal. Bidirectional recurrent neural networks. *IEEE transactions on Signal Processing*, 45(11):2673–2681, 1997.
- R. Sennrich, B. Haddow, and A. Birch. Neural machine translation of rare words with subword units. In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1715–1725, Berlin, Germany, Aug. 2016. Association for Computational Linguistics. doi: 10.18653/v1/P16-1162. URL <https://www.aclweb.org/anthology/P16-1162>.
- F. Stahlberg, J. Cross, and V. Stoyanov. Simple fusion: Return of the language model. In *Proceedings of the Third Conference on Machine Translation: Research Papers*, pages 204–211, Brussels, Belgium, Oct. 2018. Association for Computational Linguistics. doi: 10.18653/v1/W18-6321. URL <https://www.aclweb.org/anthology/W18-6321>.
- S. A. Tagliamonte and D. Denis. Linguistic ruin? lol! instant messaging and teen language. *American speech*, 83(1):3–34, 2008.
- A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin. Attention is all you need. In *Advances in neural information processing systems*, pages 5998–6008, 2017.
- J. Vig. A multiscale visualization of attention in the transformer model. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics: System Demonstrations*, pages 37–42, Florence, Italy, July 2019. Association for Computational Linguistics. doi: 10.18653/v1/P19-3007. URL <https://www.aclweb.org/anthology/P19-3007>.
- S. Welleck, J. Weston, A. Szlam, and K. Cho. Dialogue natural language inference. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 3731–3741, Florence, Italy, July 2019. Association for Computational Linguistics. doi: 10.18653/v1/P19-1363. URL <https://aclanthology.org/P19-1363>.
- S. Wiegreffe and Y. Pinter. Attention is not not explanation. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 11–20, Hong Kong, China, Nov. 2019. Association for Computational Linguistics. doi: 10.18653/v1/D19-1002. URL <https://aclanthology.org/D19-1002>.
- M. Wolters, R. Vipperla, and S. Renals. Age recognition for spoken dialogue systems: Do we need it? In *Tenth Annual Conference of the International Speech Communication Association (Interspeech)*, 2009.

- G. Zeng, W. Yang, Z. Ju, Y. Yang, S. Wang, R. Zhang, M. Zhou, J. Zeng, X. Dong, R. Zhang, H. Fang, P. Zhu, S. Chen, and P. Xie. MedDialog: Large-scale medical dialogue datasets. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 9241–9250, Online, Nov. 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.emnlp-main.743. URL <https://www.aclweb.org/anthology/2020.emnlp-main.743>.
- S. Zhang, E. Dinan, J. Urbanek, A. Szlam, D. Kiela, and J. Weston. Personalizing dialogue agents: I have a dog, do you have pets too? In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 2204–2213, Melbourne, Australia, July 2018a. Association for Computational Linguistics. doi: 10.18653/v1/P18-1205. URL <https://aclanthology.org/P18-1205>.
- Y. Zhang, S. Sun, M. Galley, Y.-C. Chen, C. Brockett, X. Gao, J. Gao, J. Liu, and B. Dolan. Dialogpt: Large-scale generative pre-training for conversational response generation. In *ACL, system demonstration*, 2020.
- Z. Zhang, S. Ren, S. Liu, J. Wang, P. Chen, M. Li, M. Zhou, and E. Chen. Style transfer as unsupervised machine translation. *CoRR*, abs/1808.07894, 2018b. URL <http://arxiv.org/abs/1808.07894>.
- Y. Zheng, G. Chen, M. Huang, S. Liu, and X. Zhu. Personalized dialogue generation with diversified traits. *arXiv preprint arXiv:1901.09672*, 2019.

Appendix A

Supplementary material

A.1 Prompts used for controlled dialogue response generation

Prompt	Average P_Y	Average P_O	Standard deviation
Young prompts			
<i>What are your plans this week?</i>	0.938	0.062	0.056
<i>What do you wanna eat?</i>	0.958	0.042	0.026
<i>Do you have any hobbies?</i>	0.963	0.037	0.021
<i>Can I add you on Facebook?</i>	0.995	0.005	0.005
<i>Awesome! I actually haven't been there. When did you go?</i>	0.995	0.005	0.004
Neutral prompts			
<i>Good weather we're having.</i>	0.489	0.511	0.133
<i>Can we talk?</i>	0.412	0.588	0.116
<i>Hi, how's it going?</i>	0.437	0.563	0.079
<i>Hello, tell me about your latest holiday.</i>	0.496	0.504	0.059
<i>Hey.</i>	0.574	0.426	0.046
Old prompts			
<i>Tell me about your family.</i>	0.061	0.939	0.050
<i>Good afternoon.</i>	0.040	0.960	0.044
<i>I had a splendid weekend.</i>	0.030	0.970	0.034
<i>Hello, how are you?</i>	0.041	0.959	0.031
<i>Hello, tell me about yourself.</i>	0.025	0.975	0.026

Table A.1: Young (red rows), neutral (yellow rows), and old (blue rows) prompts used for controlled dialogue generation in Experiment 2. Average P_Y and P_O are the averaged young and old probabilities assigned to the prompts by three initializations of BERT_{FT}.

A.2 Wordlists for BoW-based Approaches

[L: TODO - Censor foul language?]

100 Most Informative Unigrams - Young (19-29) um, cool, shit, hmm, uni, cute, tut, massive, awesome, gym, bitch, lol, grand, pizza, like, excited, yawn, Korea, cigarette, fuck, fairness, Jesus, annoying, Facebook, quicker, definitely, guess, Sunderland, oo, wanna, mountain, scared, piss, love, miss, Middlesbrough, mhm, specifically, ooh, website, roundabout, photo, nope, blanket, management, ridiculous, mental, pregnant, beers, hate, log, fucking, cry, cheaper, skinny, plural, burger, hilarious, hint, drunk, fridge, cousin, coke, genuinely, James, mates, smaller, option, balance, saving, basically, leather, nev, shut, frig, mate, yay, invite, maid, nickname, badly, garlic, CD, jokes, Uzbekistan, boyfriend, date, added, Manchester, blah, shitty, lang, tempted, stadium, wee, eh, baking, city, honestly, exam

100 Most Informative Unigrams - Old (50 plus) ordinary, Chinese, wonderful, yes, tend, father, photographs, vegetables, hospice, operation, shed, pension, areas, mother, hanging, hospices, glasses, chap, anyhow, tank, surgery, container, cheers, born, church, pain, several, workshop, right, horses, building, extraordinary, vegetarian, biscuit, americano, engine, luck, paint, emperor, lipsy, trombone, occasional, supper, lord, architect, council, roast, schools, bath, asbestos, endometrial, concrete, poodle, recall, diabetes, misty, report, heavens, enormous, lawn, potatoes, email, junk, scabies, mousse, Ebola, churches, sewing, plants, rackets, marmalade, engineering, furniture, photograph, sandwiches, unemployment, xylophone, Piccadilly, flu, claim, arab, nineteen, forgotten, sensible, blancmange, spencer, yards, emails, yellow, scruffy, fungi, garden, boiler, lodge, mostly, Robson, tricky, shark, robin, contracture

Frequency-based Young (19-29) um, shit, cool, fucking, definitely, guess, friends, everyone, literally, dad, sounds, weekend, loads, watch, fair, fuck, amazing, friend, ha, huh, hate, fun, stay, girl, holiday, blah, hours, uni, month, horrible, massive, Friday, stupid, film, parents, thirty, spend, mate, honest, change, hope, yourself, annoying, wear, wait, ridiculous, anyone, Saturday, tea, dinner, sit, crazy, hell, pound, nine, expensive

Frequency-based Old (50 plus) building, may, water, mother, perhaps, door, lots, business, cancer, area, although, worked, open, cut, number, under, young, nineteen, everybody, garden, church, case, shop, children, certainly, set, coffee, email, gave, white, along, doctor, hear, often, possibly, group, father, outside, wonderful, taken, seem, places, green, given, hand, early, women, space, front, language, dear, light, huge, supposed, country, hospital, otherwise, asked, putting, bits, gosh, wall, woman, almost, particularly, across, word, age, rest, flat, turned, decided,

finished, needed, red, bin, hospice, running, slightly, its, middle, local, percent, Chinese, paper, check, high, milk, piece, near, nobody, usually

A.3 Supplementary Material about the Effects of Prompt Style on PPLM Output

A.4 Where to put these?

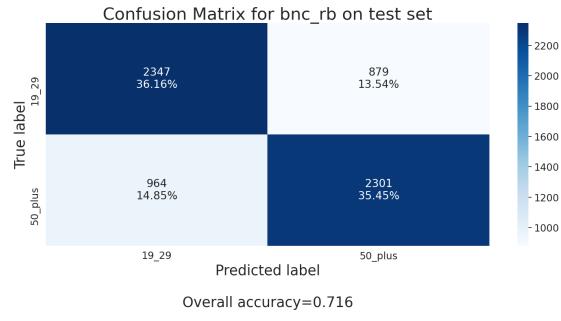


Figure A.1: Confusion matrix BERT age classifier on balanced BNC **test** set.

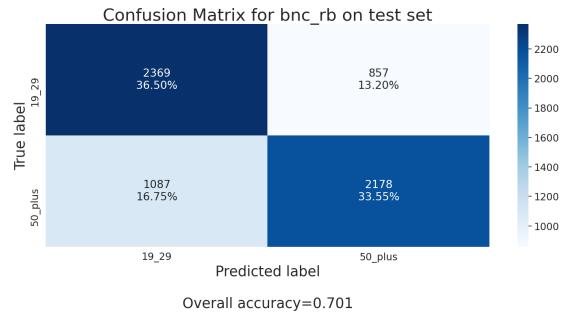


Figure A.2: Confusion matrix LSTM age classifier on balanced BNC **test** set.

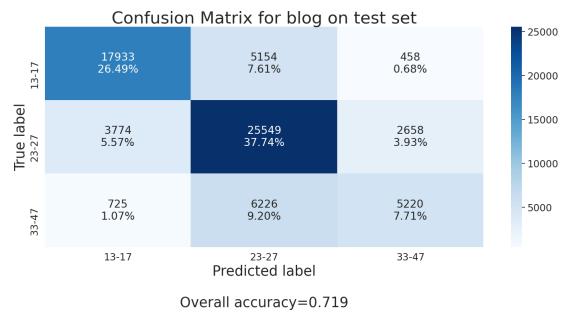


Figure A.3: Confusion matrix bi-LSTM age classifier on blog corpus **test** set.

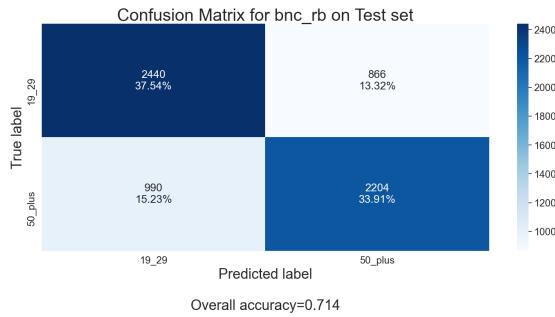


Figure A.4: Confusion matrix for best trigram age classifier on **balanced BNC test set**.

A.5 Age Discrimination on the Imbalanced British National Corpus [L: Do we even need these plots?]

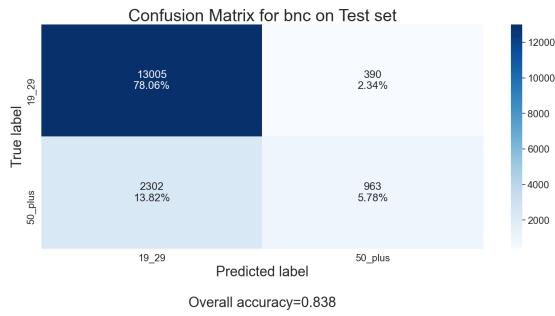


Figure A.5: Confusion matrix for best bigram age classifier on BNC test set.

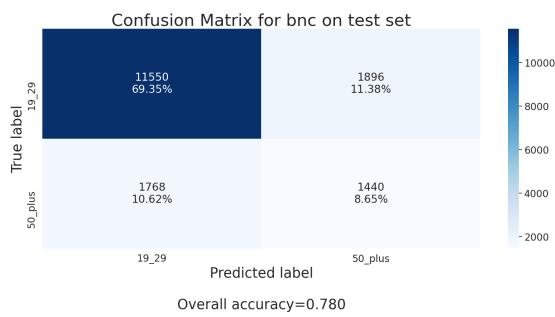


Figure A.6: Confusion matrix bi-LSTM age classifier on BNC test set.

19-29		50+	
coef.	n-gram	coef.	n-gram
-3.19	um	2.29	yes
-2.91	cool	2.21	wonderful
-2.70	s***t	1.91	building
-2.25	cute	1.86	right right
-2.15	uni	1.80	something like
-2.14	hmm	1.73	garden
-1.97	wanna	1.69	right
-1.93	f***k	1.68	ordinary
-1.91	like	1.67	shed
-1.85	massive	1.63	operation
-1.83	yeah course	1.58	born
-1.81	love	1.57	mother
-1.79	tut	1.55	photographs
-1.74	b***h	1.51	email
-1.68	like oh	1.08	anything like

Table A.2: [L: Excluding stopwords.] For each age group, top 15 most informative n -grams used by the trigram model. **coef.** is the coefficient (and sign) of the corresponding n -gram for the logistic regression model: the higher its absolute value, the higher the utterance's odds to belong to one age group. * indicates masking of foul language.

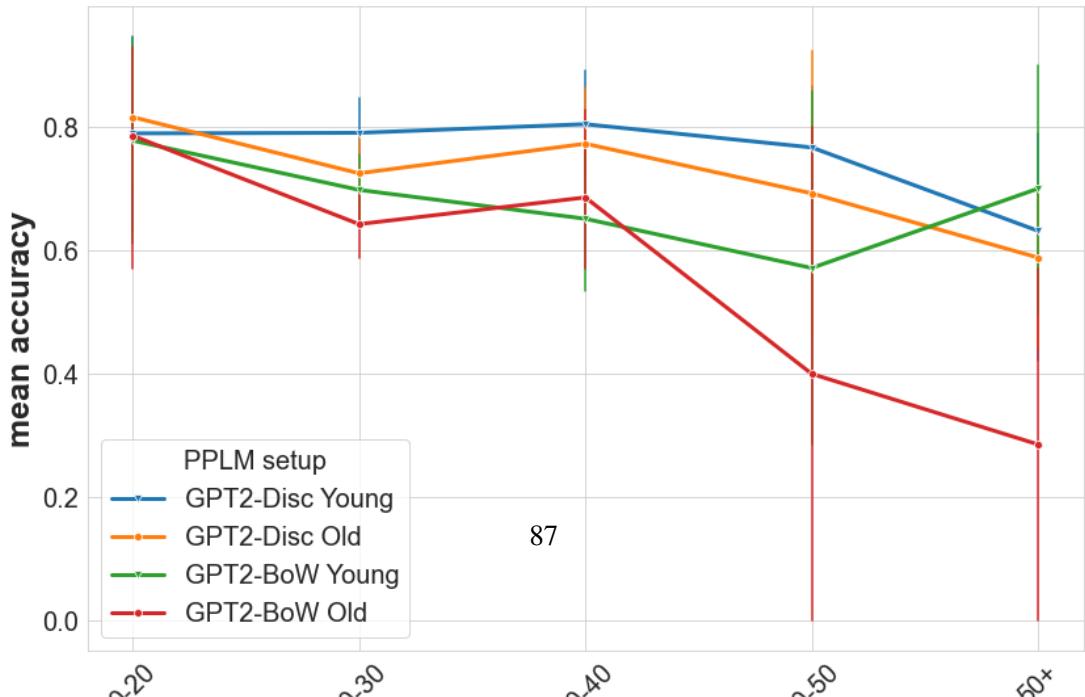
Model	Accuracy ↑ better	$F_1^{(13-17)}$ ↑ better	$F_1^{(23-27)}$ ↑ better	$F_1^{(33+)}$ ↑ better
Majority class	0.472	*	0.642	*
Schler et al. [2006]	0.762	0.860	0.748	0.504
unigram	0.601 (0.001)	0.764 (0.001)	0.704 (0.001)	0.498 (0.003)
bigram	0.625 (0.001)	0.790 (0.001)	0.712 (0.001)	0.518 (0.001)
trigram	0.623 (0.001)	0.790 (0.001)	0.712 (0.002)	0.498 (0.002)
LSTM	0.663 (0.005)	0.748 (0.003)	0.664 (0.010)	0.502 (0.004)
BiLSTM	0.618 (0.008)	0.732 (0.003)	0.579 (0.016)	0.509 (0.004)
$BERT_{frozen}$	0.623 (0.002)	0.658 (0.006)	0.678 (0.007)	0.256 (0.041)
$BERT_{FT}$	0.742 (0.010)	0.813 (0.007)	0.749 (0.013)	0.592 (0.009)

Table A.3: Discourse dataset. [L: Including stopwords] Test set results averaged over 5 random initializations. Format: *average metric (standard error)*. Values in **bold** are the highest in the column; in **blue**, the second highest. *: F_1 is actually 0/0.

Model	Accuracy ↑ better	$F_1^{(19-29)}$ ↑ better	$F_1^{(50+)}$ ↑ better
Random	0.500	0.500	0.500
unigram	0.702 (0.006)	0.713 (0.006)	0.690 (0.006)
bigram	0.703 (0.006)	0.713 (0.005)	0.693 (0.008)
trigram	0.709 (0.007)	0.718 (0.007)	0.700 (0.008)
LSTM	0.696 (0.005)	0.689 (0.018)	0.701 (0.016)
BiLSTM	0.684 (0.007)	0.688 (0.018)	0.679 (0.016)
$BERT_{frozen}$	0.673 (0.005)	0.679 (0.013)	0.667 (0.018)
$BERT_{FT}$	0.710 (0.006)	0.717 (0.007)	0.703 (0.014)

Table A.4: Dialogue dataset [L: Excluding stopwords]. Test set results averaged over 5 random initializations. Format: *average metric (standard error)*. Values in **bold** are the highest in the column; in **blue**, the second highest.

A.6 Placeholders for Final CTG Results Tables

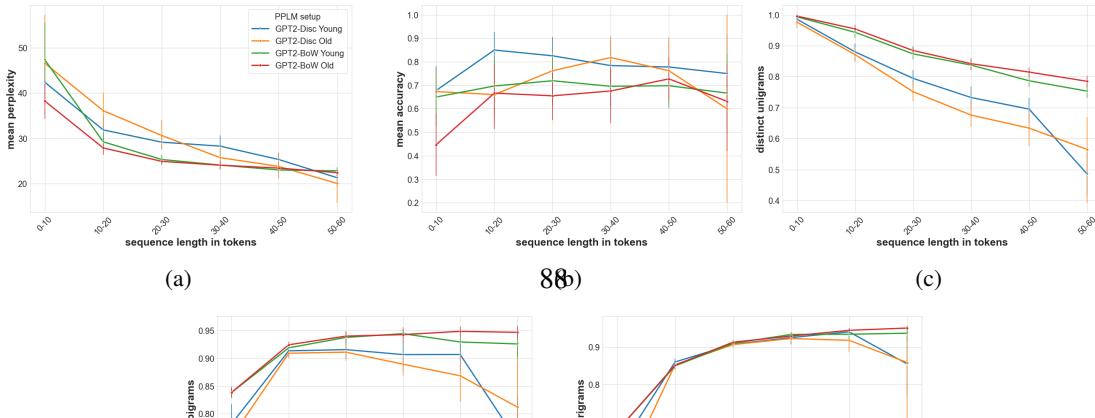


Model	ppl. ↓ better	Dist-1 ↑ better	Dist-2 ↑ better	Dist-3 ↑ better	Acc. ↑ better
Baseline***	27.45 (± 7.27)	0.90 (± 0.10)	0.92 (± 0.05)	0.86 (± 0.09)	57.5%
B_{100MCW}^{***}	26.68 (± 8.77)	0.89 (± 0.10)	0.92 (± 0.05)	0.86 (± 0.09)	51.7%
$B_{Y,FB}$	27.11 (± 7.45)	0.91 (± 0.09)	0.92 (± 0.04)	0.87 (± 0.09)	68.3%
$B_{O,FB}$	25.99 (± 6.41)	0.88 (± 0.11)	0.92 (± 0.05)	0.86 (± 0.09)	62.5%
$B_{Y,100MIU}$	28.48 (± 11.96)	0.88 (± 0.12)	0.91 (± 0.06)	0.86 (± 0.10)	69.2%
$B_{O,100MIU}$	25.57 (± 7.44)	0.88 (± 0.11)	0.92 (± 0.05)	0.87 (± 0.09)	58.3%
$D_{Y,GPT2}$	33.02 (± 12.24)	0.85 (± 0.16)	0.89 (± 0.07)	0.83 (± 0.12)	73.9%
$D_{O,GPT2}$	32.86 (± 18.08)	0.80 (± 0.21)	0.84 (± 0.13)	0.79 (± 0.19)	63.3%

Table A.5: [L: Excluding stopwords.] Results of age-controlled language generation. Perplexity is perplexity w.r.t. GPT-1. Dist-n is number of distinct n-grams normalized by text length, as a measure of diversity. Young and old accuracy are the assigned probabilities of belonging to the young or old age categories.

Model	ppl. ↓ better	Dist-1 ↑ better	Dist-2 ↑ better	Dist-3 ↑ better	\bar{P}_Y	\bar{P}_O	Acc. ↑ better
G-baseline	27.50 (± 6.58)	0.87 (± 0.09)	0.94 (± 0.04)	0.90 (± 0.06)	0.62 (± 0.42)	0.38	-
G-100MCW	27.56 (± 6.60)	0.86 (± 0.10)	0.93 (± 0.04)	0.90 (± 0.05)	0.63 (± 0.42)	0.37	-
G-B _{FB,Y}	27.91 (± 7.18)	0.87 (± 0.10)	0.93 (± 0.05)	0.90 (± 0.06)	0.69 (± 0.41)	0.31	70.4%
G-B _{FB,O}	27.58 (± 7.07)	0.86 (± 0.10)	0.93 (± 0.04)	0.90 (± 0.06)	0.58 (± 0.42)	0.42	43.0%
G-B _{100MIU,Y}	28.37 (± 7.31)	0.87 (± 0.09)	0.93 (± 0.04)	0.90 (± 0.06)	0.67 (± 0.41)	0.33	67.4%
G-B _{100MIU,O}	27.25 (± 6.15)	0.87 (± 0.09)	0.93 (± 0.04)	0.90 (± 0.06)	0.62 (± 0.42)	0.38	37.4%
G-D _Y	32.09 (± 18.98)	0.77 (± 0.20)	0.86 (± 0.13)	0.84 (± 0.15)	0.66 (± 0.43)	0.34	67.8%
G-D _O	47.15 (± 47.56)	0.73 (± 0.24)	0.75 (± 0.28)	0.75 (± 0.27)	0.24 (± 0.36)	0.76	74.3%
D-baseline	37.52 (± 12.06)	0.86 (± 0.13)	0.90 (± 0.08)	0.85 (± 0.10)	0.76 (± 0.37)	0.24	-
D-100MCW	37.80 (± 10.89)	0.85 (± 0.14)	0.89 (± 0.10)	0.85 (± 0.10)	0.82 (± 0.33)	0.18	-
D-B _{FB,Y}	38.53 (± 12.64)	0.87 (± 0.12)	0.90 (± 0.08)	0.86 (± 0.10)	0.82 (± 0.33)	0.18	83.0%
D-B _{FB,O}	37.85 (± 11.17)	0.87 (± 0.12)	0.90 (± 0.08)	0.86 (± 0.09)	0.78 (± 0.35)	0.22	21.5%
D-B _{100MIU,Y}	38.67 (± 11.70)	0.88 (± 0.11)	0.91 (± 0.07)	0.86 (± 0.10)	0.87 (± 0.28)	0.13	88.5%
D-B _{100MIU,O}	37.91 (± 12.27)	0.87 (± 0.11)	0.90 (± 0.07)	0.85 (± 0.10)	0.79 (± 0.34)	0.22	21.9%
D-D _Y	42.01 (± 16.94)	0.90 (± 0.12)	0.86 (± 0.14)	0.77 (± 0.22)	0.86 (± 0.29)	0.14	85.9%
D-D _O	41.17 (± 20.72)	0.87 (± 0.12)	0.89 (± 0.13)	0.83 (± 0.16)	0.43 (± 0.41)	0.57	56.7%

Table A.6: [L: Neutral prompt] Results of age-controlled language generation. Perplexity is perplexity w.r.t. GPT-1. Dist-n is number of distinct n-grams normalized by text length, as a measure of diversity. \bar{P}_Y and \bar{P}_O are the respective average young and old probabilities assigned by the best BERT_{FT}. Acc. is the best BERT model’s accuracy when classifying the row’s samples.



Model	ppl. ↓ better	Dist-1 ↑ better	Dist-2 ↑ better	Dist-3 ↑ better	\bar{P}_Y	\bar{P}_O	Acc. ↑ better
G-baseline	28.05 (± 6.12)	0.85 (± 0.13)	0.91 (± 0.08)	0.88 (± 0.08)	0.80 (± 0.33)	0.20	-
G-100MCW	27.71 (± 6.20)	0.85 (± 0.12)	0.91 (± 0.09)	0.88 (± 0.09)	0.75 (± 0.37)	0.25	-
G-B _{FB,Y}	28.81 (± 7.09)	0.86 (± 0.12)	0.92 (± 0.08)	0.89 (± 0.08)	0.82 (± 0.32)	0.18	83.3%
G-B _{FB,O}	28.54 (± 6.45)	0.86 (± 0.12)	0.92 (± 0.08)	0.89 (± 0.08)	0.77 (± 0.36)	0.23	22.6%
G-B _{100MIU,Y}	28.49 (± 6.49)	0.86 (± 0.12)	0.91 (± 0.08)	0.88 (± 0.08)	0.83 (± 0.32)	0.17	83.0%
G-B _{100MIU,O}	28.18 (± 5.70)	0.87 (± 0.11)	0.92 (± 0.08)	0.89 (± 0.09)	0.79 (± 0.34)	0.21	21.5%
G-D _Y	39.32 (± 37.49)	0.84 (± 0.21)	0.61 (± 0.40)	0.57 (± 0.40)	0.70 (± 0.40)	0.30	70.7%
G-D _O	85.40 (± 150.28)	0.67 (± 0.30)	0.62 (± 0.31)	0.62 (± 0.32)	0.29 (± 0.40)	0.71	70.5%
D-baseline	36.69 (± 9.11)	0.87 (± 0.10)	0.91 (± 0.06)	0.87 (± 0.08)	0.90 (± 0.24)	0.10	-
D-100MCW	36.93 (± 9.18)	0.86 (± 0.11)	0.91 (± 0.06)	0.88 (± 0.07)	0.90 (± 0.25)	0.10	-
D-B _{FB,Y}	37.35 (± 8.60)	0.88 (± 0.10)	0.91 (± 0.06)	0.87 (± 0.08)	0.90 (± 0.26)	0.10	90.0%
D-B _{FB,O}	37.25 (± 9.45)	0.87 (± 0.11)	0.91 (± 0.06)	0.87 (± 0.08)	0.88 (± 0.29)	0.12	11.1%
D-B _{100MIU,Y}	37.87 (± 8.32)	0.88 (± 0.10)	0.91 (± 0.07)	0.87 (± 0.09)	0.91 (± 0.24)	0.09	92.6%
D-B _{100MIU,O}	37.04 (± 8.78)	0.88 (± 0.10)	0.91 (± 0.05)	0.88 (± 0.07)	0.85 (± 0.32)	0.15	15.2%
D-D _Y	39.22 (± 14.96)	0.89 (± 0.12)	0.86 (± 0.19)	0.79 (± 0.23)	0.89 (± 0.25)	0.11	91.1%
D-D _O	38.46 (± 14.91)	0.82 (± 0.15)	0.87 (± 0.15)	0.83 (± 0.17)	0.52 (± 0.44)	0.48	47.4%

Table A.7: [L: Young prompt] Results of age-controlled language generation. Perplexity is perplexity w.r.t. GPT-1. Dist-n is number of distinct n-grams normalized by text length, as a measure of diversity. Acc. is the best BERT model’s accuracy when classifying the row’s samples.

A.7 CTG Results for Unprompted Setup [L: Redundant?]

Table A.10 reports the automated evaluation results of our controllable text generation models. It can be seen that uncontrolled GPT-2 baseline has a slight bias towards generating "young-sounding" language (57.5% accuracy). Furthermore, it appears that perturbing GPT-2’s output distribution with the 100 most common words across all ages results in a slight de-biasing of the generated text (54.1% accuracy). Achieving detectable control seems possible, because all GPT-2-based models surpass both baselines in terms of accuracy, with the exception of both BoW-setups using the 100 most informative unigrams.

Frequency-based BoW-models outperform those using the most informative unigrams, as illustrated by their higher average accuracy (66.75% versus. 53.1%), and lower average perplexity (27.48 versus 27.90). Discriminator-based models achieve noticeably better accuracies, with an average improvement of 8.45% over the best performing BoW-based models. However, discriminator-based models do show more signs of disfluency and repetitiveness compared to the BoW-models, as depicted by the worse perplexities and $\text{Dist-}n|_{n=1,2,3}$ scores.

Model	ppl. ↓ better	Dist-1 ↑ better	Dist-2 ↑ better	Dist-3 ↑ better	\bar{P}_Y	\bar{P}_O	Acc. ↑ better
G-baseline	29.34 (± 10.30)	0.86 (± 0.09)	0.94 (± 0.04)	0.90 (± 0.06)	0.60 (± 0.43)	0.40	-
G-100MCW	29.14 (± 10.11)	0.86 (± 0.10)	0.93 (± 0.04)	0.90 (± 0.06)	0.60 (± 0.44)	0.40	-
G- B_Y,FB	29.61 (± 10.28)	0.86 (± 0.10)	0.93 (± 0.04)	0.91 (± 0.06)	0.62 (± 0.43)	0.38	61.1%
G- B_O,FB	28.81 (± 10.10)	0.86 (± 0.10)	0.93 (± 0.05)	0.90 (± 0.06)	0.59 (± 0.43)	0.41	41.1%
G- $B_{Y,100MIU}$	29.51 (± 0.09)	0.87 (± 0.09)	0.93 (± 0.05)	0.90 (± 0.06)	0.68 (± 0.42)	0.32	68.5%
G- $B_{O,100MIU}$	29.05 (± 9.80)	0.86 (± 0.09)	0.93 (± 0.04)	0.90 (± 0.06)	0.60 (± 0.43)	0.40	39.6%
G-D $_Y$	32.34 (± 19.88)	0.77 (± 0.20)	0.84 (± 0.19)	0.80 (± 0.23)	0.65 (± 0.43)	0.35	65.4%
G-D $_O$	95.21 (± 174.42)	0.65 (± 0.27)	0.78 (± 0.18)	0.78 (± 0.18)	0.10 (± 0.25)	0.90	90.3%
D-baseline	38.18 (± 12.03)	0.86 (± 0.12)	0.90 (± 0.08)	0.86 (± 0.09)	0.72 (± 0.38)	0.28	-
D-100MCW	37.73 (± 11.88)	0.85 (± 0.13)	0.90 (± 0.08)	0.86 (± 0.09)	0.73 (± 0.39)	0.27	-
D- B_Y,FB	38.24 (± 11.53)	0.86 (± 0.12)	0.90 (± 0.08)	0.86 (± 0.10)	0.81 (± 0.34)	0.19	82.6%
D- B_O,FB	37.8 (± 11.74)	0.86 (± 0.12)	0.90 (± 0.07)	0.87 (± 0.08)	0.72 (± 0.39)	0.28	29.3%
D- $B_{Y,100MIU}$	38.66 (± 11.57)	0.85 (± 0.12)	0.90 (± 0.07)	0.86 (± 0.09)	0.81 (± 0.33)	0.19	80.7%
D- $B_{O,100MIU}$	36.93 (± 11.68)	0.87 (± 0.12)	0.90 (± 0.09)	0.86 (± 0.09)	0.69 (± 0.41)	0.31	29.6%
D-D $_Y$	42.93 (± 20.18)	0.90 (± 0.14)	0.79 (± 0.22)	0.68 (± 0.28)	0.84 (± 0.30)	0.16	85.2%
D-D $_O$	45.58 (± 38.59)	0.86 (± 0.14)	0.86 (± 0.13)	0.79 (± 0.20)	0.40 (± 0.42)	0.60	59.7%

Table A.8: [L: Old prompt] Results of age-controlled language generation. Perplexity is perplexity w.r.t. GPT-1. Dist-n is number of distinct n-grams normalized by text length, as a measure of diversity. Acc. is the best BERT model’s accuracy when classifying the row’s samples.

The accuracies of our uncontrolled DialoGPT baseline (78.1%) and the 100MCW baseline (80.7%), suggest that DialoGPT is heavily biased towards producing young-sounding language. This can be attributable to DialoGPT having been fine-tuned on Reddit threads, as the majority of Reddit users are between the ages 20 and 29¹ [Zhang et al., 2020]. DialoGPT’s strong propensity for generating younger sounding language makes it a less desirable choice for our human evaluation experiments, because it requires non-standard parameter settings to produce detectably older sounding text.

Overall, the results show that, for most models, a plug-and-play approach to controlling generated dialogue responses to possess detectable age-specific linguistic features is achievable. The most promising models being either discriminator-based, or frequency-based bag-of-words models. Discriminator-based models achieve more detectable levels of control than their BoW-based counterparts, at the cost of perplexity and repetitiveness. This could be attributable to the more complex activation-space updates that are used by discriminator-models. Furthermore, GPT-2’s preference to generate young-sounding language is severely less pronounced than that of DialoGPT, making it easier to control, given equal parameter settings.

¹<https://www.statista.com/statistics/1125159/reddit-us-app-users-age/>

Model	ppl. ↓ better	Dist-1 ↑ better	Dist-2 ↑ better	Dist-3 ↑ better	Acc. ↑ better
GPT-2 baseline	... (±...)	... (±...)	... (±...)	... (±...)	... %
GPT-2 100MCW baseline	... (±...)	... (±...)	... (±...)	... (±...)	... %
GPT-2 $B_{Y,FB}$... (±...)	... (±...)	... (±...)	... (±...)	... %
GPT-2 $B_{O,FB}$... (±...)	... (±...)	... (±...)	... (±...)	... %
GPT-2 $B_{Y,100MIU}$... (±...)	... (±...)	... (±...)	... (±...)	... %
GPT-2 $B_{O,100MIU}$... (±...)	... (±...)	... (±...)	... (±...)	... %
GPT-2 D_Y	... (±...)	... (±...)	... (±...)	... (±...)	... %
GPT-2 D_O	... (±...)	... (±...)	... (±...)	... (±...)	... %
DGPT baseline	... (±...)	... (±...)	... (±...)	... (±...)	... %
DGPT-100MCW	... (±...)	... (±...)	... (±...)	... (±...)	... %
DGPT $B_{Y,FB}$... (±...)	... (±...)	... (±...)	... (±...)	... %
DGPT $B_{O,FB}$... (±...)	... (±...)	... (±...)	... (±...)	... %
DGPT $B_{Y,100MIU}$... (±...)	... (±...)	... (±...)	... (±...)	... %
DGPT $B_{O,100MIU}$... (±...)	... (±...)	... (±...)	... (±...)	... %
DGPT $_Y$... (±...)	... (±...)	... (±...)	... (±...)	... %
DGPT $_O$... (±...)	... (±...)	... (±...)	... (±...)	... %

Table A.9: [L: Unprompted] Results of age-controlled language generation. Perplexity is perplexity w.r.t. GPT-1. Dist-n is number of distinct n-grams normalized by text length, as a measure of diversity. Acc. is the best BERT model’s accuracy when classifying the row’s samples.

Model	ppl. ↓ better	Dist-1 ↑ better	Dist-2 ↑ better	Dist-3 ↑ better	Acc. ↑ better
GPT-2 baseline	29.60 (± 17.57)	0.89 (± 0.09)	0.93 (± 0.05)	0.88 (± 0.10)	57.5%
GPT-2 100MCW baseline	27.80 (± 16.44)	0.83 (± 0.12)	0.91 (± 0.06)	0.88 (± 0.09)	54.1%
$B_{Y,FB}$	28.16 (± 14.52)	0.87 (± 0.10)	0.92 (± 0.06)	0.88 (± 0.10)	69.3%
$B_{O,FB}$	26.79 (± 8.89)	0.88 (± 0.09)	0.92 (± 0.05)	0.88 (± 0.10)	64.2%
$B_{Y,100MIU}$	29.16 (± 14.91)	0.89 (± 0.09)	0.92 (± 0.06)	0.88 (± 0.11)	52.5%
$B_{O,100MIU}$	26.63 (± 8.36)	0.87 (± 0.10)	0.92 (± 0.07)	0.88 (± 0.10)	53.7%
$D_{Y,GPT2}$	31.95 (± 14.29)	0.82 (± 0.17)	0.87 (± 0.14)	0.83 (± 0.16)	77.7%
$D_{O,GPT2}$	33.63 (± 24.40)	0.80 (± 0.18)	0.87 (± 0.11)	0.81 (± 0.21)	72.7%
DGPT baseline	35.20 (± 10.01)	0.87 (± 0.11)	0.90 (± 0.07)	0.87 (± 0.08)	78.1%
DGPT-100MCW	35.64 (± 9.72)	0.86 (± 0.10)	0.90 (± 0.06)	0.87 (± 0.08)	80.7%
$D^*_{Y,DGPT}$	41.54 (± 10.87)	0.91 (± 0.11)	0.91 (± 0.06)	0.86 (± 0.09)	84.1%
$D^*_{O,DGPT}$	38.16 (± 10.77)	0.87 (± 0.11)	0.91 (± 0.06)	0.87 (± 0.08)	55.6%

Table A.10: [L: Unprompted] Results of age-controlled language generation. Perplexity is perplexity w.r.t. GPT-1. Dist-n is number of distinct n-grams normalized by text length, as a measure of diversity. Acc. is the best BERT model’s accuracy when classifying the row’s samples.